


FLAGSHIP INSTITUTIONAL STUDY

Can Global Investors Actually Capture India's Economic Rise?

India, emerging markets and the competition for global capital



Independent India and emerging-markets research for global capital allocators. Defined questions on India's economy and markets, answered in writing, with the reasoning and the data shown.

PUBLICATION

Flagship Institutional Study
India & Emerging-Markets Strategy
August 2026

DATA CUT-OFF

Core market and macro data to 31
July 2026; post-cut-off develop-
ments labelled where used.

EVIDENCE STANDARD

Every material figure carries a
confidence class — **GREEN** primary,
AMBER secondary/derived, **RED**
excluded from headlines.

CONTACT

inquiries@kakderesearch.com · kakderesearch.com

This document is independent research for professional and institutional investors. It is not investment advice, an offer, or a solicitation, and it contains no buy or sell recommendation. Figures are sourced and dated; derived and illustrative values are labelled as such. Past performance does not indicate future results.

Contents

	The Investment-Committee Page	4
	Executive Summary	7
1	India, Seen From the Benchmark	17
2	The Growth Machine	25
3	The State, Infrastructure and Formalisation	31
4	Can India Manufacture Its Next Leap?	37
5	The Indian Consumer	43
6	From GDP to Corporate Profit	48
7	The India a Shareholder Can Own	56
8	India's Financial Machine	66
9	The Price of the Story	73
10	The Rupee, Oil and the Outside World	81
11	India Against the Alternatives	88
12	The Allocator's Answer	96
A	Appendix A — Methodology	105
B	Appendix B — Data	108
C	Appendix C — Bibliography & Research Literature	110
D	Appendix D — Company Compendium	112

The Investment-Committee Page

If you read nothing else, read this. Everything below is defended in the body; nothing here is a rating.

11.66%

INDIA'S MSCI EM WEIGHT

Down from ~20% (2024). The round-trip is the story.

8.03% vs 9.19%

MSCI INDIA VS MSCI EM · 10Y USD NET

India lost to the index it stars in.

-3.8%/yr

THE RUPEE'S DRAG

The 20-year bridge's FX leg: it converted a world-class earnings machine into a below-EM dollar return.

36.0 vs ~24

INDIA CAPE VS LONG-RUN MEAN

A mean-reversion headwind, not a tailwind.

5.2%

NIFTY-500 PROFIT/GDP, FY26

An all-time high. The engine is real *and* the base is elevated.

4.6×

FY25 DOMESTIC ABSORPTION

DII buying outran FPI selling (+₹6.07L cr vs -₹1.31L cr) — the deepest domestic buyer base in EM.

THE POSITION

Hold India near benchmark; earn the overweight, do not assume it. Funded from developed-market or cash sleeves, not from Taiwan or Korea.

India is **structurally attractive, selectively investable, and valuation-sensitive** – a market to own near its benchmark weight, not above it, until the price improves or the earnings prove they can carry the multiple. It is more robust on currency than its reputation, and less investable through the index than the size of its economy implies. This is not a Buy and not a Sell. It is a statement about *how much* and *out of what*, which is the only question an allocator actually has.

Where the benchmark puts you today

India is **11.66% of MSCI Emerging Markets** and **14.84% of MSCI EM ex-China** (both as of 31 July 2026). Both figures are *down* from a brief top-of-the-table peak near 20% of EM in late 2024. India is now the fourth-largest EM weight, behind Taiwan (26.63%), China (21.38%) and Korea (20.33%). A neutral allocator therefore already holds India at roughly one dollar in nine of an EM book. The decision is the deviation, and the benchmark has spent two years telling allocators to hold *less*.

What India wins on / loses on

Wins: a genuine earnings engine (+9.7%/yr EPS contribution over 20 years, among the strongest in EM); the deepest domestic buyer base in EM (DII buying outran FPI selling 4.6× in FY25); a currency that depreciates but does so in the most orderly way in the EM set. **Loses:** the multiple contributed *nothing* over 20 years (−0.1%/yr); the currency handed back −3.8%/yr; MSCI India **trailed MSCI EM over both 5 and 10 years in USD** (8.03% vs 9.19%, 10y); and CAPE at 36 sits ~1.5× its long-run mean.

The expected-return hurdle

Over 20 years India delivered ~7.1% in USD (single-universe Nifty), built almost entirely on earnings net of currency, with **zero help from the multiple**. From a rich CAPE, the multiple is more likely to subtract than add. So the forward case requires India to **sustain double-digit EPS growth at an already-elevated ROE** just to hold the multiple, and something *more* than that to re-rate. GDP growth is nearly a spectator: across 13 markets, growth showed no reliable positive link to returns. The marginal dollar bets on margins and multiple, not on GDP.

Principal risks

- **Valuation compression** from CAPE 36 toward its mean – the single largest identifiable downside.
- **Earnings disappointment** against a fresh all-time-high profit/GDP base that assumes normalisation is permanent.
- **Currency acceleration** beyond the orderly ~3.1%/yr trend, which unhedged turns a respectable return into a poor one.
- **Crowding and issuance:** ~₹3.05 lakh cr of FY26 equity supply is a real sponge on the celebrated domestic bid.

What would change the view (three triggers)

1. **Re-rating trigger:** CAPE falls toward its long-run mean (~24, roughly a third de-rating) while ROE holds – India moves from *hold at weight* toward *overweight*.
2. **Earnings-durability trigger:** two-plus years of double-digit Nifty-50 consolidated EPS growth *with* profit/GDP holding above 5% – confirms the elevated base is structural, not a peak.
3. **Break trigger (to the downside):** profit/GDP reverts below ~4.5% *or* the rupee's annualised depreciation accelerates past ~5%/yr – the bear leg fires, and India moves to *underweight*.

View: hold India near benchmark; earn the overweight, do not assume it. Funded from developed-market or cash sleeves, not from Taiwan or Korea.

EXECUTIVE SUMMARY

Executive Summary

A CIO who reads only this section understands the report. Every number here is drawn from the single canonical table; anything not so drawn is labelled illustrative.

11.66%

INDIA'S WEIGHT IN MSCI EM · 31 JULY 2026

The thesis, in one paragraph

A dollar invested in India captures **part** of India's rise – smaller than the GDP number implies, and concentrated in one leg of the return. The marginal dollar into India is not a bet on GDP growth. It is a bet on **margins and the multiple**: that Indian corporate earnings keep compounding near their historic double-digit pace, that profitability holds at an already-elevated level, and that a rich valuation does not compress. The evidence for this framing is blunt and it is the fact most India research omits: **MSCI India underperformed MSCI Emerging Markets over both the last five years and the last ten, measured in US dollars** (8.03% vs 9.19% annualised over ten years, net). The celebrated "India decade" paid a dollar investor *below* the emerging-market average it was part of. That is not a rounding artifact or one bad window. It is the durable centre of the allocator's case, and everything in this report is read against it.


The benchmark round-trip

Start where allocators actually meet India: as a weight in an index, not as a country.

EXHIBIT 1

India round-tripped its benchmark weight – and is now back below Taiwan and China

India's weight in MSCI Emerging Markets: 9% (2020) → ~20% (2024 peak) → 11.7% (Jul-2026). Taiwan is now the #1 EM weight.



Source: MSCI EM factsheets (GREEN, 31-Jul-2026); weight history AMBER (media citing MSCI). Kakde Research.


In 2020 India was about 9.2% of MSCI EM. By late 2024 it had climbed to roughly 20%, briefly the single largest country in the MSCI EM Investable Market Index. As of 31 July 2026 it sits at **11.66%** – fourth, behind Taiwan (26.63%), China (21.38%) and Korea (20.33%). Strip China out, the construction many allocators now run, and India is **14.84%** of EM ex-China, third behind Taiwan (33.9%) and Korea (25.9%), not the 30%-plus a decade of "India is the EM story" commentary would lead you to expect.

The round-trip is not a footnote to the thesis. It *is* the thesis, seen from the only vantage point funds occupy. India re-rated, over-earned its weight on a wave of new issuance and rising profits, and then de-rated and underperformed while Taiwan and Korea, levered to the AI and semiconductor capex cycle, delivered the returns India's price had already embedded. The benchmark is a relative-value machine refreshed daily. India's falling weight is not a verdict on India's economy; it is a verdict on India's *price relative to its alternatives*, and for two years that verdict has been "hold less."

The bull reads the collapse as a buying signal: the froth is out, now is the entry point. Before accepting that, look at what the weight was built on and what the returns actually were.

The Return Bridge

Any equity return in a foreign currency reduces to four moving parts: earnings growth, the change in the valuation multiple, dividends, and the currency. Build that bridge on a single consistent universe (Nifty earnings to Nifty price, so nothing is distorted by splicing indices) over twenty years, and the answer is clean.



Twenty years, single universe, US-dollar investor: **+9.7 (EPS) – 0.1 (P/E) + 1.3 (dividend) – 3.8 (FX) ≈ 7.1% a year in USD**. Read the bars in order and the mechanism is unambiguous:

- **Earnings compounded at +9.7%/yr** – a genuinely powerful engine, among the strongest in the EM set.
- **Dividends added +1.3.**

- **The multiple contributed essentially nothing, -0.1.** Over two decades the market paid roughly the same multiple at the end as at the start; all the intermediate re-rating and de-rating netted to zero. A multiple that expands and then contracts is a round-trip too.
- **The rupee subtracted -3.8.** The currency handed roughly 40% of the earnings compounding back to the foreign investor before it reached their account.

The residual between this decomposition and the published Nifty total return is about 1.8 percentage points, explained by anchor and endpoint mismatch (calendar-year versus August-2026 alignment), not by a missing economic factor. There is no mystery leg. **The return is earnings, minus currency, with the multiple washing out.** Two integrity flags travel with this chart: NSE switched Nifty earnings from standalone to consolidated in April 2021, so any EPS or P/E leg spanning that break measures an accounting change as performance; and the five-year window is start-point sensitive (December 2021 was a local valuation peak), which is why the ~3.8% five-year single-universe USD return is the least robust number in the set and is never headlined. The 20-year (~7.1%) and 10-year (~9.5%) figures carry the argument.

The decisive implication: because the multiple added nothing over the full cycle, the marginal dollar into India today is **not** a bet on GDP and **not** a passive bet on re-rating. It is a bet that margins hold and the multiple does not compress from a rich level. GDP growth is nearly a spectator to that bet.

Growth is not the reason, and the data says so

Because the thesis rests on displacing GDP as the reason to own India, it is worth killing the GDP argument cleanly. Across thirteen markets, ten-year real GDP growth shows **no reliable positive relationship** with ten-year equity return in USD: the fitted slope is negative (-0.49), R^2 is 0.03, correlation -0.16. Growth does not buy returns. The named points enforce the discipline. Vietnam grew fastest of the thirteen at 6.2% real and returned poorly; Taiwan grew at just 3.8% and returned the most, ~22.5%/yr in dollars; India grew at 5.8% and returned 8.0%, sitting almost exactly on a gently downward-sloping line. India is the fastest-growing *large* EM after Vietnam – a precise claim that survives the data, unlike the looser "India is the fastest-growing market," which does not. The reason the ordering scrambles is the transmission chain: growth gets diluted by issuance, taxed by currency, and pre-paid in the multiple, so a market everyone knows will grow trades at a price that already embeds the growth. India spent most of the last decade as exactly that market. The scatter is not an argument that India will do badly; it is an argument that "India grows fast" is not, by itself, a reason for the marginal dollar.


Four Indias, and the foreign-investable ladder

The economy and the market are not the same object, and the most common error in India research is to treat the chain from GDP to foreign ownership as a tidy waterfall in which each layer is smaller than the last. It is not. These are **adjacent analytical objects**, and one of them breaks the descent.

EXHIBIT 3

The Four Indias are different objects — flows and stocks, not one shrinking waterfall

Adjacent magnitudes, ~2025. A \$3.9tn economy; a \$5.35tn listed market; ~\$0.88tn actually held by foreigners.



- **The economy:** ~\$3.9 trillion of nominal GDP (₹330.7 lakh crore, FY25). Note the figure: *not* \$4.1 trillion, which comes from a different base and period; mixing the two corrupts every downstream ratio.
- **The listed market:** ~\$5.35 trillion of market capitalisation — *larger than the economy*, which is exactly why the waterfall framing fails. Market cap is a claim on future profits and global revenue, not a slice of this year's GDP.
- **The free float:** ~\$2.67 trillion, because promoters hold 50.0% of the listed market and that half is not for sale.
- **Foreign-held:** ~\$0.84–0.93 trillion of actual foreign custody.

Source: MoSPI, NSE, NSDL, Motilal Oswal [GREEN; float derived AMBER]. Foreign-held ≠ foreign-ownable — see investability ladder. Kakde Research.

The load-bearing distinction is **held ≠ ownable**. Foreign custody is what foreigners *currently* hold, not the maximum they *could* own. The true ownable maximum runs through a longer ladder — total cap, then free float, then the foreign-investable cap (which excludes strategic and promoter blocks), then actual holdings, and only then the *remaining* foreign room. Using custody as the ownable ceiling, as much India coverage does, overstates how much of India a global investor can actually buy. Ownership today: promoter 50.0%, FPI 17.3%, DII 16.1%, retail 18.5%. The investable India a foreign dollar reaches is narrower than the listed India, which is itself larger than the economy — the layers do not nest.

FY26 profit: a fresh all-time high, read two ways

Motilal Oswal (11 June 2026): Nifty-500 profit/GDP reached **5.2% in FY26, an all-time high** (broader listed India Inc. 5.7%, an 18-year high), with corporate profit up **+15.6% YoY against nominal GDP of +8.9%**. Profit is again outpacing the economy, and the ratio has made a *new* high.

This cuts two ways and the report holds both. **Bull reading:** the profit share is now structurally higher than the 2003–19 range; the post-2020 surge (tax cut, falling credit costs, operating leverage) was not a one-off that is mean-reverting – it is a level shift, and FY26 proves it by making a new high.

Bear reading: a fresh all-time high is a demanding base from which to underwrite further growth, and it sits underneath an elevated CAPE. The honest synthesis: the *level* is probably real and durable, but the *forward asymmetry* from a record profit share and a rich multiple is unattractive. (This figure traces to a single source and is flagged as such at every use.)

Valuation: rich on CAPE, cheap-ish on P/B, and the premium is not what it looks like

India's valuation sends two signals and the report shows both rather than cherry-picking. **CAPE says rich:** IIMA India CAPE (Nifty-500) is **36.0** as of July 2026, roughly **1.5× its long-run mean of ~24**, near 2021 highs – a mean-reversion headwind. **P/B says cheap-ish:** Nifty P/B of **2.95** is **~18% below its 10-year average (3.62)**, and trailing P/E is ~12% below its 10-year average – largely because the 2021 standalone-to-consolidated earnings reset lowered headline P/E. India trades at P/E 23.9, P/B 3.32, dividend yield 1.27%.

The subtle and important point concerns the *premium* India commands over EM. It is often attributed to superior return on equity. The direct ROE comparison is a **data gap** in this report, not a settled parity; the derived ROE (~14.4%) sits close to EM's ~13.8%, but a 5-basis-point superiority claim built from a P/B-to-P/E identity cannot bear weight. The honest characterisation: India's premium is a **P/B-persistence and growth-durability premium**, not a demonstrated ROE-superiority premium. It rests on the market's willingness to keep paying up for a track record of compounding, which is precisely the thing a mean-reverting CAPE puts at risk.


Currency: a low-volatility depreciator you can size around

The rupee is the second-largest force in the return bridge, and allocators routinely mishandle it.

EXHIBIT 4

The rupee depreciates steadily – but is one of the least volatile EM currencies

EM currency depreciation vs volatility, ~2016–2026. INR ~3.1%/yr at 8.1% vol vs Brazil 19.1%, South Africa 15.7%.



Source: BIS/market data via Kakde Research & Fable dataset [AMBER; RBI-corroborated]. Hedge cost ~2.8%/yr. Kakde Research.

The INR has depreciated about **3.1%/yr** against the dollar over the last decade – a persistent drag, and the reason a strong local-currency earnings machine converted into a below-EM dollar return. But it is a remarkably *orderly* depreciator. Its annualised volatility is **8.1%**, the lowest of any depreciating currency in the EM set, against Brazil at 19.1%, South Africa at 15.7%, Mexico at 11.6% and Korea at 11.4%. Only actively managed currencies sit lower. The rupee bleeds; it does not detonate. That predictability makes it hedgeable, and the hedge is cheap relative to the drag: the 12-month forward premium is about **2.8%/yr** (August 2026). An investor who hedges pays ~2.8% to remove a ~3.1% expected drag – close to a wash on the central estimate, converting an uncertain loss into a known, smaller cost. Whether to hedge is a live decision that moves realised returns by more than most stock-selection choices inside the India sleeve. This replaces the dismissive "bottom-tier rupee" framing: the currency is a manageable, quantified headwind, not a disqualifier.

Risk, so the return is priced against something

MSCI India's realised volatility is about **17.3% annualised** over three years (17.0% over five, 19.1% over ten) – not the ~14% softer treatments imply. Three-year Sharpe is 0.12. The worst drawdown in the modern record was **-72.6%** (global financial crisis). So the 8.03% ten-year USD return was earned at ~17% volatility, below the EM average, with tail risk that has historically taken more than 70% off the top. That is the risk-adjusted object on the table. Not a bad object – but not the one the headline growth rate advertises.

The domestic bid: floor and ceiling from the same pool

One structural feature genuinely differentiates India from most of EM, and it deserves front-matter billing: the domestic buyer base is deep enough to absorb foreign selling. In FY25, domestic institutions bought +**₹6.07 trillion** of equities while foreign portfolio investors *sold* -**₹1.31 trillion** – domestic buying outran foreign selling by **4.6×**. The monthly SIP run-rate is ~₹31,781 crore, but that is only a slice of the machine. Behind it sit vast pools: equity mutual-fund AUM of ~₹37.3 lakh crore, LIC's ~₹57 lakh crore of assets, EPFO's ₹24.8 lakh crore, and NPS at ~₹16 lakh crore. This is why India's market held up through a foreign-selling episode that would have broken a thinner market.

The bull stops there. The report does not, because the *same* savings pool that provides the floor also feeds the ceiling. Equity supply – IPOs, QIPs and offers-for-sale – ran **₹3.05 lakh crore in FY26** and a record **₹3.71 lakh crore in FY25**. New issuance is a real sponge that soaks up domestic inflows and dilutes the per-share earnings claim, which is one reason the earnings *engine* (+9.7%/yr aggregate) does not translate one-for-one into per-share returns. The domestic bid is a proven shock-absorber; it is not yet a proven permanent floor, and it comes with an issuance offset that the celebratory version omits. Floor and ceiling are drawn from one pool.

One dollar, ten markets

Put India in its opportunity set. A global allocator's dollar does not compete against India's own past; it competes against everything else it could buy today.

EXHIBIT 5

One Dollar, Ten Markets – no single winner; India leads on growth, lags on price and return

Allocator scorecard, 1 (weak) to 5 (strong), India-relative. Read across a row for a market; down a column for who wins each dimension.

	10y USD return	Real GDP growth	Valuation (P/E)	Div+buyback yield	FX stability	Foreign access	Governance
India	3	5	2	2	4	3	3
China	2	5	4	4	3	2	2
Taiwan	5	4	2	3	5	4	4
S. Korea	5	3	4	3	3	4	3
Indonesia	1	4	4	5	3	3	3
Brazil	4	2	5	5	1	4	3
Mexico	4	1	4	4	2	4	3
S. Africa	4	1	4	4	2	3	3
Saudi	4	3	4	4	5	2	2
Vietnam	3	5	3	3	5	2	2

SCALE 1 2 3 4 5 1 weak · 5 strong

Source: Return/growth/valuation/FX from MSCI, World Bank, market data [GREEN]; access/governance = Kakde Research judgment, labelled [INTERP]. Kakde Research.

The matrix scores the EM set across the dimensions that actually drive a foreign dollar's outcome – earnings engine, valuation, currency behaviour, investability, realised return. India's row is distinctive but not dominant: a top-tier earnings engine, a rich valuation, a benign-but-persistent currency, middling investability, and a *below-median realised dollar return*. Taiwan delivered the returns. Korea is cheaper. China carries a governance discount that the ex-China construction is built to avoid. The honest read of the matrix is that **India is a quality earnings story trading at a full price, and the dollar that funds an India overweight has to come out of a specific alternative** – a choice the report treats as a first-order decision, not a residual.

The transmission chain, named

The return bridge is the compressed version of a longer chain, and naming it explains where the GDP number leaks before it becomes a foreign return: **GDP growth → corporate revenue → listed-company revenue → margins and earnings → per-share earnings after dilution → the multiple the market pays → dividends and buybacks → the share a foreign investor can own → the currency it converts back through**. India leaks at several links. Not all of the economy is corporate; not all of the corporate sector is listed; not all listed earnings survive dilution, because a fast-growing market issues heavily; not all of the listed market is foreign-ownable; and whatever reaches a foreign holder converts through a rupee that depreciates ~3.1%/yr. The bridge already isolated the two decisive links –

earnings (large and positive) and currency (a persistent drag) – and revealed that the multiple contributed nothing across the cycle. The chain is why "a dollar captures part of India's rise": each link is a place the story thins before it reaches the account.

Portfolio implications, in brief

India is already ~11.66% of an EM book and ~14.84% of EM ex-China; a neutral allocator holds it there. The overweight case cannot rest on GDP (growth does not buy returns), nor on re-rating (the multiple added nothing over 20 years). It rests on a specific, testable claim about earnings and margins. The rupee should be sized explicitly, and hedging is close to a wash. The funding source matters as much as the sizing: an India overweight funded from Taiwan or Korea sells the very things that beat it, while one funded from the developed-market or cash sleeve is a cleaner expression of the bet. Sizing should respect ~17% volatility and a below-EM expected return. The mandate determines the size: a benchmark-aware EM fund holds at or modestly below 11.66%; an EM ex-China fund carries materially more India by construction; a global fund reaches India through its EM sleeve; and only an unconstrained, long-horizon pool can express India the efficient way, through selected single stocks rather than the 52.9%-concentrated index. In every case the position is a sizing-and-funding decision, not a rating.

Bull and bear, as falsifiable claims

The report commits to one canonical falsification dashboard rather than a house view dressed as certainty. **The bull is falsified** if profit/GDP reverts below ~4.5%, if the rupee's depreciation accelerates past ~5%/yr, or if two-plus years of double-digit EPS fail to materialise from the elevated base. **The bear is falsified** if CAPE de-rates toward its ~24 mean while ROE holds, if EPS compounds double-digit for two-plus years with profit/GDP above 5%, or if the currency drag durably narrows. Both legs carry explicit thresholds, reconciled into a single dashboard in Section 12. The Kakde Research position sits between them by design: **structurally attractive, selectively investable, valuation-sensitive, more robust on currency than assumed, and less investable through the index than the economy implies.**

SECTION 1

India, Seen From the Benchmark

A global allocator does not meet India as a country. They meet it as a line in an index, a weight that has to be over- or under-held relative to a benchmark someone else defined.

8.03% vs 9.19%

MSCI INDIA VS MSCI EM · 10-YEAR USD NET


A global allocator does not meet India as a country. They meet it as a line in an index, a weight that has to be over- or under-held relative to a benchmark someone else defined. Start there, because the weight has done something over the last six years that most India research never mentions: it round-tripped.

In 2020, India was about 9% of MSCI Emerging Markets. By late 2024 it had climbed to roughly 20%, and for a stretch it was briefly the single largest country in the MSCI EM Investable Market Index, at the top of the emerging-market table for the first time. As of 31 July 2026 it sits at 11.66%. Not first. Fourth. Taiwan is now the largest EM weight at 26.6%, China 21.4%, Korea 20.3%, and India behind all three. Strip China out of the index (the construction many allocators now run because of China's own governance discount) and India is 14.8%, again third, behind Taiwan at 33.9% and Korea at 25.9%. Not the 30%-plus that a decade of "India is the EM story" commentary might lead you to assume.

EXHIBIT 6

India round-tripped its benchmark weight — and is now back below Taiwan and China

India's weight in MSCI Emerging Markets: 9% (2020) → ~20% (2024 peak) → 11.7% (Jul-2026). Taiwan is now the #1 EM weight.



Source: MSCI EM factsheets (GREEN, 31-Jul-2026); weight history AMBER (media citing MSCI). Kakde Research.

The chart is the whole opening argument in one frame. India's orange line rises steadily from 2015, accelerates hard into a late-2024 peak, then falls just as hard while Taiwan's line, riding the AI semiconductor cycle, crosses above it and keeps going. The round-trip is not a footnote to the India thesis. It is the thesis, seen from the only vantage point that funds actually occupy. India re-rated, grew into and then past its "fair" benchmark weight on a wave of new issuance and rising profits, and then de-rated and underperformed while the money that had crowded in found faster-moving stories elsewhere.

That single picture reframes every question that follows. The interesting question about India in 2026 is not whether it is a large, fast-growing economy. It plainly is. India is the world's fifth-largest equity market by capitalisation, around \$5.05 trillion, sitting behind only the United States, China, Japan and (depending on the day and the FX

print) a shifting cast at the top. An economy of that scale with a domestic growth rate near the front of the large-EM pack is, definitionally, interesting. The question an allocator has to answer is narrower and harder: **if a global investor has one more dollar to put to work, why should it go to India – and, just as important, what does it come out of?**

That second half is the part India bulls skip. A marginal dollar into India is a marginal dollar out of something. Out of Taiwan, which has delivered the returns. Out of Korea, which is cheaper. Out of the US, which is the funding source for most EM overweights in practice. Out of cash, which now yields something again. India does not compete for capital against its own past. It competes against whatever else that dollar could buy today, and the honest starting point is that the benchmark has been quietly telling allocators to hold *less* India, not more, for two years running.

The weight is not the return

It is tempting to read the weight collapse as a buying signal – India got cheap relative to itself, the froth came out, now is the entry point. Before accepting that, look at what the weight was built on and what the returns actually were.

Here is the fact most India research buries, and this report promotes to the front: **India underperformed the MSCI Emerging Markets index over both the last five years and the last ten, measured in US dollars.** MSCI India returned 8.03% annualised over ten years in USD net terms. MSCI EM returned 9.19%. India did not beat the benchmark it is supposedly the star of. It trailed it, over five years and over ten. For a dollar investor, the celebrated "India decade" delivered a rate of return below the emerging-market average it was part of.


This is not a rounding artifact or a single bad window. It is the durable, awkward centre of the allocator's case. Everything else in this report (the growth machine, the profit cycle, the currency, the ownership funnel) has to be read against the plain fact that the transmission from India's real economy to a foreign investor's dollar returns has, so far, leaked away most of what the domestic story promised.

To see where it leaked, decompose the return. Any equity return in a foreign currency reduces to four moving parts: earnings growth, the change in the valuation multiple the market pays for those earnings, dividends, and the currency. Build that bridge on a single consistent universe (Nifty earnings to Nifty price, so nothing is distorted by splicing different indices together) over twenty years, and the answer comes out clean.

EXHIBIT 7

20 years of Indian equity: earnings built it, the rupee took ~40% of it back

India Return Bridge, single universe (Nifty→Nifty), USD investor, 20 years to 2026. Multiple re-rating added nothing.



Source: NSE (Nifty level/EPS/PE), RBI (USD/INR); Kakde Research decomposition [GREEN inputs, derived legs]. Residual ~1.8pp = timing, not a missing factor.

Twenty years, single universe, US-dollar investor. Earnings growth contributed +9.7 percentage points a year, a genuinely powerful engine, among the strongest in the emerging-market set. Dividends added +1.3. The valuation multiple contributed essentially nothing, -0.1: over two decades the market paid roughly the same multiple at the end as at the start, so all that intermediate re-rating and de-rating netted to zero. And the rupee subtracted -3.8 points a year. The sum is about 7.1% a year in dollars.

Read those bars in order and the mechanism is unambiguous. India's companies compounded earnings at a rate most markets would envy. Then the currency handed roughly 40% of that compounding back to the foreign investor before it ever reached their account. Multiple expansion, the thing that made the 2021–24 period feel so good, added nothing across the full cycle, because a multiple that expands and then contracts is a round-trip too. The residual between this decomposition and the published Nifty total return is about 1.8 percentage points, and it is explained by anchor and endpoint mismatch, the calendar-year versus August-2026 alignment, not by a missing economic factor. There is no mystery leg. The return is earnings, minus currency, with the multiple washing out.

Two data-integrity flags belong on this chart wherever it is used, because they are where naive versions of this analysis go wrong. First, NSE changed the Nifty earnings basis from standalone to consolidated in April 2021; any EPS or P/E leg that spans that break is measuring an accounting change as if it were performance. Second, the five-year window is start-point sensitive (December 2021 was a local valuation peak with Nifty P/E near 23.7), which is why the five-year single-universe USD return of about 3.8% is the least robust number in the set and should never be headlined. The twenty-year and ten-year figures (roughly 7.1% and 9.5% in USD) carry the argument. The five-year is a cautionary tale about anchoring.

The transmission chain

The return bridge is the compressed version of a longer chain, and naming the chain now sets up the rest of the report. A dollar of GDP growth in India has to survive a sequence of steps before it becomes a dollar of return in a foreign portfolio:

GDP growth → corporate revenue growth → listed-company revenue → margins and earnings → per-share earnings after dilution → the multiple the market pays → dividends and buybacks → the share a foreign investor can actually own → the currency it converts back through.

Every link in that chain is a place where the story can leak, and India leaks at several of them. Not all of the economy is corporate. Not all of the corporate sector is listed. Not all of listed earnings survive dilution, because a fast-growing market issues a lot of new equity, and issuance is a claim on future per-share earnings. Not all of the listed market is available to foreigners. And whatever does reach a foreign holder converts back through a rupee that depreciates about 3.1% a year on trend.

The return bridge already told us the two decisive links. Earnings growth is real and large. The currency is a persistent, quiet drag. What the bridge also told us – and this is the part that should reorder an allocator's priors – is that **the multiple contributed nothing over the full cycle**. Which means the marginal dollar into India today is not a bet on GDP, and it is not a passive bet on re-rating adding to returns from here. It is a bet on two things specifically: that margins hold or expand from an already-elevated level, and that the multiple, currently rich on some measures, does not compress. GDP growth is almost a spectator to that bet. It sets the ceiling on earnings but does not determine what a foreign investor keeps.

That is the preview of the answer this report builds toward. A dollar investor can capture *part* of India's rise – the earnings part, net of currency – but not all of it, and the part they capture has almost nothing to do with the headline GDP number that dominates the India narrative. The marginal dollar bets on margins and multiple. Everything downstream in this report is an attempt to price that bet honestly.


Growth is not the reason – and the data says so plainly

If the case for the marginal dollar is not GDP, it is worth killing the GDP argument cleanly rather than leaving it to linger, because it is the argument allocators hear most often and trust least once they have seen the evidence.

EXHIBIT 8

Across 13 markets, faster growth did not buy higher returns

Equity return (MSCI net TR, USD) vs real GDP growth, 2015–2026. Vietnam grew fastest (6.2%); Taiwan returned most.



Source: MSCI factsheets 31-Jul-2026; World Bank WDI [GREEN]. Slope/R² Kakde Research. Kakde Research.

Thirteen markets, ten-year real GDP growth on the horizontal axis, ten-year equity return in USD on the vertical. If faster-growing economies delivered higher equity returns, the cloud of points would slope up to the right. It does not. The fitted line slopes gently *down*: slope -0.49, R² 0.03, correlation -0.16, across thirteen markets. That is not a relationship. It is the visual signature of no reliable relationship at all – the growth an economy posts tells you almost nothing about what its stock market paid a dollar investor.

The individual points make it concrete and they discipline a specific error. Vietnam grew fastest of the thirteen, at 6.2% real a year, and sits near the bottom of the return distribution. Taiwan returned the most, around 22.5% a year in USD, on real growth of just 3.8%. India is the orange dot: 5.8% growth, 8.0% return – fast-growing, middling-returning, sitting almost exactly on a downward-sloping line. Note precisely what this says about India. India is the fastest-growing *large* emerging market after Vietnam. It is not the fastest grower outright – Vietnam's 6.2% exceeds India's 5.8%. That distinction matters, because the entire point of this chart is that growth *rank* and return *rank* are different orderings, and getting India's growth rank wrong in the bull direction is exactly the slip that lets people believe growth buys returns.

Why does the link break? Because everything between GDP and the shareholder is a wedge. High-growth economies tend to be capital-hungry, and they fund that growth by issuing equity, which dilutes the per-share claim. They often run persistent currency depreciation, which taxes the foreign holder. And crucially, markets price growth *in advance* – a country everyone knows is growing fast trades at a multiple that already embeds the growth, so realised returns

depend on whether reality beats an already-optimistic price, not on whether growth is high. India has been, for most of the last decade, a market where the good news was in the price. The scatter is not an argument that India will do badly. It is an argument that "India grows fast" is not, by itself, a reason for the marginal dollar.

Section 2 opens on exactly that scatter, because the growth machine has to be understood before it can be discounted.

What the weight round-trip actually taught allocators

Return to the benchmark, now armed with the bridge and the scatter, and the 2020→2024→2026 round-trip reads as a case study in the whole chain firing and then reversing.

On the way up, India did over-earn its weight, and it did so honestly for a while. Corporate profits rose faster than nominal GDP. New listings and follow-on issuance expanded the investable market – MSCI weight is partly a function of how much investable float exists, so a wave of IPOs and QIPs mechanically lifts a country's weight even before any price moves. Foreign flows chased the re-rating. India rose from roughly 9% to roughly 20% of EM on a combination of real earnings, real issuance, and a rising multiple. Three of the four bridge legs were pushing the same direction at once.

On the way down, two of those reversed. The multiple stopped expanding and began to compress from a rich level. And the relative story changed: Taiwan and Korea, levered to the AI and semiconductor capital-expenditure cycle, delivered the returns India's valuation had already priced in but not delivered. Capital is relative. An allocator did not need India to fall to cut it; they only needed something else to rise faster, and something did. The weight fell to 11.66% not because India's economy faltered but because the benchmark is a relative-value machine and India lost the relative race for two years.

That is the uncomfortable lesson embedded in the weight chart. India's benchmark weight is not a verdict on India's economy. It is a verdict on India's *price relative to its alternatives*, refreshed daily, and for the recent past that verdict has been "hold less." The bull who points at the round-trip as a buying opportunity has to explain why the thing that de-rated India, a rich multiple meeting faster stories elsewhere, has now reversed. Nothing in the benchmark data says it has.

The currency, as its own decision

The rupee earns its own paragraph, because the return bridge made it the second-largest force in the whole decomposition and allocators routinely mishandle it. The INR has depreciated about 3.1% a year against the dollar over the last decade. That is the -3.8 leg in the bridge (the two differ because the bridge spans twenty years and captures a slightly steeper early stretch), and it is the single reason a strong local-currency earnings machine converted into a below-EM dollar return.

The instinct is to treat that as a disqualifier. It is not, and the reason is volatility. The rupee is a persistent depreciator, but a remarkably *orderly* one. Its annualised currency volatility is about 8.1%, the lowest of any depreciating currency in the emerging-market set, against Brazil at 19.1%, South Africa at 15.7%, Mexico at 11.6% and Korea at 11.4%. Only the actively managed currencies sit lower. So the rupee costs a dollar investor a predictable toll rather than delivering the violent, portfolio-wrecking gaps that make some EM currencies uninvestable. It bleeds; it does not detonate. For an allocator that distinction is the difference between a currency you can size around and one you cannot.

That predictability also means it can be hedged, and the hedge is cheap relative to the drag: the twelve-month forward premium is around 2.8% a year as of August 2026. An investor who hedges pays roughly 2.8% to remove a roughly 3.1% expected drag – close to a wash on the central estimate, converting an uncertain currency loss into a

known, smaller cost. Whether to hedge is a live decision, not a footnote, and it moves the realised return by more than most stock-selection choices inside the India sleeve. The point for the marginal dollar is that the currency is a manageable, quantified headwind, not a reason to stay away, but it is one that has to be sized explicitly, because leaving it unhedged is what turned the celebrated earnings engine into a below-benchmark result.

Risk, so the return is priced against something

One number keeps the return figures honest. MSCI India's realised volatility is about 17.3% annualised over three years (17.0% over five, 19.1% over ten) – not the roughly 14% that softer treatments imply. Three-year Sharpe is 0.12. The worst drawdown in the modern record was -72.6%, in the global financial crisis. So the 8.0% ten-year USD return was earned at a volatility near 17% and with tail risk that has, historically, taken more than 70% off the top. An allocator sizing India is buying a mid-single-digit-to-high-single-digit dollar return, below the EM average, at emerging-market volatility, in a currency that bleeds about 3% a year. That is the risk-adjusted object on the table. Not a bad object. But not the one the headline growth rate advertises.

What this means for the allocation

The benchmark is the cleanest possible starting statement of the problem, so let it state it. India is 11.66% of MSCI EM and 14.8% ex-China, down from a brief top-of-the-table peak, and it trailed the EM index over both five and ten years in dollars. A neutral allocator holds India at its benchmark weight and needs a reason to deviate. The bull case for an overweight cannot rest on GDP growth, because the cross-country evidence says growth does not buy returns and India's own return bridge says the multiple added nothing over twenty years. The overweight is a specific, narrow bet: that Indian corporate *earnings* keep compounding near their historic double-digit pace, that margins hold at an already-high level, that the multiple does not compress from a rich starting point, and that the rupee's roughly 3% annual drag does not accelerate. Each of those is testable, and the rest of this report tests them. The marginal dollar into India is a bet on margins and multiple, funded most plausibly out of the developed-market or cash sleeve rather than out of Taiwan or Korea, and sized to a return that history says will be respectable but below the EM average unless something in the transmission chain changes. That is the frame. The next question is whether the growth everyone points to even reaches the listed, investable economy the bet actually owns.

SECTION 2

The Growth Machine

The scatter that closed Section 1 is the right place to open Section 2, because it defines the task. Across thirteen markets, faster growth did not buy higher returns: slope -0.49 , R^2 0.03 , thirteen observations, Vietnam growing fastest and Taiwan returning most.

$$n = 13 \cdot R^2 = 0.03$$


GDP GROWTH VS EQUITY RETURN, CROSS-SECTION

The scatter that closed Section 1 is the right place to open Section 2, because it defines the task. Across thirteen markets, faster growth did not buy higher returns: slope -0.49 , R^2 0.03, thirteen observations, Vietnam growing fastest and Taiwan returning most. India sits in the middle of that cloud: real growth of 5.8% a year, a dollar return of 8.0%, roughly on the downward-sloping line.

EXHIBIT 9

Across 13 markets, faster growth did not buy higher returns

Equity return (MSCI net TR, USD) vs real GDP growth, 2015–2026. Vietnam grew fastest (6.2%); Taiwan returned most.



Source: MSCI factsheets 31-Jul-2026; World Bank WDI [GREEN]. Slope/ R^2 Kakde Research. Kakde Research.

The natural response is to dismiss the growth story as irrelevant to the allocation. That is too quick. Growth is not the reason for the marginal dollar, but it is the substrate the reason lives in: it sets the ceiling on how fast earnings *can* compound, shapes which sectors list, and determines whether the demographic dividend everyone cites turns into consumers with income or into a labour surplus with none. The job of this section is to characterise the growth machine precisely enough to answer one downstream question: does this growth reach the listed, investable economy that a foreign portfolio actually owns? Understanding the engine is the prerequisite for discounting it correctly.

The shape of the economy

India's nominal GDP was ₹330.7 lakh crore in FY25, which at the report's own consistent exchange-rate band works out to about \$3.9 trillion. (Note the figure carefully: this is not \$4.1 trillion. The \$4.1 trillion number that circulates comes from a different GDP base and calendar period, and mixing the two inside a single funnel – a common practice – corrupts every ratio built on top of it. One number, one FX, throughout.) That \$3.9 trillion makes India the world's fifth-largest economy, and its composition is unusual for a country at its income level.

Services dominate. India skipped much of the manufacturing-led development path that China, Korea and Taiwan followed, and built instead on services: IT and business-process exports at the high end, and a vast, mostly informal domestic service economy at the low end. Manufacturing's share of output has been stuck well below the government's stated ambitions for over a decade, despite successive industrial-policy pushes. Agriculture still occupies a share of employment far larger than its share of output, which is another way of saying a large fraction of the workforce is trapped in the lowest-productivity part of the economy. This matters for the allocation because the listed market looks nothing like the employment base. The Nifty is financials, IT, energy, consumer and industrials – the formal, corporate, capital-market-facing slice. The workforce is agriculture, informal services and small-scale trade. The gap between those two Indias is the first place growth can fail to reach the investable economy.

Where the growth comes from: accumulation versus productivity

Growth accounting splits any economy's expansion into three sources: more labour, more capital, and total factor productivity, the residual efficiency with which labour and capital are combined. The distinction is not academic. Growth driven by piling up factors of production runs into diminishing returns and requires ever-rising investment to sustain. Growth driven by rising productivity compounds cleanly and is what ultimately shows up in corporate margins and returns on capital.

India's growth has leaned heavily on factor accumulation, a rising working-age population and periods of strong investment, with productivity contributing a more variable and, in several stretches, disappointing share. This is a qualitative characterisation, not a precise decomposition; the point estimates on India's TFP contribution vary widely by source and methodology and should be treated as illustrative rather than canonical. But the direction is well established and it carries an allocator implication: an economy growing mainly by adding workers and capital needs those workers to find productive work and that capital to earn a return. If either fails, headline GDP can keep rising while the per-capita, per-share, per-unit-of-capital numbers that actually drive equity returns stagnate. Growth that is fast in aggregate but thin in productivity is exactly the kind that shows up big in the GDP print and small in the return bridge.

The capital side of the accounting deserves the same scepticism as the labour side. India funds its investment from a domestic savings pool that is large in absolute terms but has been shifting in composition, with household financial savings under pressure and a rising share of savings intermediated through capital markets rather than banks. An economy that grows by accumulating capital depends on that capital earning a return; if investment runs ahead of the productive opportunities to deploy it, the marginal rupee of capital earns less, aggregate growth is sustained only by investing more, and returns on capital – the thing that ultimately underwrites equity returns – drift down. This is the capital-side version of the jobs problem: quantity of factors is not the constraint, it is the efficiency with which they convert into output and profit. The productivity-light character of India's growth is precisely a statement that this conversion has been less efficient than the investment and population numbers, taken alone, would promise.

Demographics are a conditional bet, not a guarantee

The single most repeated line in India bull cases is demographic: a young and growing population, a rising working-age share, a "dividend" that will power decades of growth while China and the developed world age. The demographics are real. The dividend is conditional, and the condition is jobs.

A demographic dividend is not delivered by births. It is delivered when a rising share of working-age people is actually employed in progressively more productive work, so that output per capita rises and a consuming middle class forms. If the working-age population grows faster than the economy creates formal, productive jobs, the same demographics that are labelled a dividend become a liability: underemployment, informality, and political pressure

rather than compounding consumer demand. India adds millions to its working-age population every year. The binding question is whether the formal economy absorbs them, and on the evidence it has not been absorbing them fast enough.

This is where the growth story and the employment story pull apart, and the allocator has to hold both. GDP can grow at 6% while the quality of that growth – measured in formal-sector job creation – disappoints. A demographic tailwind conditional on jobs that are not appearing at scale is a weaker input to the earnings machine than the raw population pyramid suggests. The dividend is a call option on job creation, not a bond that pays regardless.

Employment, informality and human capital

The character of Indian employment is the growth machine's most important and least flattering feature. The overwhelming majority of Indian workers are in the informal economy (small enterprises, self-employment, casual labour, agriculture), outside the formal wage system, largely outside the tax net, and almost entirely outside the listed corporate sector. Widely cited estimates place informal employment around 90% of the total workforce; treat that as an illustrative order of magnitude rather than a precise canonical figure, because definitions and sources differ, but the qualitative fact is not in dispute. Informality is the norm, not the exception.

Two consequences follow for the allocation. First, the transmission from GDP growth to *listed* corporate revenue is weaker than in a more formalised economy, because so much economic activity happens in businesses that will never list and whose growth never touches a public earnings line. Second, formalisation itself becomes a distinct investment theme – the slow migration of activity from informal to formal, which routes revenue and profit toward the organised, often listed, players. That formalisation trade is real and has driven parts of the recent earnings story, but it is a share-shift story (organised taking share from unorganised) layered on top of the aggregate growth story, and the two should not be double-counted.

Human capital sits underneath all of this. India's headline education and skilling metrics have improved, but quality remains uneven and the gap between the credentials the labour force holds and the skills the formal economy demands is wide. A young population is an asset only to the degree it is a *productive* young population, and human-capital constraints are one reason the demographic dividend has been slower to convert than the population math implies. These are contextual, well-documented features rather than canonical point estimates, and they belong in the model as reasons the growth-to-earnings transmission is lossy, not as precise inputs.

India is not one economy

The most important correction the growth machine needs is that there is no single Indian economy to model. There are several, layered on top of each other, and the dispersion across them is enormous. The clearest way to see it is at the state level, where the income gap is wide enough to change what "India grows at 6%" even means.

METRIC	VALUE	NOTE
Richest large state per-capita income (Karnataka)	₹4.22 lakh	GREEN
Bihar per-capita income	₹0.76 lakh	GREEN
Ratio, richest large state to Bihar	~5.6×	GREEN
Share of national GDP, top-5 states	~48%	GREEN
Share of FDI, top-2 states	~51%	GREEN

The richest large state, Karnataka, has a per-capita income of about ₹4.22 lakh. Bihar's is about ₹0.76 lakh. That is a gap of roughly 5.6×, wider than the income gap between many entirely separate countries. The top five states generate close to 48% of national output. Just two states account for around 51% of foreign direct investment. India's growth is not a uniform national phenomenon; it is heavily concentrated in a handful of southern and western states with functioning cities, ports, industrial clusters and administrative capacity, while large populations in the north and east grow from a far lower base and capture a far smaller share of the capital.

For an allocator this dispersion cuts two ways. It means the "1.4 billion consumers" framing is misleading: the addressable, income-earning, formally-employed consuming market is a fraction of the population, concentrated in the high-income states, and that fraction is what listed consumer companies actually sell to. But it also means the growth runway is genuinely long *if* the lagging states converge, because convergence would pull hundreds of millions into the formal, consuming economy. The bet embedded in the demographic dividend and in much of the consumer earnings story is, at bottom, a bet on internal convergence that has been slow and uneven. The state table is the quantitative face of "India is not one economy," and it should temper any model that treats national GDP as if it flowed evenly to a national consumer.

The services engine and what it lists as

India's growth is unusually services-led, and the character of those services shapes how much of the growth an equity investor can touch. At the top sits a genuinely world-class, dollar-earning IT and business-services export sector: formal, listed, profitable, and one of the clearest channels through which Indian growth reaches a public earnings line. That sector is a real asset in the index and part of why the earnings machine in the return bridge is as strong as it is. But it is also mature, exposed to global tech-spending cycles and to the automation of the very work it sells, so it is a source of *quality* earnings rather than a source of the fastest domestic growth.

Beneath it is the far larger domestic service economy – retail trade, transport, construction-adjacent services, personal services – which is where most of the workforce sits and where informality is highest. That layer grows with the domestic economy but lists barely at all; its expansion shows up in GDP and in the incomes of the informal workforce long before, if ever, it shows up in a public company's revenue. So the services tilt cuts both ways for the allocator: it hands the index a high-quality export earner at the top, while burying the bulk of domestic service-sector growth in businesses no foreign portfolio can own. The investable expression of India's services growth is narrower than the sector's share of GDP suggests.

The concentration of investable India

The same concentration that shows up across states shows up in the listed market, and it compounds the transmission problem. The Nifty's top ten stocks are about 52.9% of the index. Half the investable large-cap market is ten companies. So even where GDP growth does reach the listed economy, a foreign investor buying the index is buying a highly concentrated slice of it: a handful of financials, an energy-and-telecom conglomerate, IT majors, and a few consumer names. The breadth of the real economy does not translate into breadth of the investable one. This is a structural reason the growth-to-return transmission is lossy at the top of the chain, before margins, dilution, multiple or currency even enter.

Growth is not returns – the point, stated once, precisely

Having characterised the engine, state the conclusion the scatter forces, and state it carefully, because the careless version is where most growth-to-returns arguments go wrong. Across thirteen markets, higher GDP growth was not associated with higher equity returns: the fitted slope is negative (−0.49), the R^2 is 0.03, the correlation −0.16, and

there are thirteen observations. The correct summary is not "growth hurts returns" – the relationship is too weak and the sample too small to claim a reliable negative link. The correct summary is that there is *no statistically reliable positive relationship*. Growth does not buy returns.

The named points enforce the discipline. Vietnam grew fastest at 6.2% real and returned poorly. Taiwan grew at 3.8% and returned the most, around 22.5% a year in dollars. India grew at 5.8% and returned 8.0%. India is the fastest-growing *large* EM after Vietnam, a precise claim that survives the data, unlike the looser "India is the fastest-growing market," which does not, because Vietnam grew faster. The reason the ordering scrambles is the transmission chain from Section 1: growth gets diluted by issuance, taxed by currency, and pre-paid in the multiple. A market everyone knows will grow trades at a price that already embeds the growth, so the return depends on beating an optimistic price, not on the growth rate itself. India spent most of the last decade as exactly that kind of market.

Does the growth reach the listed economy?

Pull the section together into the question that hands off to the rest of the report. India's growth machine is real, large and fast by large-EM standards. It is also productivity-light, jobs-constrained, heavily informal, geographically concentrated, and expressed through a listed market that is itself concentrated in ten names. Every one of those features is a place where GDP growth leaks before it becomes listed-company earnings, and listed-company earnings are the only part of the growth an equity allocator can own.

So the answer to "does this growth reach the listed, investable economy?" is: partially, and less than the headline implies. The formal, organised, listed slice of India grows faster than the aggregate in the periods when formalisation is taking share – which flatters the earnings line – but it sits on top of an informal base that the index cannot capture, in an economy whose income and capital are concentrated in a few states, and it reaches the foreign investor only after issuance, multiple and currency have each taken their cut. The growth is not the return. The growth is the raw material from which a return might, under specific conditions on margins and multiple, be manufactured. Whether those conditions hold is the profit-cycle question, and that is where the report turns next.

What this means for the allocation

The growth machine is a reason to take India seriously and a warning against pricing it on its GDP rate. For the marginal dollar, three conclusions carry forward. First, discount the demographic dividend to a conditional: it pays only to the extent the formal economy creates jobs, and job creation has lagged the population, so the consumer-earnings tailwind is weaker and slower than the raw population pyramid suggests. Second, treat "India" as a portfolio of very unequal states, not a single consumer market: the addressable, formally-employed, income-earning economy is concentrated (top-five states near 48% of output, two states around 51% of FDI, a 5.6× income gap between the richest large state and the poorest), which caps how broadly consumer growth actually flows. Third, hold the growth-versus-returns finding as a permanent prior: across thirteen markets growth showed no reliable positive link to equity returns, India sits mid-cloud, and the fastest-growing large EM after Vietnam has been a middling dollar performer. None of this argues against owning India. It argues for owning it on the strength of its earnings machine and the price paid for it, and for refusing to pay up on the GDP number alone. The next question is whether the earnings machine, the profit cycle now sitting at a fresh all-time high relative to GDP, can sustain what the allocation is really betting on.

SECTION 3

The State, Infrastructure and Formalisation

India's public sector has spent a decade building the best physical and digital plumbing in the emerging-market world, and almost none of the water it carries is billed by a listed, foreign-investable company.

₹22.08 lakh crore

GROSS GST COLLECTED, FY25 — A RECORD

THE FINDING, UP FRONT. India's public sector has spent a decade building the best physical and digital plumbing in the emerging-market world, and almost none of the water it carries is billed by a listed, foreign-investable company. Formalisation is genuine: output has moved out of the cash shadows and into the GST net, and the receipts prove it. But the taxable economy and the investable economy are two different circles, and the overlap is thinner than the formalisation narrative implies. An allocator who reads "digital public infrastructure" and "₹12 lakh crore of annual capex" as a bull case for listed Indian equity is making an inference the ownership data does not support. This section traces where the state-driven value actually lands, and it lands mostly on balance sheets you cannot buy on the NSE.

Three engines are usually cited as the state's contribution to the equity story: public capital expenditure, GST-led formalisation, and the UPI-anchored digital public infrastructure (DPI) stack. Each is real, each is large, and each monetises into listed equity far more weakly than the headlines suggest. We take them in turn, then close on what that means for a foreign-investable claim.

3.1 The capex machine is real, and mostly sits off the exchange

Central government capital expenditure is the cleanest of the three stories to measure and the hardest to convert into a listed-profit claim.

₹12.2 lakh crore — central government capex budgeted for FY27, versus ₹10.52 lakh crore actually spent in FY25. Roughly 3.1% of GDP, with "effective" capex (including grants-in-aid to states for capital assets) near 4.36%.

CENTRAL GOVT CAPEX	₹ LAKH CR	% OF GDP (APPROX)	CLASS
FY24 (implied actual)	~9.6	~3.3	AMBER
FY25 (CGA provisional actual)	10.52	~3.3	GREEN
FY26 (revised estimate)	~10.9	~3.1	AMBER
FY27 (budget estimate)	12.2	~3.1	GREEN

Sources: CGA provisional actuals (31-May-2025); Union Budget 2026-27 (Business Standard, PRS). FY20–FY23 actuals were not re-verified this cycle and are omitted rather than carried; the widely reported multi-year ramp from roughly ₹3.4 lakh cr in FY20 is directionally correct but treated here as illustrative.

The fiscal discipline underneath the ramp is worth stating because it is the part most likely to survive a downturn. FY25's deficit came in at 4.8% of GDP, below the revised estimate. The FY26 revised deficit sits near 4.4% and the FY27 target is 4.3%. So the capex has been financed inside a narrowing deficit rather than on top of a widening one, which is the difference between a durable investment programme and a pre-election sugar rush. On the state's own terms, this is a success.

The allocation problem begins one step downstream. When the Union spends ₹12 lakh crore on capital assets, the money flows to roads, railways, defence, water, and grid. Ask who books the resulting profit and the answer is dominated by entities an EM allocator either cannot own or would not want as a core position: the National Highways Authority and its EPC contractors (many unlisted or thinly traded), Indian Railways and its captive PSUs, defence public-sector undertakings, and a long tail of state-level construction firms. The genuinely listed, genuinely liquid beneficiaries, the large engineering-and-construction majors and the cement complex, capture a slice of the order book, not the whole of it. Their margins on government work are also structurally thinner than on private orders, because public tenders compete the return down toward the cost of capital by design.

There is no clean primary series that partitions ₹12 lakh crore of capex into "listed-captured profit" versus "PSU or unlisted or subsidy," and we will not manufacture one. The honest statement is a judgment (INTERP): the listed-equity share of incremental public-capex profit is a minority of the spend, concentrated in a handful of large-cap names already trading at valuations that price the order book in. The capex is a GDP story and a formalisation story. It is only a weak listed-earnings story, and it is close to a non-story for the free-float, foreign-investable sliver that this report ultimately cares about.

3.2 GST: formalisation you can measure, ownership you cannot assume

The Goods and Services Tax is the single best-measured piece of India's formalisation, and it does what the bulls claim. Since its mid-2017 launch, gross collections have compounded from a partial-year base to a genuine record, and the active taxpayer base has widened.

₹22.08 lakh crore — gross GST collected in FY25, a record, up 9.4% year on year, against roughly ₹7 lakh crore in the first (partial) year of the tax.

GST GROSS COLLECTIONS	₹ LAKH CR	NOTE	CLASS
FY18 (9 months)	~7.2	GST live from Jul-2017	AMBER
FY22	14.83	post-COVID recovery	AMBER
FY24	20.18		AMBER
FY25	22.08	record, +9.4%	GREEN
FY26	22.27	+0.9% only — see below	AMBER

Sources: PIB "Nine Years of GST" (Jul-2026); Outlook Business/PIB FY25; monthly data via Munim, CAClubIndia. Active GST taxpayers ~1.65 crore.

Two features of the FY26 print matter for anyone using GST as a proxy for economic momentum. First, the flat +0.9% headline is a policy artefact, not a demand collapse. The GST 2.0 rate rationalisation, a move to a two-slab 5/18% structure plus a 40% demerit rate, took effect on 22-Sep-2025 and cut effective rates across the middle of the year. Underlying volume was recovering by the fourth quarter (February +8.1%, March +8.8%, and July 2026 printing ₹2.11 lakh crore, up 15.4% year on year). So the tax base is still formalising even as the tax take flattened. Second, the monthly record of ₹2.43 lakh crore set in April 2026 confirms that the ceiling keeps rising. The formalisation trend is intact.

Here is where the allocator has to resist a comfortable elision. Formalisation moves activity from the informal, untaxed, unincorporated economy into the formal, taxed, incorporated economy. That is a value transfer toward the organised sector, and it is real. But "organised" is a far larger set than "listed," and "listed" is a far larger set than "free-float, foreign-investable." A kirana store that starts paying GST, or an unlisted mid-market manufacturer that registers to claim input credit, has formalised without ever issuing a share an allocator can buy. The GST base of roughly 1.65 crore taxpayers dwarfs the roughly 2,000-odd meaningfully liquid listed companies by three orders of magnitude.

The correct framing is a funnel of formalisation, and it narrows hard at the last stage: informal to formal (measured by GST, and real), formal to incorporated (partly captured by corporate-tax data), incorporated to listed (a small fraction), listed to free-float (about half, given roughly 50% promoter ownership), and free-float to foreign-investable (smaller still). Value entering the top of that funnel is genuine. The share reaching the bottom, where a global EM investor can own it, is a minority of a minority. GST formalisation lifts the tax-to-GDP ratio and the state's fiscal

capacity. Its transmission into listed corporate profit is real but diluted, and into foreign-ownable listed profit, thin. This is the same discipline the report applies to the "Four Indias" ladder in Section 2, imported here: never let a number measured at one stage of the funnel do argumentative work at another.

There is a second-order effect that does reach listed equity, and it is worth crediting because it is the strongest link in the formalisation chain. As activity formalises, market share migrates from unorganised competitors who did not pay tax to organised players who do. A branded, GST-compliant, listed manufacturer or retailer competing against a cash-economy rival gets a durable cost-and-scale advantage when the rival is forced onto the same tax footing. That share-shift is a genuine tailwind for organised-sector listed companies, and it is the mechanism most worth owning in the formalisation story. But it is a slow, sector-specific grind, most visible in building materials, branded staples, footwear, jewellery, and organised retail, and it is already well understood and largely priced by the market. It is a reason to prefer organised-sector leaders, not a reason to re-rate the index. Formalisation gives the listed consumer and industrials complex a structural share-gain story; it does not hand the aggregate market a new growth leg.

3.3 UPI and DPI: the world's best rails, monetised by almost nobody you can own

The digital public infrastructure story is where the gap between national achievement and investable claim is widest, and where the temptation to conflate the two is strongest.

₹314 lakh crore across **24,162 crore transactions** — the value and volume run through UPI in FY26, up ~20.6% and ~30.0% respectively. UPI now clears roughly **49% of the world's real-time payments** and about **85% of India's digital payments, across 703 banks.**

These are extraordinary numbers, and by July 2026 the system set a fresh monthly record of 23.66 billion transactions. As national infrastructure, UPI is a genuine wonder: near-universal, interoperable, and effectively free at the point of use. The question for this report is narrow and unsentimental. Who earns the economic rent on ₹314 lakh crore of payment flow, and can an allocator own them?

The answer is structural and it is the crux of the DPI-as-equity-thesis problem: **UPI runs at zero merchant discount rate (MDR)**. There is, by policy, essentially no per-transaction interchange to capture on person-to-merchant UPI. The rails were built as a public good and priced as one. That decision, defensible on financial-inclusion grounds, is precisely what starves the listed-equity monetisation.

Trace the value chain and the listed capture keeps evaporating:

DPI LAYER	WHO OPERATES IT	LISTED AND FOREIGN-INVESTABLE?
The UPI switch / network	NPCI (National Payments Corporation of India)	No — not-for-profit, bank-owned utility
Largest consumer PSP apps	PhonePe, Google Pay	No — controlled by Walmart and Alphabet respectively; foreign parents, and the India entities were not independently listed as of the cut-off
Listed domestic PSP	Paytm (One97)	Yes, but a small-cap with a contested path to profit, and MDR-free UPI is not its earnings engine
The identity/ consent layer	Aadhaar, Account Aggregator, DigiLocker	No — public or regulated infrastructure

Ownership characterisations of the PSP apps are widely reported facts treated here as illustrative (INTERP); the ₹314 lakh cr, 24,162 crore, 49% and 85% figures are GREEN (PIB "UPI turns 10").

The monetisation that does exist accrues largely off the exchange or abroad. NPCI, the entity that actually clears the flow, is a not-for-profit utility that distributes no shareholder return. The two apps that intermediate the majority of consumer volume are controlled by foreign strategics whose India payments arms were not separately investable at the cut-off. Value is being created at enormous scale; it is simply not being pooled into a liquid, listed, domestically taxable, foreign-ownable claim. Where DPI has produced investable businesses, it is one layer up, in the lending, insurance-distribution, and wealth-tech franchises that ride on top of the rails and monetise through spreads and fees rather than through payments interchange. Those are real, and some are listed. But that is a fintech-credit story wearing DPI's clothing, and it should be underwritten as credit risk, not as a toll on payment volume.

The contrast with how payments monetise elsewhere is the sharpest way to see the gap. In the United States, a comparable volume of consumer payments would throw off tens of basis points of interchange on every swipe, pooled into a small number of large, listed, foreign-ownable networks and processors that are among the most profitable franchises in global equity. India built rails that clear more real-time transactions than the rest of the world combined and set the toll to zero. The consumer and the merchant capture the surplus; the state and the banking system absorb the cost; the equity investor captures essentially none of it directly. This is not a criticism of the policy, which has plausibly done more for financial inclusion than any single measure in India's history. It is a warning against a specific analytical mistake: importing a Western "payments TAM times a take-rate" model onto a system whose take-rate is, by law, approximately nothing. The data-layer of India's DPI is a public good, and public goods do not have equity holders. The one place the model reasserts itself is credit: the account-aggregator and UPI data exhaust make underwriting cheaper and faster, and lenders who exploit that will earn spread. So the investable DPI thesis is, correctly stated, a lending thesis, and it lives or dies on credit costs through a cycle, not on transaction volume through a boom.

The honest summary for the allocator: DPI's listed-equity monetisation is thin. UPI is a subsidy from the state (and from the banks that absorb the switching cost) to consumers and merchants. It lowers the friction of the whole economy, which shows up eventually and diffusely in consumption and formalisation. It does not show up as a fat, ownable margin pool the way card interchange does in the United States. Anyone modelling "India's payments TAM" into a listed equity target is modelling a rent that policy has deliberately set to zero.

3.4 State capacity: the real asset, imperfectly investable

Step back from the three engines and the through-line is a state that has become materially more capable: it can spend ₹12 lakh crore of capex inside a shrinking deficit, collect ₹22 lakh crore of GST across 1.65 crore registrants, and operate payment rails that clear half the planet's real-time transactions. That capacity is itself an asset. It lowers macro tail risk, it underwrites the ~\$693 billion reserve buffer (Section 4), and it is a large part of why India's currency is the lowest-volatility depreciator in the EM set rather than a serial crisis currency.

But state capacity is an asset the allocator holds indirectly, as a reduction in the risk premium, and not as a cash flow. It compresses the discount rate a little; it does not add to the numerator. And it comes with the standard caveats that no amount of digital plumbing removes: execution still varies enormously by state (the ~5.6× income gap between the richest large state and Bihar, revisited in Section 5, is a governance gap as much as an income one), the capex programme is a political commitment that a future government could redirect, and zero-MDR is a policy that could change in either direction, each of which would rewire the DPI monetisation question overnight.

3.5 ALLOCATION TAKEAWAY

Formalisation is one of the most real things happening in India, and one of the least investable. The state has moved value into the taxed economy at scale; it has moved far less into the listed economy, and less still into the free-float, foreign-investable economy that ends this report's funnel. Public capex is a minority listed-profit story concentrated in already-priced large caps. GST proves formalisation without proving listed capture. DPI is a national triumph monetised, by deliberate policy, at close to zero rent for anyone an EM investor can own.

For the allocator, three consequences. First, discount the "DPI/formalisation flywheel" as a direct earnings thesis; treat it as a risk-premium compressor, worth a few tens of basis points on the discount rate, not a leg of EPS growth. Second, where formalisation does convert to listed profit, it concentrates in financials and organised-retail consumer names, which is where to look, and which is also where valuations already embed it. Third, be sceptical of any target price built on a payments or capex "TAM": the state has priced the toll booths at zero and keeps the order book on balance sheets you cannot buy. The formalisation of India is a fact. Its transmission into a foreign-investable claim is the weakest link in the entire chain, and it is the manufacturing sector, not the state, that carries the next test of whether India can convert national scale into ownable profit.

SECTION 4

Can India Manufacture Its Next Leap?

India is doing real manufacturing volume and capturing thin manufacturing value, and the profit that does get captured lands disproportionately on foreign or unlisted balance sheets.

\$333.19bn

FY26 MERCHANDISE TRADE DEFICIT

THE FINDING, UP FRONT. India is doing real manufacturing volume and capturing thin manufacturing value, and the profit that does get captured lands disproportionately on foreign or unlisted balance sheets. The "China+1" and production-linked-incentive (PLI) narrative is not fabricated; assembly lines exist, exports of assembled electronics have risen, and headline FDI has arrived. But at every link of the chain from announced investment to a foreign-investable profit claim, value leaks out, and the largest single fact of the trade account points the wrong way: India's goods deficit with China is at a record. This section walks the chain deliberately, because the manufacturing thesis is almost always sold at the announcement stage and almost never audited at the profit stage.

A caution on data before we start. India's manufacturing story is under-served by clean, primary, replication-grade numbers, and this section is honest about that. The trade aggregates are GREEN (official PIB release). The China bilateral deficit, PLI disbursement rates, FDI trajectory, and iPhone value-added figures are widely reported but are not in this report's canonical dataset; they are used here as illustrative context (INTERP) and flagged as such wherever they appear. We do not headline a number we did not verify.

4.1 The chain, and where value leaks

The correct way to audit a manufacturing thesis is as a chain of conversions, each with a leakage:

Announced investment → capacity → production → domestic value added → listed profit → foreign-investable claim.

The bull case lives at the left end (big announcement numbers) and the allocation case can only be settled at the right end (ownable profit). The characteristic error in this literature is to treat the announcement and capacity stages as if they proved the profit stage. They do not. Each arrow below loses value.

4.2 The trade account: volume up, deficit at a record

Start with what is measured cleanly.

\$333.19 billion — India's FY26 merchandise trade deficit, on **\$441.78bn** of goods exports against **\$774.98bn** of goods imports. Services exports of **\$418.31bn** nearly close the gap, taking total exports to **\$860.09bn** and holding the current account near balance.

INDIA TRADE, FY26	US\$ BN	CLASS
Merchandise exports	441.78	GREEN
Merchandise imports	774.98	GREEN
Merchandise trade deficit	333.19	GREEN
Services exports	418.31	GREEN
Total exports (goods + services)	860.09	GREEN
Current account balance (FY26) -0.6% of GDP (~\$25.2bn)		AMBER

Sources: PIB PRID 2252272 (15-Apr-2026); RBI BoP Q4 FY26 (via Business Standard).

Two things are simultaneously true here, and the manufacturing thesis needs both stated. The current account is close to balanced at -0.6% of GDP, which is the macro-stability win: India runs a huge goods deficit but plugs most of it with services exports (IT, GCCs, business services) and remittances. That is genuinely reassuring for the currency

and the reserve position (~\$693bn at the cut-off, against an all-time high near \$728bn in February 2026). It is also, quietly, an admission about where India's tradable competitiveness actually sits. The country closes its trade gap with services, the sector it already dominates, not with the manufactured goods the "next leap" is supposed to deliver.

The services point is not a digression from the manufacturing question; it is the answer to a version of it. Services exports of \$418bn are within touching distance of the entire \$442bn merchandise export book, and services are where India has a real, defended, high-value-added, globally competitive position. The uncomfortable implication for the manufacturing thesis is that India's proven path to earning foreign exchange and building globally competitive firms has run through services, not factories. An allocator who wants exposure to Indian tradable competitiveness has, historically, been better served by the IT and business-services complex than by any manufacturing bet, and nothing in the trade account overturns that yet. The reserves buffer that manufacturing bulls point to as evidence of strength was built substantially on services and remittances. That buffer is real, and it is a large part of why the currency is the lowest-volatility depreciator in the EM set. But it is a services-and-remittances achievement wearing a manufacturing-optimism headline, and the two should not be confused when sizing a manufacturing position.

And underneath the aggregate sits the fact that reframes the whole China+1 story:

~\$99–106 billion (illustrative) — India's goods trade deficit with China, at or near a record. The single largest bilateral imbalance in the trade account.

This figure is not in the canonical dataset and is used as illustrative context, but its direction is not in dispute and it is corroborated by the structure of the official numbers: import dependence on Chinese intermediate and capital goods has risen precisely as assembly has scaled. The "China+1" line, told as India substituting for China, collides with a trade account in which India imports more from China than ever, much of it the components, machinery, and inputs that feed the very assembly lines being celebrated as diversification. India is, at this stage, less an alternative to the Chinese supply chain than a new downstream node bolted onto it.

The mechanism is not a contradiction once you separate final assembly from the supply chain behind it. China+1 means global buyers diversifying the *location of final assembly* away from China. It does not, at this stage, mean diversifying the *source of components*, because the component ecosystem, precision tooling, chemicals, display and semiconductor supply, machinery, remains concentrated in China and East Asia and takes a decade-plus to relocate. So an India that wins final-assembly mandates mechanically increases its imports of Chinese inputs to feed those lines. Rising Indian electronics exports and a rising India-China deficit are the same phenomenon viewed from two ends of the line, not opposing forces. The optimistic reading is that assembly is the beachhead and localisation of components follows, deepening domestic value added over time; there is real evidence of that beginning in electronics. The realistic reading for an allocator underwriting today is that the deepening is early, slow, and capital-intensive, and that the trade account will get worse against China before the value-added mix improves. The leap is a decade project being sold on a two-year narrative.

4.3 Assembly is not value addition: the iPhone problem

The flagship of the manufacturing story is electronics, and specifically smartphone assembly. Exports of assembled phones have risen sharply, and the political economy treats this as arrival. The value-added arithmetic is more sobering.

Low single-digit to ~10% (illustrative) — the share of an assembled smartphone's ex-factory value that accrues as Indian domestic value added at the current stage of the supply chain. The display, the system-on-chip, the memory, the camera module, and the precision components are overwhelmingly imported; India contributes final assembly, test, and pack labour.

The number is widely cited around the high-single-digit range and is used here illustratively (INTERP), but the mechanism is robust regardless of the exact figure. A \$1,000 phone exported from India adds far less than \$1,000 to Indian GDP or to Indian corporate profit; it adds the assembly margin plus the wage bill plus a PLI subsidy, against a large imported bill of materials that shows up on the other side of the trade account as... imports from China and East Asia. This is why the export figure and the China-deficit figure move together rather than in opposition. Counting the gross export value as "India's manufacturing leap" double-counts the imported content. The domestic-value-added leg is the one that matters, and it is thin and only slowly thickening as component ecosystems localise.

This is the same category error the report demolishes elsewhere in the "domestic revenue intensity" work: end-market geography (where the phone ships) is not production geography (where the value is made) is not profit geography (whose balance sheet books the margin). For manufacturing, the three diverge sharply, and only the third is ownable.

The trajectory does matter, and the fair version of the bull case lives here. Value added rises as component supply localises, as PCB, camera modules, casings, and eventually displays and packaging get made in India rather than imported. Each step localised lifts domestic value added on the same export unit and, over time, thickens the thin assembly margin into something more like a supply-chain margin. India is genuinely at the start of that curve in electronics, which is more than can be said for most of the "next leap" sectors. But the allocator has to price the curve, not the destination. At today's position, the exported unit is mostly imported value with an Indian label and an Indian wage bill, and the subsidy that closes the economics is a government transfer, not a market margin. Underwriting the destination as if it were the present is exactly how announcement-stage numbers get mistaken for profit-stage numbers.

4.4 PLI, FDI, and the ownership question: who books the profit

The production-linked incentive schemes are the policy engine, and their record is mixed in a specific, allocation-relevant way. Announced and committed investment under PLI is large; actual disbursement has run well behind announcement, and the schemes that have delivered volume (electronics, pharma) have done so in sectors where the marginal producer is either a foreign contract manufacturer or a domestic unlisted or promoter-heavy entity. Both the disbursement shortfall and the sector pattern are widely reported and are treated here as illustrative (INTERP), because this report has no primary PLI-disbursement series.

Follow the profit and the ownership problem becomes concrete:

- **Contract electronics assembly** is dominated by foreign firms (the large Taiwanese and other East Asian contract manufacturers). Their India profit, to the extent it exists after a thin assembly margin, accrues to a foreign-listed or foreign-private parent. An EM-India allocator does not own it by owning India.
- **Domestic champions in the space** are frequently unlisted (large conglomerate electronics arms) or are recently listed small- and mid-caps trading at valuations that already price years of flawless execution. The listed, liquid, foreign-investable proxies for the manufacturing leap are a short list of expensive small caps, not a broad, cheap, large-cap opportunity.

- **Foreign strategic ownership** shows up even in the listed tape: roughly 8% of NSE-listed value sits with foreign promoters (NSE, Mar-2026), a reminder that a meaningful slice of "Indian" listed manufacturing profit is already foreign-controlled at the shareholder level.

Layer this onto the report's ownership spine: promoters hold about **50%** of the listed market, foreign portfolio investors about **17.3%**, domestic institutions **16.1%**, and retail **18.5%**. A foreign portfolio allocator is buying a claim on the free float, roughly half the listed market, of which manufacturing is a modest and expensively priced slice, and much of the genuinely fast-growing manufacturing profit is being booked one rung outside that claim, on unlisted or foreign-parent balance sheets. The ownership structure is not incidental to the manufacturing thesis; it is the thesis's binding constraint. Even if every optimistic assumption about capacity, production, and value-added localisation comes true, the profit still has to arrive on a listed, free-float, foreign-ownable balance sheet before it becomes a return for the allocator, and the prevailing structure, high promoter holding, foreign contract manufacturers, unlisted conglomerate arms, wholly-owned FDI, routes it almost everywhere else first.

FDI is the last leak worth naming, and it is the most complete. Headline gross FDI has been substantial, but net FDI has compressed as repatriation and outward investment have risen (widely reported; illustrative here). The distinction between gross and net is itself an ownership tell: rising repatriation means foreign parents pulling profit out of their Indian subsidiaries, which is exactly the profit an EM-India allocator does not share in. And the structural point is starker than the trend. Foreign direct investment that builds a wholly-owned Indian plant creates Indian jobs, Indian GDP, and Indian tax; it creates no listed, free-float, foreign-portfolio-investable equity at all. It is, from the portfolio allocator's seat, the most extreme version of the leakage: maximal economic contribution, zero direct ownable claim. The better India does at attracting wholly-owned FDI manufacturing, the more of its industrial value added sits in exactly the legal structure a portfolio investor cannot touch. The macro win and the equity non-event are, in the FDI case, the same transaction.

It is worth being explicit about what a genuinely investable manufacturing exposure would require, because the absence of it is the finding. It would need a large, liquid, domestically listed company, majority free-float rather than promoter-locked, capturing a thick and rising share of domestic value added, in a sector where the profit pool is not competed away on subsidised tenders, at a valuation that does not already discount a decade of perfect execution. Very few names clear all five bars at once. The contract manufacturers fail the "domestically listed and free-float" bar; the conglomerate electronics arms fail the "listed" bar; the capex beneficiaries fail the "profit pool not competed away" bar; and the handful of pure-play listed proxies that clear the first bars fail the valuation bar. That is not an argument that no money can be made in Indian manufacturing equity. It is an argument that it is a stock-picker's minefield, not an index-level allocation, and that the burden of proof sits on the specific name, not on the national narrative.

4.5 The chain, totalled

Put the leakages end to end:

STAGE	WHAT IT DELIVERS	THE LEAK TO THE NEXT STAGE
Announced investment	Large headline commitments (PLI, FDI MoUs)	Disbursement and realisation run behind announcement
Capacity	Real assembly lines, real plants	Capacity utilisation and margins compete down on subsidised, competitive tenders
Production	Rising export volumes (electronics)	Gross export value overstates domestic content
Domestic value added	Assembly margin + wages + subsidy	Thin (single-digit to ~10% for flagship electronics); imported bill of materials feeds the China deficit

STAGE	WHAT IT DELIVERS	THE LEAK TO THE NEXT STAGE
Listed profit	A modest, expensive small/mid-cap slice	Much manufacturing profit sits unlisted or foreign-parent
Foreign-investable claim	A short list of high-multiple names	~50% promoter-held, ~8% foreign-promoter; free float is small and priced

Stages 1, 4, 5 use illustrative (INTERP) inputs; stage 3 export values and the ownership percentages are GREEN.

4.6 ALLOCATION TAKEAWAY

The manufacturing leap is happening as a GDP and employment event and is leaking, at every conversion, out of the foreign-investable profit pool. India is capturing volume and thin value; the value it captures is disproportionately booked by foreign contract manufacturers, unlisted domestic arms, or wholly-owned FDI subsidiaries, none of which an EM-India portfolio owns. The listed proxies that do exist are a narrow set of small and mid caps already priced for perfection, and the trade account, with a record China deficit sitting behind rising electronics exports, says the supply chain is being deepened as much as diversified.

For the allocator, the discipline is to underwrite manufacturing as a bottom-up, name-by-name stock-picking problem with a high bar on valuation and on ownership structure, and to refuse it as a top-down, own-the-index thesis. "Buy India for the manufacturing renaissance" fails at the ownership stage before it fails at the valuation stage. The macro dividend of manufacturing (jobs, a stable current account, a deeper industrial base) is real and accrues to the country. The equity dividend is small, expensive, and partly foreign-owned before you arrive. Whether India's growth reaches the allocator at all now turns on the demand side, the household consumer, and on the same test that defeated the state and the factory: does the money show up as listed revenue.

SECTION 5

The Indian Consumer

Household demand is the revenue line of the entire Indian equity market, and the part of that demand that shows up in listed revenue is the top of a K, not the middle of a bell curve.

7.0% of GNDI

NET HOUSEHOLD FINANCIAL SAVINGS, FY25

THE FINDING, UP FRONT. Household demand is the revenue line of the entire Indian equity market, and the part of that demand that shows up in listed revenue is the top of a K, not the middle of a bell curve. Aggregate consumption is large and growing; listed consumer revenue has recently grown slower than nominal GDP; and the gap between the two is the mass market, which spends in the informal and unlisted economy and buys the volume categories that have been flat. The household balance sheet, meanwhile, has just staged a rebound in net financial savings that looks reassuring and is, on inspection, a deleveraging, itself a demand-cautious signal. This section separates what is measured from what is judgment, because the consumer story is where India's narrative and India's data diverge most.

A note on what is measurable. Household financial-savings data is real and official (RBI), but it carries a denominator trap that much published commentary falls into, and that this section corrects. Income distribution at the granular level, and the "K-shape / premiumisation" framing, are largely judgment (INTERP) supported by fragments; they are labelled as such. The link from demand to listed revenue is the one thing we can quasi-measure, via the corporate revenue-versus-GDP gap, and it is where the allocation weight sits.

5.1 The household balance sheet: a rebound that is really a retrenchment

Start with the most widely mishandled number, because handling it correctly changes its meaning.

7.0% of GNDI — net household financial savings in FY25, up from **5.8%** in FY24 and a multi-decade low of **5.1%** in FY23. The rebound is real. Its cause is not rising thrift.

HOUSEHOLD SAVINGS (RBI)	FY23	FY24	FY25	DENOMINATOR
Net financial savings	5.1	5.8	7.0	% of GNDI*
Gross financial savings	—	12.1	11.8	% of GNDI
Financial liabilities (flow)	—	6.4	4.8	% of GNDI
Gross domestic savings	—	32.3	34.2	% of GNDI

The FY23 5.1% figure was originally released on a % of GDP basis; the FY24–FY25 revised series is % of GNDI. The denominators differ and should not be spliced naively. Source: RBI Annual Report 2025-26 via Business Standard (29-May-2026); earlier vintage via Business Standard (Sep-2023).

Two corrections to the common treatment. First, the denominator: the RBI net-household-financial-savings series is expressed as a percentage of **Gross National Disposable Income (GNDI)**, not GDP. GNDI is larger than GDP (it adds net factor income and net current transfers from abroad, which for India means remittances), so a given rupee of saving is a slightly smaller percentage of GNDI than of GDP. The figure is commonly quoted against GDP and headlined as a "50-year low." The level was roughly right for FY23; the "50-year" is an overstatement (the press characterisation was a 47-year low, and it rests on a series that cannot be fully reconstructed), and the number still circulating is in any case two years stale.

Second, and more important for demand: the FY25 rebound to 7.0% did not come from households saving more. Gross financial savings actually *fell* (12.1% to 11.8% of GNDI). The net number rose because households *borrowed less*: the liabilities flow dropped from 6.4% to 4.8% of GNDI. Net savings went up because deleveraging outran the decline in gross saving. That is a materially different signal than a thrift revival. Households pulling back on new borrowing is, mechanically, households pulling back on the credit that funds discretionary consumption, the car loan, the consumer-durable EMI, the credit-card spend. The balance-sheet number that looks like household strength is, read correctly, a household caution light for the demand line. It is neither a distress signal (households are not dis-saving to survive, as the FY23 low implied) nor a boom signal. It is retrenchment.

The FY23 episode and the FY25 rebound are, taken together, a two-act warning rather than a low-then-recovery reassurance. Act one: households ran their net financial savings down to a multi-decade low by leaning on credit to sustain consumption, the debt-stock reaching roughly 37.6% of GDP. Act two: households stopped leaning on credit, which repaired the savings ratio but did so by throttling the very borrowing that had propped up discretionary demand. Both acts are consumption-negative at the margin relative to a healthy baseline, the first because it was unsustainable, the second because its correction removes a demand crutch. The reading an allocator should carry forward is that the Indian household entered the period with less spare financial capacity than the aggregate consumption headlines implied, and has spent the subsequent year rebuilding balance-sheet caution rather than spending power. Gross domestic savings at 34.2% of GNDI remain high by global standards, so this is not a fragile household sector in absolute terms. It is a household sector whose *marginal* propensity to fund new consumption with credit has cooled, which is precisely the margin that drives incremental listed consumer revenue.

5.2 Income distribution: the K is a judgment, and the judgment is well-supported

The granular income distribution India would need to nail the "K-shape" claim is not in this report's primary dataset, and we will not invent a Gini or a decile table. What is measurable is the geographic dispersion, and it is severe.

~5.6x — the per-capita income gap between the richest large state (Karnataka, ~₹4.22 lakh) and the poorest (Bihar, ~₹0.76 lakh).

A gap of that magnitude between two states of the same country is, on its own, most of the K-shape argument. India does not have one consumer; it has an affluent, urban, formal-sector consumer whose income and consumption look like an upper-middle-income country, and a mass consumer whose income looks like a low-income one, and the two are separating rather than converging. The premiumisation thesis, that the top of the distribution is spending strongly on aspirational and premium categories while the bottom's volume consumption stagnates, is a judgment (INTERP) rather than a measured finding in this dataset. But it is corroborated from several independent directions: the deleveraging in the household data (the marginal borrower is the mass consumer), the geographic income gap, and the corporate revenue pattern below, in which premium-skewed listed names have out-earned mass-market volume categories. The report states it as judgment, and rates the judgment as well-supported.

It is worth naming what the K-shape does to the "billion consumers" headline, because the headline is where allocation mistakes originate. A market of 1.4 billion people is not a market of 1.4 billion customers for any given listed company. The addressable top of the K, the households with discretionary income to spend on branded, formal, premium goods and services, is perhaps a tenth to a fifth of the population by most reasonable cuts, still an enormous absolute number, and it is this cohort, not the median Indian, that populates the revenue base of the listed consumer complex. The investment implication is that India's consumer growth is best understood as the growth of that top cohort plus the slow formalisation of the layer beneath it, not as the arithmetic of aggregate population times rising average income. Two economies with the same headline GDP-per-capita path can deliver very different listed-consumer revenue depending on where in the distribution the marginal rupee of income lands. In India it has been landing, disproportionately, at the top, which is good for premium listed names and poor for the mass-volume thesis that periodically excites the market.

5.3 The load-bearing question: does demand reach listed revenue?

This is the only place in the consumer story where the allocation weight actually rests, and it is quasi-measurable. If household demand is the market's revenue line, the test is whether that demand converts into *listed* revenue growth. Recently, it has not kept pace with the economy.

Listed revenue has lagged nominal GDP. In FY25, the 1,000 largest non-financial listed companies grew revenue **+6.4%** against nominal GDP of roughly **+9.8%** (an independent Business Standard computation); the Nifty-500 revenue line was weaker still on the broker series. FY26 shows a recovery, with RBI's listed private-corporate sample posting sales **+10.1%**.

REVENUE VS NOMINAL GDP	GROWTH	SOURCE	CLASS
BS1000 non-financial revenue, FY25	+6.4%	Business Standard (independent)	AMBER
Nominal GDP, FY25	~+9.8%	MoSPI (old base)	GREEN
RBI listed private-corporate sales, FY26	+10.1%	RBI (24-Jun-2026)	AMBER
Nominal GDP, FY26	+8.9%	MoSPI (new base)	GREEN

The FY25 gap is the finding, and it is the mirror image of the profit story in Section 6. India's listed corporate *profit* hit a fresh all-time high (profit/GDP 5.2% in FY26, an all-time high, profit +15.6% against nominal GDP +8.9%), while listed corporate *revenue* grew slower than the economy in FY25. Profit rose faster than sales because of margin expansion, not because the top line captured the economy's growth. That combination, record profit share on lagging revenue share, tells the consumer story precisely: the demand that reaches listed companies is being converted to profit efficiently, but a shrinking share of the economy's demand is reaching listed companies in the first place. The mass consumer's spending is happening in the informal, unlisted, and unbranded economy, or in categories (staples volumes, entry-level two-wheelers) that were flat, while the listed tape is weighted toward the premium consumer and the financials that lend to them.

The durability question follows immediately, and it is where the allocator earns or loses money. Margin-led profit growth on lagging revenue is, by construction, harder to repeat than volume-led profit growth: margins mean-revert, revenue compounds. A market at a fresh all-time-high profit share, built on margin rather than on the top line capturing the economy, is a market whose forward earnings depend on either margins holding at an elevated level or revenue finally accelerating to carry the baton. The FY26 recovery in listed sales to +10.1% is the first evidence the baton might pass, and it is why the profit high is not simply a bubble to be faded. But the revenue acceleration has to be sustained and broad, reaching the mass consumer and not just the premium cohort, for the profit level to be defensible rather than a cyclical peak. The consumer data, a household sector deleveraging and saving cautiously, is not yet corroborating a durable, broad-based revenue acceleration. It is corroborating a narrow, premium-led one. That distinction is the difference between the profit high being a new plateau and being a top.

For the allocator this resolves a common confusion. "India has 1.4 billion consumers and rising incomes" is true and is *not* the same statement as "listed Indian consumer companies will grow revenue faster than GDP." The first is demographics; the second is a claim about where in the K the listed companies sit and whether the mass market is formalising into their revenue fast enough to offset premiumisation's narrow base. The FY25 revenue gap says: not yet. The FY26 recovery to +10.1% sales says the gap can close cyclically. Neither says the structural leak, mass demand bypassing listed revenue, has been fixed.

5.4 What the consumer means for the domestic bid

There is a second channel through which the consumer reaches the market, and it is the one that has actually moved prices: the household as an equity buyer rather than as a customer. The same households whose financial savings just retrenched are the source of the domestic institutional bid that has become the market's dominant flow.

₹31,781 crore per month of SIP inflows (June 2026), feeding a domestic institutional bid that bought **+₹6.07 lakh crore** in FY25 against foreign portfolio *selling* of **-₹1.31 lakh crore**, a **4.6×** absorption.

This is the consumer's most direct and most measurable contact with listed equity, and it runs backwards from the revenue story. The household is a stronger buyer of the market than the market is a seller to the household. Mutual-fund equity assets stand near ₹37 lakh crore, and the monthly SIP flow, while only a slice of the total domestic pool, is the marginal price-setter. The caution here connects back to the balance sheet: the SIP flow is drawn from the same household financial-savings pool that just showed gross financial savings *falling*. A retrenching household sector that is simultaneously routing a record share of its financial savings into equities is a composition shift (from deposits and gold toward funds) more than a rising-tide story, and composition shifts can stall or reverse. The domestic bid is a genuine shock absorber for the market, evidenced in FY25. It is not yet a proven permanent floor, and its fuel is the same household saving that the demand story needs to see spent.

The composition shift deserves one more turn, because it is where the consumer-as-saver and the consumer-as-customer stories interlock in a way that is easy to miss. Indian households have historically parked wealth in physical assets, property and gold, and in bank deposits. The financialisation of the last decade, the migration of that saving into equity funds, insurance, and pensions, is a genuine structural change and a large part of why the domestic bid has grown. But financialisation of a broadly flat savings pool is a one-time reallocation, not a perpetual-motion inflow. If gross household financial saving keeps drifting down, the equity flow becomes a fight for share of a shrinking pie against deposits, gold, and debt service, rather than a rising tide lifting all of them. The bull case for the domestic bid requires either the savings rate to re-accelerate or the equity share of savings to keep climbing from an already-elevated base; the bear case is that a cautious, deleveraging household eventually slows both its consumption and its SIP. The two channels are fed from one tank. That is the single most important thing to hold about the Indian consumer: the money that is supposed to grow listed revenue and the money that is holding up listed prices come from the same household balance sheet, and that balance sheet is currently signalling caution, not abundance.

5.5 ALLOCATION TAKEAWAY

The Indian consumer is real, large, and growing, and reaches listed equity through two channels that point in opposite directions. As a customer, the consumer under-delivers to the listed revenue line: mass demand leaks into the informal and unlisted economy, listed revenue grew slower than GDP in FY25, and the profit record was built on margin rather than on capturing the top line. As a saver, the same consumer over-delivers to the listed market's *price*, via a domestic institutional bid that dwarfs foreign selling. The household balance sheet sits under both, and its recent "improvement" is a deleveraging that cautions the demand side even as it funds the flow side.

For the allocator, three consequences. First, buy the consumer through the top of the K and through the financials that intermediate it, because that is where household demand actually converts to listed revenue and profit; be sceptical of "mass consumption" theses that assume the informal spender formalises into your revenue on your timetable. Second, treat the domestic bid as a support for valuation, not as a substitute for earnings; it can hold the market above where fundamentals alone would place it, which is a risk to entry price as much as a comfort. Third, watch the household savings mix, not just the level: gross financial savings falling while net rises on deleveraging is a signal that the fuel for both consumption and the SIP flow is tighter than the headlines read. The consumer determines the revenue line; the revenue line, converted at whatever margin the cycle allows, determines the corporate profit that Section 6 takes up next, where the record profit share and the lagging revenue share finally meet.

SECTION 6

From GDP to Corporate Profit

India's listed-corporate profit share of GDP made a fresh all-time high in FY26, and that single fact forces a reassessment of the bear case, not a footnote to it.

5.2%

NIFTY-500 PROFIT/GDP, FY26 — ALL-TIME HIGH

THE FINDING, UP FRONT. India's listed-corporate profit share of GDP made a fresh all-time high in FY26, and that single fact forces a reassessment of the bear case, not a footnote to it. Nifty-500 aggregate profit reached 5.2% of GDP, marginally above the pre-crisis peak of FY08, with profit growing 15.6% against nominal GDP of 8.9%. A reasonable prior treated the FY20-to-FY25 profit surge as a largely one-off normalisation, a tax cut and a credit-cost rebound working themselves out, near exhaustion. FY26 does not fit that story. Profit outran the economy again, and the ratio set a new high rather than mean-reverting toward its historic range. So the "cyclical, about to revert" leg has to be genuinely reassessed. The honest reassessment lands in an uncomfortable place for both camps: the profit level is now structurally higher than the 2003-19 range and is real, not an accounting mirage, but a fresh all-time-high profit share sitting under an elevated CAPE is a poor place to buy forward asymmetry. And the profit that does exist lands mostly on listed balance sheets a foreign portfolio only partly owns, which is where Section 7 takes over.

A caution on sourcing before the argument, because it governs how hard the headline can be pushed. The entire profit-to-GDP architecture in this section, the 5.2% ratio, the FY08 comparison, the profit-versus-GDP growth gap, traces to a single series: Motilal Oswal's India strategy work (11-Jun-2026). Single-source dependence of this kind is exactly what an evidence standard exists to catch. We disclose it at every use. The direction is corroborated by the broader listed aggregate and by the RBI listed-private-corporate sales recovery cited in Section 5, but the precise ratio is one house's construction, and a reader should treat the second decimal place as that house's, not as an independently replicated fact.

6.1 The number that forces the reassessment


Start with the level, because the level is the finding.

5.2% of GDP — Nifty-500 aggregate profit in FY26, a fresh all-time high, marginally above the FY08 pre-crisis peak. On the broader listed India Inc. measure the figure is **5.7%**, an 18-year high. Profit grew **+15.6%** in FY26 against nominal GDP of **+8.9%**.

EXHIBIT 10

Profit-to-GDP made a fresh all-time high in FY26 – not an exhausted cycle

Nifty-500 aggregate profit as % of GDP. FY26 5.2% > FY08 peak 5.2%; profit +15.6% vs nominal GDP +8.9%.



Source: Motilal Oswal India Strategy, Jun-2026 [GREEN, single-source – flag]. Kakde Research.

The shape of the series matters as much as the endpoint. Profit/GDP ran near 5.2% at the FY08 peak, collapsed through the post-crisis and pre-COVID years to a trough of about 2.1% in FY20, and has climbed almost without interruption since: 4.2% by FY22, 4.7% across FY24 and FY25, and 5.2% in FY26. Two features of that path are load-bearing for the allocation.

First, the trough-to-peak move is enormous. Profit share more than doubled from FY20 to FY26, from roughly 2.1% to 5.2%. A doubling of the profit-to-GDP ratio inside six years is the largest single tailwind Indian equities enjoyed over the period, larger in its earnings effect than any re-rating, and it is the mechanical reason the earnings leg of the twenty-year return bridge in Section 1 came out as strong as it did. Earnings compounded not only because the economy grew but because corporate profit took a rising share of that economy.

Second – and this is the point the exhausted-cycle reading misses – the ratio did not stall at its old ceiling. The FY08 peak of 5.2% could have been read as a natural cap, the level at which India's profit share tops out before competition and cost inflation erode it. FY26 broke through it. That is the fact that weakens the "near-exhausted normalisation" thesis. If the FY20-to-FY25 recovery were purely the reversal of a depressed trough, the ratio would have converged back toward its long-run average, somewhere in the 3.5-to-4.5% band, and stopped. Instead it is making new highs. Something structural, not merely cyclical, is holding profit share above its historic range.

The reassessment has to be stated plainly because it changes the argument. The old bear leg was that the profit cycle had done most of its work and was about to give some back, which, combined with a rich multiple, made the forward return case weak. Half of that leg is now wrong. The profit cycle is not obviously about to give back; it is extending. But the conclusion the leg supported, that forward returns from here are unattractive, survives on a different and arguably firmer basis, developed in 6.4: a profit share at a fresh all-time high is a demanding base from which to grow,

and it is being capitalised at an elevated multiple. The old reason for caution was "profit will revert." The corrected reason is "profit is already at a record and priced as if the record is permanent." The caution is the same; its foundation moved.

6.2 What actually lifted profit, decomposed honestly


A widespread shortcut explains the profit surge with a level statistic doing a change job: it takes banking-and-financials' 1.84-percentage-point share of the FY25 profit level and uses it to explain the FY20-to-FY25 change in profit. A share of a level cannot explain a change unless you also know what that share was at the start, and the trough-year financial share is rarely computed. That is an attribution error, and it has to be replaced with a proper decomposition of the change.

Here is the honest structure of what moved. The levels are anchored to the canonical Nifty consolidated EPS series; the attribution weights across the legs are indicative and derived, and are labelled as such, because no single primary source partitions the profit increase cleanly into these buckets. The chart shows the shape; the discipline is that the shape, not the exact rupee of each bar, is what the argument rests on.

EXHIBIT 11

What actually lifted profits: revenue and margins did most; the tax cut was one-off

Indicative FY20→FY26 Nifty EPS attribution. Revenue and operating leverage dominate; the 2019 tax cut was a one-time step.



Revenue growth and margin expansion did most of the work; the 2019 tax cut was a one- time step, not an engine.

On an indicative FY20-to-FY26 Nifty-50 EPS attribution, from a base near ₹450 to roughly ₹1,160: revenue growth contributes the largest single leg, operating-margin expansion the second, credit-cost normalisation a meaningful third, the corporate-tax cut a modest one-off step, and sector mix a small residual.

Source: NSE/Motilal Oswal EPS series [GREEN levels]; attribution weights INDICATIVE/derived [AMBER] – labelled per §12. Kakde Research.

FY20→FY26 EPS LEG (INDICATIVE)	CONTRIBUTION	CHARACTER	CLASS
FY20 base	~₹450	starting level	GREEN (level)
Revenue growth	+₹330	recurring, cyclical	AMBER (derived)
Margin expansion	+₹150	partly structural, partly cyclical	AMBER (derived)
Tax cut (2019)	+₹55	one-off step, does not repeat	AMBER (derived)
Credit-cost normalisation	+₹120	one-off normalisation off a bad-loan trough	AMBER (derived)
Other / sector mix	+₹55	residual	AMBER (derived)
FY26	~₹1,160	ending level	GREEN (level)

Levels: NSE/Motilal Oswal consolidated EPS series (GREEN). Attribution weights: indicative/derived (AMBER), constructed per the audit's §12 discipline. The bars are directional, not a sourced partition.

Read the decomposition and three conclusions follow that the level-statistic version obscured.

The first is that the tax cut was a step, not an engine. India cut its headline corporate tax rate in September 2019, and that reset the after-tax profit base upward in a single move. It is real, it lifted the level, and it does not repeat. A tax cut raises the plateau once; it does not compound. Any forward earnings model that implicitly extrapolates the tax cut's contribution is double-counting a one-off. Isolating it, as the decomposition does, is exactly what the level-statistic shortcut cannot do when a single financials-share figure stands in for the whole change.

The second is that credit-cost normalisation was also, in large part, a one-off, and it is the piece most relevant to whether the profit high is durable. Indian banks entered the period working off a severe bad-loan cycle. As provisions fell from crisis levels back toward normal, reported bank profit rose sharply, and because financials are a large weight in the index, that lifted aggregate profit meaningfully. But normalisation is a return to a baseline, not a permanent tailwind. Once credit costs reach a healthy floor, they cannot fall further, and the leg stops contributing. Part of the FY20-to-FY26 profit surge is therefore the reversal of a depressed starting point, precisely the "normalisation" the exhausted-cycle reading emphasises. The correction is not that normalisation was unimportant; it is that normalisation is one leg among several, not the whole story, and the report must neither ignore it nor let it swallow the recurring legs.

The third, and the most important for the forward case, is that the two largest legs, revenue growth and margin expansion, are the ones that determine whether the profit high holds. Revenue growth is recurring but cyclical; it slowed below nominal GDP in FY25 (Section 5) before recovering in FY26. Margin expansion is partly structural (operating leverage, formalisation, mix toward higher-margin businesses) and partly cyclical (input-cost cycles, capacity utilisation). If the durable part of the profit high is revenue plus the structural slice of margins, and the one-off part is the tax step plus the credit-cost normalisation, then the sustainable profit share is somewhere below the reported 5.2%, because the two one-off legs are baked into the current level and will not push it higher from here. That is the quantitative content of "structurally higher but not infinitely extendable." The profit level is real; a portion of it is a plateau raised by steps that have already happened.

This decomposition is what the exhausted-cycle reading needs and lacks, and its absence is why that bear leg was mis-specified. Once you separate the recurring legs from the one-off legs, the reassessment becomes precise: the recurring legs justify a permanently higher profit share than the pre-2019 range, which weakens the "it all reverts" bear case; the one-off legs are spent, which weakens the "it keeps rising" bull case. The truth is a higher, flatter plateau, not a continuing climb and not a reversion.

6.3 One EPS series, not a menu

India earnings narratives carry a second, subtler defect that this section repairs by construction: the cherry-picked endpoint. A "+2.0% FY25 EPS stall" can be quoted as if it were the earnings fact when it is merely the bottom of a 980-to-1,047 estimate range whose top implies +9.1%; an FY26 estimate labelled "~+8%" in fact spans +1.2% to +22.8% across commonly cited sources; and a Nifty-50 print of +7.2% is rarely reconciled against a Nifty-500 print of +2%. A reader can be handed a stall or a boom depending on which line of which table is quoted.

The discipline this report adopts is one canonical series, stated once, with its methodology, and used everywhere. The canonical earnings series is Nifty-50 consolidated EPS, the same series that anchors the decomposition above and the return bridge in Section 1. On that series, FY26 EPS grew roughly 10%, consistent with the profit-versus-GDP gap (+15.6% profit growth at the broad aggregate translates to a smaller per-share figure once you account for the index's dilution and composition). The single most important methodological fact about this series is a break that most published multiples gloss over: NSE switched the Nifty EPS basis from standalone to consolidated in April 2021. Any EPS or P/E leg that spans that break is comparing two different definitions, which is one reason headline trailing P/ E looks optically cheaper than it should (the consolidated switch lowered the denominator-to-price relationship). We flag the break wherever a series crosses it and never headline a growth rate computed across it.

The point is not that one number is right and the others wrong. It is that an earnings figure has to come with its universe (Nifty-50), its basis (consolidated), its period (fiscal year), and its source, and then the same figure has to be used in the text, the exhibits, the summary, and the model. A report that quotes +2% in one paragraph and +8% in another, both true of some line in some table, has not measured earnings; it has offered a menu. The canonical answer for FY26 is roughly +10% on the Nifty-50 consolidated basis, which is neither the stall nor the boom that could be cherry-picked from the range, and it is the figure that reconciles with the profit-to-GDP high rather than contradicting it.

6.4 The valuation counterweight, and why the level being real is not enough

The profit level being real and structurally higher is a genuine improvement to the bull case. It is also, on its own, insufficient, because the question for an allocator is never "is profit high" but "is profit high relative to what I am paying for it." Here the two standard valuation lenses point in opposite directions, and the report shows both rather than choosing the flattering one.

CAPE 36.0 versus a long-run mean near 24; Nifty P/B 2.95 versus a 10-year average of 3.62. On cyclically adjusted earnings, India is roughly 1.5× its own history and near 2021 highs. On price-to-book, it is about 18% *below* its 10-year average. Trailing P/E of 20.5 is about 12% below its 10-year average, partly an artefact of the 2021 standalone-to-consolidated earnings reset.

INDIA VALUATION LENS	VALUE	VERSUS OWN HISTORY	SIGNAL	CLASS
CAPE (Nifty-500, IIMA)	36.0	~1.5× long-run mean (~24)	rich	GREEN
Trailing P/E 20.5-23.9		~12% below 10y avg	cheap-ish (reset artefact)	GREEN
Price-to-book (Nifty)	2.95	~18% below 10y avg (3.62)	cheap-ish	GREEN
Dividend yield	1.27%	low	rich	GREEN

The divergence is not noise, and resolving it is where the forward case is decided. CAPE compares today's price to a ten-year average of real earnings, so it is high precisely because earnings have surged: the denominator is smoothed over a decade that includes the depressed FY20 trough, while the price reflects the current 5.2%-of-GDP profit peak. A high CAPE is, in part, the mathematical signature of a profit share at an all-time high being capitalised at a full multiple. P/B looks cheaper because book value grew with retained profits and because the 2021 earnings-basis reset distorted the P/E-linked measures. Both are true. The reconciliation for an allocator is this: the cheap-looking lenses (P/B, trailing P/E) are cheap partly because profit and book have risen, and the rich lens (CAPE) is rich partly for the same reason, which means the divergence is largely a statement about *earnings being high*, not a genuine disagreement about value. When earnings are at a cyclical-plus-structural peak, low P/B and high CAPE describe the same fact from two ends.

That is why the profit high, however real, produces poor forward asymmetry rather than a green light. Work through the arithmetic of the forward return. Earnings are at a record share of GDP, lifted by two one-off legs that will not repeat and by recurring legs that grow with a cyclical economy. The multiple on cyclically adjusted earnings is roughly 1.5× its own mean. For the forward return to be attractive, you need either earnings to keep taking share of GDP from an already-record base, which the decomposition says is hard because the one-off legs are spent, or the multiple to expand from a rich level, which the twenty-year bridge says has added nothing across a full cycle and is more likely to compress from here than to extend. The asymmetry is unattractive not because profit is fake but because profit is real, high, and fully priced simultaneously. The good news is in the level and in the price at the same time. What is left to pay the forward investor is the recurring growth of earnings minus whatever the multiple gives back, minus the currency, which is exactly the thin, earnings-driven, re-rating-free return the bridge in Section 1 already described.

The reassessment, complete, reads like this. The exhausted-cycle reading was wrong that the profit cycle was about to revert; it has extended and made new highs, and the level is structurally higher than the old range for reasons the decomposition can name. But that reading's conclusion, that forward returns are unattractive, is correct on the corrected reasoning: a record profit share capitalised at an elevated CAPE offers the marginal dollar a poor entry, because both the earnings and the valuation are pricing the good news at once. A higher plateau is not the same as a higher return.

6.5 Who captures the profit

There is one more question the profit-to-GDP high does not answer, and it is the hinge into Section

1. Whose profit is it? The 5.2% figure is *listed* corporate profit as a share of GDP, the Nifty-500 aggregate. It is not all of Indian corporate profit, and it is emphatically not the slice of profit a foreign portfolio owns. Total Indian corporate profit, listed and unlisted, is larger than the listed figure; the listed figure is larger than the free-float figure; and the free-float figure is larger than the foreign-ownable one. The profit high is measured at the widest of those rings that this report cares about, listed India, and it thins at every ring inward.

The dollar sizing makes the point concrete and sets up the next section's opening exhibit. Listed corporate profit at 5.2% of a \$3.9 trillion economy is roughly \$0.2 trillion a year. That is the "Corporate India" object in the Four Indias picture: about \$200 billion of annual listed profit flow, against a \$3.9 trillion economy, a \$5.35 trillion listed market, and a foreign-held stock of under \$0.9 trillion. Whether a fresh all-time-high profit share is good news for the allocator depends entirely on how much of that \$200 billion flow, and the market that capitalises it, is actually ownable. That is a stocks-and-ownership question, not a profit question, and it is where Section 7 goes.

6.6 ALLOCATION TAKEAWAY

The profit-to-GDP high is the strongest genuinely-new fact in favour of Indian equities, and it changes the bear case honestly: the profit cycle is not reverting, and the level is structurally higher than history because recurring revenue and margin legs, not just one-off tax and credit-cost steps, drove it. An allocator still carrying a "profit about to normalise down" view should retire it.

But three disciplines convert that fact into the correct, and still cautious, allocation. First, do not extrapolate the one-off legs: the 2019 tax cut and the credit-cost normalisation raised the plateau and are spent, so the sustainable profit share is below the reported record even if it stays well above the old range. Second, price the level, do not just celebrate it: a record profit share capitalised at a CAPE near 1.5× its mean offers poor forward asymmetry, because the earnings and the valuation are both discounting the good news, and the twenty-year bridge says the multiple is more likely to subtract than add from here. Third, use one earnings number: FY26 EPS grew roughly 10% on the canonical Nifty-50 consolidated series, neither the cherry-picked stall nor the cherry-picked boom, and every forward model should run off that single figure with its basis break disclosed. The profit is real and high. The return you can earn on it depends on the price you pay and, decisively, on how much of it you are allowed to own, which is the next section's entire subject.

SECTION 7

The India a Shareholder Can Own

The gap between "India" and "the India a foreign shareholder can actually own" is the widest and most misused number in the whole allocation, and most treatments get its ceiling wrong.

\$1.65tn

FOREIGN-INVESTABLE POOL (FIF-ADJUSTED)

THE FINDING, UP FRONT. The gap between "India" and "the India a foreign shareholder can actually own" is the widest and most misused number in the whole allocation, and most treatments get its ceiling wrong. There is not one India of a single size; there are several adjacent objects of very different magnitudes: a \$3.9 trillion economy, about \$200 billion of annual listed profit, a \$5.35 trillion listed market, roughly \$2.67 trillion of free float, and under \$0.9 trillion actually held by foreigners. Those are not five stages of a single shrinking pipe. Two of them are annual flows and three are stocks of market value, and arranging flows and stocks in one descending waterfall is a category error this report refuses. More consequentially, the roughly \$0.85 trillion foreigners *hold* is routinely used as the foreign-*ownable* ceiling. Holdings are not the ceiling. The investable pool is larger than current holdings, which means there is real remaining room for foreign capital to own more India, not less. This section builds the ownership ladder properly, grounds it in the real shareholding register, and then redesigns the domestic-revenue measure that most coverage botches, so that "how Indian is this profit" stops conflating where a company sells with where it produces and books its margin.


7.1 The Four Indias are adjacent objects, not a waterfall

The most persistent framing error in the prior work was the funnel: a single narrowing pipe from economy to foreign-ownable, with each stage smaller than the last, and language ("narrows," "contracts," "funnels") that treated the descent as monotonic. It is not monotonic, and it cannot be, because the objects are not commensurable.

EXHIBIT 12

The Four Indias are different objects — flows and stocks, not one shrinking waterfall

Adjacent magnitudes, ~2025. A \$3.9tn economy; a \$5.35tn listed market; ~\$0.88tn actually held by foreigners.



Economic India ~\$3.9tn (GDP, annual flow). Corporate India ~\$0.2tn (listed profit, annual flow). Listed India ~\$5.35tn (market cap, a stock). Free-float India ~\$2.67tn (a stock). Foreign-held India ~\$0.85-0.93tn (a stock). The listed market is larger than the economy; the profit flow is a fraction of both.

Source: MoSPI, NSE, NSDL, Motilal Oswal [GREEN; float derived AMBER]. Foreign-held ≠ foreign-ownable — see investability ladder. Kakde Research.

Look at what the funnel framing has to suppress to survive. The listed market, at \$5.35 trillion, is *larger* than the \$3.9 trillion economy. A market capitalisation bigger than GDP is unremarkable, it reflects the market pricing future profits and is common in markets with high multiples, but it is fatal to a "narrowing pipe" story, because the second object is bigger than the first. Funnel graphics survive by quietly deleting the \$5.35 trillion listed panel, precisely because it breaks the monotonic descent. Deleting the inconvenient object to preserve the shape of the argument is exactly backwards.

The correct framing is that these are adjacent analytical objects, each measured in its own unit, placed side by side so their magnitudes can be compared without pretending one flows into the next. The chart makes the distinction structural: the two annual flows (the \$3.9 trillion economy and the \$0.2 trillion listed-profit flow) sit on one side, the three stocks of market value (the \$5.35 trillion listed market, the \$2.67 trillion free float, the \$0.88 trillion foreign-held) on the other, with a divider between flows and stocks that the eye cannot cross by accident.

Why this matters for the allocation, beyond tidiness, is that the funnel framing invites a specific mistake: reading a ratio computed across two non-commensurable objects as if it meant something. Dividing foreign-held stock (\$0.85 trillion) by GDP flow (\$3.9 trillion), or listed profit flow (\$0.2 trillion) by listed market cap (\$5.35 trillion, which is roughly a 3.7% earnings yield and *is* meaningful), or foreign-held by free float, all produce numbers, but only some of those numbers are economically interpretable, and the funnel encourages quoting all of them with equal confidence. The trap is real: "a fifth" gets used to mean both $\$0.85\text{tn}/\5.35tn (15.9%, foreign share of the listed market) and $\$0.85\text{tn}/\2.67tn (31.8%, foreign share of free float) – the same phrase for two different ratios. Adjacent objects, each labelled with its unit, make that impossible. You cannot casually divide a flow by a stock when the chart has just told you which is which.

The Four Indias, correctly drawn, deliver one clean allocation message before any ladder is built: the object a foreign shareholder actually owns, the roughly \$0.88 trillion foreign-held stock, is small relative to the \$5.35 trillion listed market and tiny relative to the \$3.9 trillion economy, but it is *not* the same object as the economy, the profit, or even the free float, and the relationship between them is not subtraction down a pipe. It is five different measurements of five different things that happen to all be called "India."


7.2 The investability ladder: foreign-held is not the ceiling

Now build the one descent that *is* legitimate, because it stays within a single unit (stock of market value) and each rung is genuinely a subset of the rung above. This is the ladder that replaces custody data as a ceiling.

EXHIBIT 13

Foreign-held is not the ceiling: there is real remaining room to own India

The investability ladder. ~\$0.88tn is held today; the FIF-adjusted investable pool is larger — room exists (illustrative FIF).



Total market cap \$5.35tn → free float (ex-promoter) \$2.67tn → foreign-investable (FIF-adjusted) ~\$1.65tn → actually foreign-held ~\$0.88tn → remaining foreign room ~\$0.77tn. Foreign investors hold under \$0.9tn today; the foreign-investable pool is roughly \$1.65tn; the difference is room, not a wall.

Source: NSE, NSDL [GREEN holdings]; FIF-adjusted cap illustrative pending MSCI IMI pull [AMBER]. Held ≠ ownable — see the method note. Kakde Research.

INVESTABILITY RUNG	VALUE	WHAT IT REMOVES	CLASS
Total market cap \$5.35tn		nothing	GREEN
Free float (ex-promoter) \$2.67tn		~50% promoter holding	AMBER (derived)
Foreign-investable (FIF-adjusted) ~\$1.65tn		float foreigners cannot or do not access	AMBER (illustrative)
Actually foreign-held ~\$0.88tn		the part not yet bought	GREEN (NSDL/NSE)
Remaining foreign room ~\$0.77tn		—	DERIVED

Free float derived from ~50% aggregate promoter holding (NSE). FIF-adjusted figure is illustrative pending a formal MSCI IMI foreign-inclusion-factor pull (AMBER). Holdings are not the ownable maximum.

The correction is the whole point. The common move is to take the roughly \$0.85 trillion that foreigners *hold in custody* (NSDL data) and use it as the foreign-*ownable* float, the maximum foreign capital could own. That is a definitional error with a real allocation consequence. What foreigners hold today is a fact about the past (how much they have bought); what foreigners *could* own is a fact about the structure (how much is investable to them). The two differ by the remaining room, and the remaining room is exactly what an allocator deciding whether to add India needs to know.

Trace the legitimate rungs. Total market cap is \$5.35 trillion. Remove promoter holdings, roughly half the market on the aggregate register, and free float is about \$2.67 trillion. Not all of that free float is accessible to foreign investors: some is held in instruments or structures outside foreign reach, some sits in names below foreign-ownership headroom limits or outside index-inclusion criteria. Applying a foreign-inclusion factor of the kind MSCI uses to build its investable indices brings the foreign-investable pool to roughly \$1.65 trillion. That figure is illustrative here, flagged AMBER, because a precise number requires a formal MSCI IMI foreign-inclusion-factor pull that this report has not completed; the *structure* of the rung is sound even if the exact value awaits that pull. Against a foreign-investable pool near \$1.65 trillion, foreigners actually hold about \$0.88 trillion. The difference, roughly \$0.77 trillion, is remaining foreign room: investable India that foreign capital has not yet bought.

Two internally-inconsistent versions of the foreign figure circulate, and reconciling them is part of the fix. One is \$0.85 trillion (a custody snapshot); the other computes 17.3% of the \$5.35 trillion listed market and gets \$0.93 trillion, and the two are rarely reconciled. The honest statement is a range: foreign holdings are approximately \$0.85 to \$0.93 trillion depending on whether you take the custody snapshot or the ownership-percentage computation, and the report uses the range rather than pretending to a false single-point precision. Either way, the number is *holdings*, and holdings sit well below the \$1.65 trillion investable pool.

The allocation implication runs opposite to the way the funnel was usually read. The funnel, ending at "\$0.85 trillion foreign-ownable," was typically read as a constraint: India is nearly full, foreigners already own the ownable slice, there is little room. The corrected ladder says the reverse. Foreigners hold roughly half the foreign-investable pool. There is on the order of \$0.77 trillion of investable headroom, which is one reason the flow story matters (Section 5 and the section that follows): foreign positioning in India is not maxed out, it round-tripped down from the 2024 peak, and the room to re-add is real. Whether foreign capital *should* re-add is the valuation-and-return question the rest of the report answers, and that answer is cautious. But "there is no room left" was never true; it was an artefact of using holdings as the ceiling. The ceiling is higher than the holdings, and the gap is the opportunity or the overhang depending on which way flows turn.

7.3 Who actually holds the shares

The ladder is built on an ownership structure, and that structure is now grounded in the real register rather than in illustrative placeholders. The aggregate ownership of the listed market divides four ways.

Promoters ~50.0%, foreign portfolio investors ~17.3%, domestic institutions ~16.1%, retail ~18.5% of the listed market (NSE). Promoters, the single largest block, are the reason free float is only about half of market cap.

The single most consequential number here is the promoter share. At roughly 50%, promoter holding is why the first rung of the ladder cuts the ownable market in half before foreign access is even considered. India is a market of controlled companies: founding families, conglomerate groups, and the state hold half the listed value and do not trade it. That is not a defect, controlled ownership has governance costs and benefits both, but it is a structural fact that halves the free float and therefore halves the base from which foreign ownership is measured.

The real company-level register shows how much that structure varies, and the variation is itself an allocation signal. Drawing on the shareholding data for thirty large caps (exchange filings, June 2026, flagged AMBER as filing-sourced):

- **The genuinely widely-held core is the private banks.** HDFC Bank and ICICI Bank have *zero* promoter holding; their shares are owned by institutions and the public outright. HDFC Bank is roughly 41.8% foreign portfolio investors and 41.7% domestic institutions; ICICI Bank is about 33.8% foreign and 42.3% domestic. These are the most freely-ownable large caps in India, which is a large part of why financials dominate foreign portfolios: they are where the free float actually is. Larsen & Toubro is similar, zero promoter, 19.1% foreign, 43.0% domestic.
- **The promoter-locked names sit at the other extreme.** TCS is 71.8% promoter (Tata group), Avenue Supermarts is 74.5% promoter, Hindustan Unilever 61.9%, HCL Technologies 60.9%. In these names the free float is a quarter to a third of the company, and however attractive the business, the ownable slice is small.
- **The state-controlled names are a distinct category.** State Bank of India is 55.5% government, Hindustan Aeronautics 71.6% government, Bharat Electronics 51.1%. Here the "promoter" is the state, which will not sell control, and the free float is what remains.
- **Reliance sits near the aggregate.** Promoter 50.5%, foreign 17.2%, domestic 21.1%, close to the market-wide split, which is fitting for the index's largest single weight.

The takeaway from the real register is that "buy India" resolves, at the company level, into a market where the ownable free float is concentrated in a specific and identifiable set of names, the zero-promoter private financials and the lower-promoter large caps, while a large part of the index by value is locked up in promoter and government hands that a portfolio investor cannot access at any price. The free-float rung of the ladder is not evenly distributed across the market; it is bunched in financials and a handful of widely-held names, which is one more reason the foreign portfolio's actual composition tilts so heavily toward banks. Ownership structure, not just sector view, shapes what a foreign allocator ends up holding.


7.4 The DRI redesign: where revenue is earned is not where profit is made

There is a genuinely important question here – how much of "Indian" listed profit is actually earned in India versus abroad – and the standard measure answers it with a conflation that produces nonsense. It computes a company's "India share" as 100% minus its export share, which counts the output, labour, capital, and tax of a domestic facility as "not India" the moment its product is sold abroad. On that logic Reliance's Jamnagar refinery, one of the largest and most Indian industrial assets in the country, employing Indian workers, built with Indian capital, paying Indian tax, becomes "not India" to the extent it exports refined products. That is wrong, and it has to be rebuilt from the definitions up.

EXHIBIT 14

Two different 'Indias' inside the index: where revenue is earned, and who controls the shares

Domestic revenue share vs promoter holding, large caps (FY26). IT earns abroad; banks are domestic; promoter control varies widely.



Source: Company filings / exchange shareholding, Jun-2026 [AMBER]. Kakde Research.

The redesign separates three geographies that the standard measure collapses into one:

- **End-market revenue geography:** where the customer is, where the sale is booked. This is what "exports" measures.
- **Production, cost, and asset geography:** where the factory, the workforce, the capital expenditure, and the balance-sheet assets sit.
- **Profit geography:** where the margin is actually earned and taxed.

For a services exporter these three coincide in a specific way: Infosys and TCS *produce* in India (Indian engineers, Indian campuses, Indian cost base) and *sell* abroad (roughly 94-97% of revenue is from foreign clients). Their end-market revenue geography is overwhelmingly foreign; their production and cost geography is overwhelmingly Indian; and their profit is earned on an Indian cost base serving foreign demand. Calling them "not Indian" because they export is exactly the error, they are among the most Indian-*production* companies in the index, monetising Indian labour for foreign customers.

For a domestic bank the three coincide the other way: HDFC Bank earns essentially 100% of its revenue in India, from Indian borrowers and depositors, produced and booked entirely in India. It is domestic on every geography.

For Reliance the discipline matters most. Its end-market revenue is roughly split between domestic sales and exports, so on the old measure it looked "half not-India." But its production, assets, workforce, capex, and tax are overwhelmingly Indian; the exported barrels are refined at Jamnagar by Indian workers and taxed in India. Its *profit* geography is far more domestic than its *revenue* geography, because the value added and the margin are captured in India even when the customer is abroad.

Roughly 28% of large-cap revenue is earned outside India (end-market), so domestic revenue intensity is about 72%. But domestic profit intensity, where the margin is booked and taxed, is higher, because export revenue from Indian production is Indian profit. Universe: large-cap sample, free-float-cap weighted; method: aggregate end-market revenue split from segment disclosures.

The chart pairs each company's domestic revenue share against its promoter holding, and the two "Indias inside the index" are visible at a glance. The IT names anchor the low-domestic-revenue end (Infosys near 3% domestic revenue, TCS near 6%), because they sell abroad, even though they produce at home. The banks anchor the high-domestic end (HDFC Bank at 100%). In between sit the genuinely mixed businesses: Dr Reddy's near 16% domestic revenue, Tata Motors near 28% (its profit swings on Jaguar Land Rover, a genuinely foreign business), Sun Pharma near 33%, Hindalco near 40%, Reliance near 49%, and Bharti Airtel near 77%.

Bharti Airtel is worth a specific note, because it carries a correction of its own. A widely repeated figure puts Airtel's India-versus-Africa split at roughly 43/57 – a number that comes from confusing an EBITDA split with a revenue split. Corrected, Airtel is roughly 77% India and 23% Africa on revenue. The direction of the error is instructive: it made an Indian telecom look majority-foreign, the same category of mistake as calling Jamnagar "not India." Getting the definition right moves Airtel from "mostly foreign" to "mostly Indian," which is what it is.

To be explicit about method, because measures of this kind rarely state their universe or weights: the domestic-revenue measure here is built on a large-cap sample, free-float-market-cap weighted, using end-market revenue splits from company segment disclosures, and it is labelled AMBER because those disclosures vary in granularity and the aggregation involves judgment. The headline, roughly 28% of large-cap revenue earned abroad, is a *revenue-geography* statement. The *profit-geography* statement, domestic profit intensity, is higher than the revenue figure, because a material share of that 28% is Indian-produced output sold to foreign customers, whose margin and tax are Indian. This report states both and never lets the revenue figure stand in for the profit figure. That distinction is the entire content of the redesign: an allocator who wants to know "how exposed am I to India" needs the profit geography; an allocator who wants to know "how exposed am I to global demand" needs the revenue geography; and the two are different numbers that the old measure fused into one wrong one.

7.5 ALLOCATION TAKEAWAY

The India a shareholder can own is a specific, buildable object, and getting its size and shape right changes the allocation in three concrete ways.

First, stop reading the ownership picture as a shrinking funnel and stop using foreign holdings as the ceiling. The Four Indias are adjacent objects in two different units; the listed market is larger than the economy, not a stage downstream of it. And foreigners hold roughly \$0.88 trillion against a foreign-investable pool near \$1.65 trillion, so there is on the order of \$0.77 trillion of remaining room. The correct mental model is a ladder with headroom, not a pipe that has run dry. Whether to use that headroom is a valuation call, and Section 6 says be cautious; but the room exists, and positioning is not maxed.

Second, respect the promoter wall when sizing what is actually ownable. Half the listed market is promoter-held and does not trade, and the free float is bunched in a specific set of names, the zero-promoter private banks and the lower-promoter large caps, rather than spread evenly across the index. A foreign allocator's realistic India is disproportionately financials and a handful of widely-held names, because that is where the free float physically is. Sector preference and ownership structure point to the same overweight, and an allocator should know that the market is choosing financials for them through the float.

Third, use profit geography, not revenue geography, to judge India exposure, and never count Indian production as foreign because it is exported. Roughly 28% of large-cap revenue is earned abroad, but domestic profit intensity is higher, because the IT majors and Reliance's export book monetise Indian production and pay Indian tax. The companies that are genuinely foreign in profit terms, Tata Motors through Jaguar Land Rover, Airtel's Africa arm, the pharma export franchises, are identifiable and can be sized deliberately, rather than smeared across the index by a measure that mistakes a customer's address for a factory's.

The India a shareholder can own is smaller than the economy, concentrated in the free float, tilted toward financials by ownership structure, and more domestic in its profit than its revenue geography suggests. It is a real, ownable object with genuine remaining room. What turns that room into a position, or leaves it empty, is the flow of capital into and out of it, the domestic bid that has become the market's dominant buyer and the foreign flow that round-tripped from the 2024 peak. That financial machine, who is buying, who is selling, and whether the domestic floor is a genuine floor or a fair-weather one, is where the report turns next.

SECTION 8

India's Financial Machine

The most important flow in Indian equity is no longer foreign; it is domestic, and it is far larger than the SIP number everyone quotes.

4.6x

FY25 DII BUYING VS FPI SELLING

THE FINDING, UP FRONT. The most important flow in Indian equity is no longer foreign; it is domestic, and it is far larger than the SIP number everyone quotes. In FY25 the domestic institutional bid bought ₹6.07 lakh crore of equity while foreign portfolio investors sold ₹1.31 lakh crore – a 4.6× absorption that held the market up through a genuine foreign exit. Behind that flow sits a savings machine most allocators still size by the ₹31,781-crore monthly SIP: LIC alone holds roughly ₹57 lakh crore of assets, equity mutual funds ₹37.3 lakh crore, EPFO ₹24.8 lakh crore, NPS and APY ₹16 lakh crore. The SIP run-rate is a slice of the tank, not the tank. But two corrections keep this from becoming a bull case. First, the same savings pool that bids for equity also funds the *supply* of it – record IPO, QIP and offer-for-sale issuance is a real sponge that partly offsets the inflow. Second, the popular claim that domestic buying protects the rupee is a mechanism error: domestic absorption supports the equity *price*, not the balance of payments, and a supported exit price can raise the dollars a departing foreigner repatriates. The domestic bid is a proven shock-absorber. It is not yet a proven permanent floor, and it does not defend the currency. This section builds the machine from the banks up, then follows the money to those two conclusions.

The financial sector is where India's equity market is most itself. Financials are the single largest weight in the index, the sector whose profit did the heaviest lifting in the FY20–FY26 earnings recovery, and the conduit through which household savings become corporate capital. To underwrite Indian equity is, disproportionately, to underwrite Indian finance. So the section starts with the lenders, moves to the credit cycle they sit in, and ends on the flow story that has quietly become the market's dominant fact.

8.1 Two banking systems under one flag

India runs two banking systems that share a regulator and almost nothing else. The public-sector banks, led by the State Bank of India, still hold a majority of system deposits and carry the state's policy objectives on their balance sheets. The private banks – the large-cap franchises that dominate the *investable* financial index – hold a minority of deposits but the majority of the sector's market capitalisation, profitability, and foreign-investor interest. The gap between deposit share and market-cap share is the whole story: the listed, foreign-ownable financial claim is skewed heavily toward private franchises that punch far above their balance-sheet weight in the index and far below it in the deposit base.

This split matters for an allocator in a way the aggregate "Indian banks" label conceals. Public-sector banks trade at low multiples for structural reasons that are unlikely to fully reverse: lower return on equity, periodic recapitalisation calls on the exchequer, directed lending, and a governance discount that reflects state control of capital-allocation and management appointments. Private banks trade at premium multiples for the mirror reasons: higher and more stable ROE, cleaner underwriting through the last cycle, and the structural share gain that formalisation hands to organised, compliant lenders – the same share-shift mechanism that Section 3 identified as formalisation's strongest transmission into listed profit. The multi-year through-line of the sector has been market share migrating from public to private, from unorganised credit to formal credit, and from the branch to the digital rail. That migration is the most reliable earnings tailwind in the Indian financial complex, and it is already substantially reflected in the relative multiples of the private leaders.

The concentration of the index compounds this. The Nifty top-ten names account for **52.9%** of the index (GREEN), and financials are the plurality of that top slice. An allocator buying the Indian index is buying, to a first approximation, a handful of large private banks, one dominant public bank, and the largest private financial conglomerates – plus a diversified tail that moves the needle far less. There is no owning "India" without owning Indian finance, and there is no owning Indian finance at the index level without a heavy tilt to the private-sector premium franchises that already carry the sector's optimism in their price.

8.2 The credit cycle: a clean balance sheet, a maturing upswing

The reason financials carried the profit recovery is a credit cycle that turned decisively in the lenders' favour. India came through the 2016–2020 asset-quality crisis – the corporate bad-loan reckoning that followed the previous capex boom – and emerged, by the mid-2020s, with the cleanest bank balance sheets in a generation. Gross non-performing assets fell to multi-decade lows, provision coverage rose, and credit costs – the charge lenders take for bad loans – collapsed from crisis levels to something near a cyclical floor. Those specific figures are not in this report's canonical dataset and are treated here as illustrative context (INTERP), but the direction is not in dispute and it is corroborated by the profit data: a large part of the FY20–FY26 bank earnings surge was not new lending, it was the *disappearance of the credit-cost drag* that had suppressed reported profit through the bad-loan years.

This is a crucial distinction for anyone extrapolating financial-sector earnings, and it is the same level-versus-delta discipline the report applies to corporate profit in Section 6. A bank whose earnings jumped because credit costs fell from crisis levels to a cyclical low has captured a one-time normalisation, not a permanent new growth rate. Credit costs cannot fall below zero; the tailwind of a cleaning balance sheet is, by construction, exhausted once the balance sheet is clean. From here, financial-sector earnings have to be driven by loan growth, net interest margins, and fee income against a *normal* – not a falling – credit-cost line, and the next credit cycle will eventually turn the other way, as every credit cycle does. The pristine current asset quality is a reason the sector is not fragile. It is not a reason to project the recent earnings growth rate forward, because a meaningful part of that growth rate was the credit-cost line healing, and it has largely healed.

The lending upswing itself has been real but is maturing. System credit growth ran in the high teens through the recent expansion before cooling as the cycle aged and the regulator leaned against unsecured retail exuberance (illustrative; the RBI's macroprudential tightening of risk weights on unsecured consumer credit is widely reported). The cooling connects directly to Section 5's household finding: the same household deleveraging that repaired the net-financial-savings ratio is, from the bank's side of the ledger, softer demand for new consumer credit. A household borrowing less is a bank lending less to households. So the retail-credit engine that powered private-bank growth is running into a household sector that is, at the margin, choosing to de-lever. That is not a crisis; it is the ordinary maturing of a credit upswing, and it caps the pace at which even the best-run private lenders can compound from here.


8.3 The domestic bid, sized properly

Here the section reaches the fact that has changed the Indian market's microstructure, and the first job is to size it correctly, because the popular framing – "SIPs are holding up the market" – undersells the phenomenon by an order of magnitude.

EXHIBIT 15

The domestic bid is far bigger than the SIP headline

Indian institutional pools vs the annual SIP run-rate. LIC alone (~₹57L cr) dwarfs one year of SIPs (~₹3.8L cr).



LIC ~₹57 lakh crore · equity mutual funds ₹37.3 lakh crore · EPFO ₹24.8 lakh crore · NPS+APY ₹16 lakh crore. The domestic institutional savings pool dwarfs the ~₹31,781 crore monthly SIP run-rate. One year of SIPs (~₹3.8 lakh crore) is a slice of a tank measured in tens of lakh crore.

Source: AMFI, LIC filing, PIB/EPFO, PFRDA [GREEN/AMBER], 2025-26. Kakde Research.

The SIP figure is the most-quoted domestic number and the least representative. It is a *flow*, and a marginal one, drawn off the top of a vast and growing *stock* of institutionalised household savings. The Life Insurance Corporation alone, at roughly ₹57 lakh crore of assets, is a bigger pool than the entire equity-mutual-fund industry, and a fraction of its book is already in equities that it manages with a long horizon and a countercyclical hand. EPFO, the provident-fund body, holds ₹24.8 lakh crore and routes a policy-set share – on the order of a tenth – into equity ETFs. NPS and APY add ₹16 lakh crore of long-horizon retirement money with a rising equity allocation. These are not hot-money pools. They are contractual, sticky, and, in the pension and insurance cases, structurally biased toward buying more equity over time as mandates and allocations rise. The mutual-fund SIP is the visible, retail-branded tip of an institutional iceberg that is far larger and, in aggregate, far stickier than the monthly flow number implies.

What that pool did in FY25 is the demonstration:

+₹6.07 lakh crore (DII) versus –₹1.31 lakh crore (FPI) – a 4.6× absorption. Domestic institutions bought more than four-and-a-half times what foreign portfolio investors sold, and the market did not break.

This is the single most important microstructure fact in the contemporary Indian market, and it is a genuine change from the pre-2020 regime, when foreign flows set the marginal price and a foreign exit meant a market decline. In FY25 the foreigners left in size and the domestics simply absorbed them. The ownership arithmetic now reflects it: domestic institutions hold **16.1%** of the listed market against foreign portfolio investors' **17.3%** (GREEN, NSE), a

gap that has narrowed to near-parity from a chasm a decade ago, with retail direct holding at **18.5%** and promoters at **50.0%**. The domestic institution has become a co-equal marginal buyer alongside the foreigner, and on the FY25 evidence, the more committed one.

The right framing is worth stating in full, because both halves are load-bearing: the domestic bid is a **proven shock-absorber, not yet a proven permanent floor**. It absorbed a real foreign exit in a real year – that is proven, and it retires the old "foreigners set the price" model. But it has not been tested through a domestic recession, a domestic equity drawdown deep enough to frighten the SIP investor, or a multi-year period of poor returns that would test whether the retail saver keeps his standing instruction in place when his statement is red. The flow is drawn from the household balance sheet that Section 5 showed to be deleveraging and saving cautiously, and a large part of the equity-flow growth is a *composition shift* – households moving existing savings from deposits and gold into funds – rather than a rising savings tide. A one-time reallocation is not a perpetual inflow. The honest statement is that the domestic bid has passed one hard test and has not yet faced its hardest one, and an allocator who treats it as a guaranteed floor is extrapolating from a single favourable observation.

8.4 The floor and the ceiling come from the same pool

The correction that keeps the domestic-bid story honest is on the supply side, and it is the fact the bull case routinely omits. Equity is not a fixed quantity that inflows chase; it is manufactured to meet demand, and India has been manufacturing it at a record pace.

₹3.05 lakh crore (FY26) and ₹3.71 lakh crore (FY25, a record) – equity supply through IPOs, qualified institutional placements, and offers-for-sale. A real sponge that partly offsets the domestic inflow.

The mechanism is straightforward and it is the counterweight to the flow narrative. When domestic savings bid aggressively for equity, they lift prices; higher prices are precisely the signal that pulls promoters, private-equity sponsors, and companies to *sell* equity – to float new listings, place shares with institutions, and reduce promoter stakes into the strength. Roughly ₹3 to ₹3.7 lakh crore of that selling has met the domestic bid in each of the last two years. Set that against the annual SIP run-rate of about ₹3.8 lakh crore and the picture sharpens: a very large share of the marginal retail inflow has been absorbed not by pushing up the prices of existing shares but by funding the creation of new ones. The savings pool that provides the floor under prices simultaneously calls forth the supply that caps them. The floor and the ceiling come from the same savings pool.

This does not negate the domestic bid; it disciplines it. The flow is real and the shock-absorption is real, but the reflexive conclusion – that relentless SIP inflows must mechanically inflate the index – ignores that the market meets excess demand by issuing more claim. In an efficient primary market, that is healthy: it channels household savings into productive corporate capital-raising rather than into pure price inflation of a fixed share count. For the allocator it means the domestic bid is better understood as a lower cost of equity for Indian issuers – cheap, abundant, patient domestic capital – than as a guaranteed one-way lift to secondary-market prices. Cheap capital for issuers is a genuine structural positive for the Indian corporate sector's ability to fund growth. It is a much weaker argument for the proposition that the index is destined to rise because money keeps coming in, because the money keeps calling forth the paper to absorb it.

8.5 The sign error: domestic buying supports price, not the balance of payments

The final correction is a mechanism error this report explicitly retires, because it matters for how an allocator thinks about the currency in the next section. It is often argued that domestic buying reduces pressure on the rupee, on the reasoning that rupee-based investors "aren't demanding dollars" when they buy equity. The logic inverts the actual balance-of-payments mechanics.

Trace the transaction. A foreign portfolio investor decides to exit ₹1 of Indian equity. He must sell that equity for rupees, then convert the rupees to dollars to repatriate – and *that* dollar demand is what pressures the rupee. Now introduce the domestic buyer. The domestic institution buys the foreigner's ₹1 of equity. Has the foreigner's dollar demand changed? No. The foreigner still sells, still converts, still repatriates. The domestic bid did not reduce the repatriation; it *enabled* it at a better price, by giving the foreigner a willing buyer so he did not have to crash the market to exit. If anything, the effect on repatriated dollars runs the opposite way to that claim: a supported exit price means the foreigner sells the same number of shares for *more* rupees, converts *more* rupees, and repatriates *more* dollars than he would have into a falling, illiquid market. The domestic bid can, at the margin, *increase* the dollar outflow associated with a given foreign exit, because it improves the price at which that exit clears.

The correct decoupling is this: **domestic absorption supports the equity price; it does not support the balance of payments.** The two are separate channels, and the popular argument fuses them. The equity-price buffer is real and valuable – it is the shock-absorption of §8.3, and it means a foreign exit no longer automatically means a market crash. But it provides no defence for the currency, and it may modestly worsen the FX effect of an exit by improving the exit price. The rupee is defended by the current account, the reserve buffer (~\$693bn, Section 4), and the RBI's intervention – not by domestic equity buyers. Anyone modelling the domestic bid as a currency stabiliser is double-counting a benefit that lives entirely on the equity-price side of the ledger and does not cross to the FX side. This matters directly for Section 9 and Section 10, because the return an allocator earns is the equity return *net of* the currency, and the domestic bid touches only the first term.

8.6 ALLOCATION TAKEAWAY

India's financial machine is the core of the equity claim and the clearest place to see both the strength and the limits of the domestic story. The banks are two systems: an investable private complex that carries the sector's premium and its optimism, and a public complex that carries the state's objectives and a structural discount – and the index buys you a heavy dose of the premium private franchises, whose share-gain and clean-balance-sheet story is already largely in the multiple. The credit cycle that powered financial earnings has substantially normalised; a large part of the recent bank-profit surge was credit-cost healing that cannot repeat, so the forward earnings rate must come from loan growth and margins against a household sector that is choosing to de-lever.

The flow story is the genuine regime change, and it deserves to be sized properly and bounded honestly. The domestic bid is far larger than the SIP headline – LIC, EPFO, NPS and the fund industry together dwarf the monthly flow – and it absorbed a 4.6× foreign exit in FY25 without breaking the market. That is a proven shock-absorber and it retires the old foreign-flow-sets-the-price model. But three disciplines travel with it. First, it is not a proven permanent floor: it has not been tested through a domestic drawdown, and its fuel is a cautious, deleveraging household balance sheet. Second, the same savings pool funds record equity *supply*, so the flow is better read as cheap capital for issuers than as a one-way lift to prices – the floor and the ceiling come from the same pool. Third, the domestic bid supports the equity price and not the balance of payments; it is no defence of the rupee, and may even raise the dollars a departing foreigner takes home. For the allocator, the practical consequence is to treat the domestic bid as a lower cost of equity and a lower drawdown risk from foreign exits – real and valuable – but not as a reason to pay any price, and not as a currency hedge. The price you pay still has to clear a valuation test the flow cannot excuse, and that test is next.

SECTION 9

The Price of the Story

India is expensive on the measures that predict long-run returns and only middling-expensive on the measures that do not, and an honest allocator has to hold both facts at once rather than pick the flattering one.

CAPE 36

VS A LONG-RUN MEAN NEAR 24

THE FINDING, UP FRONT. India is expensive on the measures that predict long-run returns and only middling-expensive on the measures that do not, and an honest allocator has to hold both facts at once rather than pick the flattering one. The cyclically adjusted price-earnings ratio (CAPE) is 36 – about one-and-a-half times its long-run mean near 24 – which is a real mean-reversion headwind. Yet the price-to-book ratio at 2.95 sits roughly 18% *below* its ten-year average, and trailing P/E is below its ten-year average too, because a 2021 change in how Nifty earnings are compiled reset the headline series. The divergence is not a paradox; it is a data artefact plus a genuine growth premium, and this section shows both without cherry-picking. India's premium over emerging markets is a price-to-book and growth-persistence premium – the market pays up for a belief that Indian companies will *sustain* high returns on equity for longer – and it is emphatically *not* a spot return-on-equity premium, because direct ROE data is a gap in this report's evidence and the derived figure (~14.4%) is statistically indistinguishable from the EM average (~13.8%). Building the price back into an implied-expectations model exposes the core tension: today's price does not merely require India to sustain its current return on equity; it requires that return to *rise*. The observed ROE does not yet support what the price assumes. And the escape hatch of active management is shut – roughly three-quarters of active large-cap funds have underperformed their benchmark over both one and ten years. This section prices the story honestly and hands the residual risk to the currency.

Valuation is where a research report either keeps its discipline or spends it. The temptation is to reach for whichever multiple supports the view already held – the bull cites P/B, the bear cites CAPE – and to present a single tidy number as if valuation were a solved quantity. It is not. This section presents the divergence first, explains it, then does the harder work of turning price into an expectation and testing that expectation against what the companies actually earn.


9.1 Two multiples, two verdicts

Start with the divergence, stated plainly, because it is the fact most reports suppress.

EXHIBIT 16

Is India expensive? CAPE says clearly yes; P/B says not really

MSCI/Nifty India valuation vs history. CAPE 36 is $\sim 1.5\times$ its mean; P/B 2.95 is $\sim 18\%$ below average (the 2021 earnings-basis reset).



CAPE 36 ($\approx 1.5\times$ its mean of ~ 24) says rich. Price-to-book 2.95 ($\approx 18\%$ below its 10-year average of 3.62) says cheapish. Trailing P/E 20.5 sits below its 10-year average of ~ 23.4 . Same market, three multiples, two verdicts.

Source: IIMA India CAPE, craytheon Nifty P/E-P/B, MSCI [GREEN]. Kakde Research.

The two verdicts are both correct on their own terms, and the reconciliation is the 2021 earnings-basis reset. In April 2021 the NSE switched the Nifty earnings series from *standalone* to *consolidated* – from parent-company accounts to group accounts that fold in subsidiaries. Consolidated earnings are higher, so the same price divided by higher earnings produces a *lower* headline P/E. Any price-earnings comparison that spans the 2021 break is therefore comparing a post-2021 consolidated multiple against a pre-2021 standalone history, and the mechanical effect is to make current P/E look cheap relative to its own past – not because the market de-rated, but because the denominator was redefined upward. This is the same data-integrity flag the report raises against the return-bridge decomposition: the standalone-to-consolidated break distorts any earnings-multiple series that crosses it, and it must be flagged, never headlined.

CAPE sidesteps part of this problem, which is why it deserves the most weight of the three. The cyclically adjusted P/E averages a decade of *real, inflation-adjusted* earnings in its denominator, smoothing the cyclical and one-off distortions that a trailing multiple carries – including, to a degree, the reset. At 36 against a long-run mean near 24, India's CAPE is close to its historical highs and about half again its own average. On the historical relationship between starting CAPE and subsequent long-run returns – one of the more robust regularities in equity data across markets – a starting point this elevated is associated with below-average forward real returns. That is not a market-timing signal; CAPE is useless for timing and this report makes no timing claim. It is a statement about the *starting yield* an allocator buys at, and the starting yield is poor.

Price-to-book tells the softer story, and it is not nothing. At 2.95 versus a ten-year average of 3.62, India's book multiple has actually compressed – the market pays less per rupee of net assets than it did on average over the past decade. Book value is far less sensitive to the 2021 earnings reset than P/E, because the reset changed how *earnings* were compiled, not the underlying equity base, so the P/B signal is cleaner than the P/E signal here. The P/B compression is real, and it is the strongest single fact in the "India is not as expensive as you think" case. But P/B has to be read together with the return on that book – a high P/B is justified by a high ROE, and a falling P/B against a stable ROE is a genuine cheapening, while a falling P/B against a *falling* ROE is just the market repricing worse economics. Which of those describes India is the question §9.3 has to answer, and it is where the evidence runs thin.

The honest synthesis: on the smoothed, reset-resistant measure that best predicts long-run returns (CAPE), India is expensive – clearly. On the book measure that resists the reset but needs an ROE to interpret it (P/B), India is middling to slightly cheap versus its own history. On the trailing measure most contaminated by the reset (P/E), India looks cheap, and that cheapness is the least trustworthy of the three. An allocator who quotes only P/E to argue India is reasonably valued is leaning on the one number the data specifically warns against. An allocator who quotes only CAPE to argue India is a bubble is ignoring a real book-value compression. The truth is the divergence, and the divergence tilts expensive once you weight the measures by how much you can trust them.

9.2 What the premium actually is – and is not

India has, for most of the past decade, traded at a premium to the emerging-market complex, and the nature of that premium is routinely mis-stated. The precise version matters because it determines what has to be true for the premium to be justified.

The premium is a **price-to-book and growth-persistence premium**. The market pays a high multiple of book value for Indian equity because it believes Indian companies will generate high returns on equity, and – more importantly – *sustain* them for longer than the average emerging-market company, whose high returns tend to be competed away faster. That belief in persistence, in a long runway of above-cost-of-capital returns, is the economic content of the premium. It is a defensible thesis: India's demographic runway, its formalising domestic market, and the franchise quality of its best private-sector companies are real reasons to expect return persistence. The premium is a bet on the *durability* of Indian corporate returns.

The premium is **not** a spot return-on-equity premium, and this is where careless treatments overreach. Near-parity ROE between India and EM – a figure like 13.9% versus 13.8%, a five-basis-point claim – is sometimes asserted and used to argue India earns more per unit of book, justifying the price. Two problems. First, that ROE figure is derived, not observed: it is backed out of the identity $ROE = (P/B) \div (P/E)$, which is arithmetically circular when the P/E is itself distorted by the 2021 reset – you cannot prove a return premium using a ratio of the two multiples whose relationship is the thing in question. Second, a derived, reset-contaminated number cannot carry the confidence of a verified fact.

The corrected statement is deliberately weaker and more honest. Direct, independently sourced ROE for India versus EM is a **gap** in this report's evidence – the report does not have a clean, primary, matched ROE comparison and will not manufacture one from an identity. The *derived* Indian ROE is approximately **14.4%**, and the EM figure is approximately **13.8%**. Those are close enough that the correct conclusion is *rough parity, not a demonstrated premium* – and certainly not a premium provable to five basis points. India does not visibly out-earn EM on spot ROE in the data available. What India is priced for is not earning more *today*; it is earning it *for longer*. That reframing is not a demotion of the thesis – return persistence is a perfectly good reason to pay up – but it relocates the entire premium

onto an assumption about the *future path* of ROE rather than onto a *current* fact about its level. A premium that rests on the future is a premium that the future can disappoint, and the next subsection tests exactly how much future the price is demanding.

9.3 What the price is really assuming: an implied-expectations view

The disciplined way to interrogate a valuation is to run it backwards: instead of forecasting cash flows to derive a price, take today's price as given and solve for the expectations embedded in it, then ask whether those expectations are reasonable. A single-stage version of this exercise produces a single perpetual-growth figure – around 11.5–12% forever – which this report rejects as economically implausible (no company grows at 11.5% in perpetuity; a perpetual growth rate above the economy's nominal growth rate eventually makes the company larger than the economy). A single-stage model is the wrong tool. The right tool is a multi-stage reverse discounted-cash-flow model that lets growth fade, and that is anchored to the *return on equity* actually required to fund the growth it assumes.

The building blocks, stated so the reader can see every assumption (the derived figures below are illustrative, built to reconcile the canonical multiples and cost of equity; they are not independently observed):

INPUT	VALUE	BASIS
Starting point	CAPE 36 / P/B 2.95	GREEN (canonical)
Near-term earnings growth	~15%, then fading	AMBER (broker consensus, illustrative)
Fade horizon	high growth for ~5 years, decaying toward terminal over the next ~5	illustrative structure
Cost of equity (INR)	~13.3%	AMBER (risk-free ~7% + equity risk premium)
Cost of equity (USD)	~10.5–11%	derived (INR CoE less ~2.5–3% currency drag)
Payout ratio	~30% (retention ~70%)	AMBER
Terminal growth	≤ nominal GDP (~5–6% real-plus-inflation, USD terms lower)	constraint

The mechanics are less important than the single relationship they expose, so state it directly. In a residual-income or justified-multiple framework, a price-to-book of 2.95 against a cost of equity of 13.3% *requires* a return on equity comfortably above the cost of equity, sustained across the high-growth stage, to be justified. Back the required ROE out of the price and it lands **above** the ~14.4% derived observed ROE – on the order of the mid-to-high teens through the near-term stage. The single-stage arithmetic implies roughly the same thing by a cruder route: an 11.8% perpetual growth on a 70% retention rate implies a required ROE near 16.9% (growth = retention × ROE, so ROE = 11.8% ÷ 0.70 ≈ 16.9%). Whether one uses the crude single-stage figure or the multi-stage justified-P/B logic, the arrow points the same way.

That is the core tension, and it is the sharpest internal check in the entire valuation:

The price assumes ROE rises. Observed/derived ROE is ~14.4%. The multiple implies a required ROE in the mid-to-high teens sustained through the growth phase. The gap between the two is the bet an allocator is making, whether or not they know it.

These two numbers – a ~14% observed ROE and a ~17% implied ROE – are usually left unreconciled. This report resolves the contradiction by naming it: today's Indian equity price is not priced for the ROE India *has*; it is priced for an ROE India would have to *earn its way up to*. The premium of \$9.2 is precisely this – the market is paying not for the current return but for a higher, persistent future return. The reverse-DCF makes the implicit assumption explicit and quantifies it. And it clarifies what would validate or break the price. The price is validated if margins hold at their elevated level and revenue accelerates to carry the earnings baton (the Section 5 and Section 6 tension), lifting ROE from ~14% toward the mid-teens and holding it there. The price is broken if ROE mean-reverts toward the cost of capital as the credit-cost tailwind exhausts (Section 8) and revenue keeps lagging GDP (Section 5). The valuation is a leveraged bet on the ROE path, and the ROE path is exactly the variable the rest of this report has flagged as unsupported at the current price.

One more discipline the multi-stage model enforces: the currency. The cost of equity that matters to a foreign allocator is the *dollar* cost of equity, roughly 10.5–11% once the persistent rupee drag is netted against the INR cost of equity of ~13.3%. But the offset cuts both ways – the same currency drag that lowers the dollar discount rate also lowers the dollar cash flows, because rupee earnings translate into fewer dollars each year. The reverse-DCF cannot be run in rupees and then admired in dollars; the currency has to be carried consistently through both the discount rate and the cash flows, and when it is, the ~3.1%/year rupee depreciation of Section 10 is a direct subtraction from the dollar return regardless of how elegant the rupee-based valuation looks. The price of the story is a rupee price. The allocator pays in dollars.


9.4 The escape hatch is shut: active has not rescued the investor

Confronted with an expensive index, the natural allocator response is to seek refuge in selection – if the index is dear, own the cheap, high-quality subset and let a skilled active manager avoid the priced-for-perfection names. The Indian data says that hatch is shut.

EXHIBIT 17

Active selection has not rescued the India investor either

Share of India active large-cap funds underperforming their benchmark, to end-2025. A majority lag at every horizon.



~75% of active large-cap funds underperformed their benchmark over one year, and ~76% over ten years (S&P SPIVA India, year-end 2025). At every horizon measured, a clear majority of active large-cap managers lagged the index they were paid to beat.

Source: S&P DJI SPIVA India Year-End 2025 [GREEN]. Kakde Research.

The persistence of the underperformance across horizons is the point. A one-year figure could be noise – a single bad year for the average stock-picker. A ten-year figure of the same magnitude is structural: it says that over a full market cycle, selecting active large-cap managers has been, for three investors in four, a losing proposition against simply owning the index. This is not an India-specific anomaly; the same SPIVA pattern holds across most developed and emerging markets, and it is one of the most replicated findings in all of finance. But it lands with particular force on the Indian thesis, because the India bull case frequently retreats to selection when confronted with index-level valuation – "yes the index is expensive, but a good manager can find value." The data says the good manager, on average and over a decade, has not.

For the allocator the implication compounds the valuation problem rather than relieving it. The expensive index is the honest benchmark, because the active alternative has underperformed it three-quarters of the time over ten years, net of the fees that make the underperformance worse. An allocator cannot assume away the CAPE-36 starting point by outsourcing to a stock-picker, because the stock-picker has, in aggregate, delivered less than the expensive index – the fees are a certain cost and the alpha is an improbable benefit. This does not prove that *no* manager can beat the Indian index; a minority did, and identifying them in advance is the whole game and a legitimately hard one. It proves that the *expected* outcome of the selection strategy has been underperformance, and that an allocation case built on "we'll pick the winners" is building on a base rate that runs three-to-one against. The disciplined position is to treat the index valuation as the valuation that matters, to access it at the lowest cost available, and to hold any active tilt to a high burden of proof – the same burden of proof the manufacturing section demanded of individual names. Selection is not a solution to an expensive market; it is a second, independent bet that the Indian evidence says usually loses.

9.5 ALLOCATION TAKEAWAY

The price of the Indian story is high on the measure that matters most and defensible on the measures that matter less, and the allocator's job is to hold that tension without resolving it in favour of the comfortable reading. CAPE at 36 – half again its mean, reset-resistant, and the best long-run-return predictor of the three multiples – says the starting yield is poor and forward real returns are likely below average from here. Price-to-book at 2.95, roughly 18% below its own decade average, says India has genuinely compressed on the cleanest reset-resistant measure, and that is the strongest fact in the not-a-bubble case. Trailing P/E looks cheap and should be trusted least, because the 2021 standalone-to-consolidated reset flatters it. Weighted by trustworthiness, the balance tilts expensive.

The premium is a growth-persistence premium, not a spot-ROE premium: India is priced for sustaining high returns on equity for longer, not for earning more today, and the derived ROE (~14.4%) is close enough to EM (~13.8%) that the old claim of a demonstrated Indian ROE premium is retired to rough parity. Run the price backwards through a multi-stage reverse-DCF and the embedded assumption surfaces as the core tension: the multiple requires an ROE in the mid-to-high teens, above what India currently earns, sustained through the growth phase. Today's price is not priced for the India that exists; it is priced for an India whose ROE rises and holds. That is a coherent bet, but it is a bet, and the rest of this report has flagged the ROE path – margin durability, lagging revenue, exhausting credit-cost tailwinds – as the least supported variable at the current price. And the selection escape is closed: three-quarters of active large-cap managers have trailed the index over a decade, so the expensive index is the honest benchmark and cannot be assumed away by hiring a stock-picker.

For the allocator, three consequences. First, size the position to a poor starting yield: at CAPE 36, India is an underweight-to-neutral valuation case dressed in a strong growth narrative, and the growth narrative does not change the arithmetic of the entry multiple. Second, underwrite the ROE path explicitly, because that – not GDP growth, not demographics – is the variable the price is levered to; if you do not believe ROE rises toward the mid-teens and holds, you do not believe the price. Third, access the market cheaply and hold active tilts to a high bar, because selection has been a losing bet three times in four. Every number in this section, though, is a *rupee* number, and the allocator earns *dollars*. The valuation tension resolves – or fails to – only after the currency takes its cut, and the rupee is where this report goes next.

SECTION 10

The Rupee, Oil and the Outside World

The rupee is not the problem the bears make it out to be, and it is not the free option some bulls pretend it is. Measured properly, the Indian rupee is a persistent but genuinely low-volatility depreciator: it has slid against the dollar at roughly 3.1% a year over the last decade, and it has done so more smoothly than almost any other emerging

-3.8%/yr


FX LEG OF THE 20-YEAR RETURN BRIDGE

THE FINDING, UP FRONT. The rupee is not the problem the bears make it out to be, and it is not the free option some bulls pretend it is. Measured properly, the Indian rupee is a *persistent but genuinely low-volatility* depreciator: it has slid against the dollar at roughly 3.1% a year over the last decade, and it has done so more smoothly than almost any other emerging-market currency. Its annualised volatility, about 8.1%, is less than half of Brazil's (19.1%) and South Africa's (15.7%), and comfortably below Korea's (11.4%) and Mexico's (11.6%). Only the two overtly managed currencies in the peer set, China's renminbi and Vietnam's dong, are steadier. So the rupee is the quietest depreciator in the free-floating EM world. That is a real and underappreciated feature. But it does not make the currency a source of return, and it does not make hedging free. The forward market prices the depreciation in: the twelve-month forward premium sits near 2.8% a year, which is close enough to the trend depreciation that hedging the rupee costs you roughly what leaving it unhedged costs you. Hedged and unhedged are approximately a wash on trend; the difference is that hedging buys you lower volatility, not higher return. Anyone who writes "hedge the rupee" as though it were a costless improvement has not looked at the forward curve. This section measures the currency, traces the oil-to-inflation-to-policy chain that drives it, and then places it where it belongs in the return arithmetic: as the negative FX leg of the Return Bridge, dragging roughly 3.8 percentage points a year off the dollar return of Indian equity over twenty years.

EXHIBIT 18

The rupee depreciates steadily – but is one of the least volatile EM currencies

EM currency depreciation vs volatility, ~2016–2026. INR ~3.1%/yr at 8.1% vol vs Brazil 19.1%, South Africa 15.7%.



Source: BIS/market data via Kakde Research & Fable dataset [AMBER; RBI-corroborated]. Hedge cost ~2.8%/yr. Kakde Research.

10.1 Retire the "bottom-tier rupee" and measure what is actually there

The rupee debate is usually conducted in slogans: a "thirty years of trend" claim asserted without a multi-decade series ever being pulled, or a dismissive framing that lumps the rupee in with the EM basket's worst depreciators. Both are wrong in opposite directions. The rupee does depreciate, reliably, so the bulls who wave it away are wrong. But it depreciates *slowly and smoothly*, so the bears who file it next to the real and the rand are also wrong. This report retires both framings and replaces them with measurement.

Here is what the measurement shows, from the peer FX statistics assembled to 31-Jul-2026.

~3.1% a year — the rupee's annualised depreciation against the US dollar over the trailing ten years, at an annualised volatility of about **8.1%**. Over twenty years the currency has lost close to half its dollar value (a cumulative ~100% rise in USD/INR), but it has done so in a low-variance grind rather than in the step-devaluations that mark most of its peers.

EM CURRENCY (VS USD)	10Y ANN. DEPRECIATION	ANNUALISED VOLATILITY	WORST SINGLE YEAR
Vietnam (dong)*	1.4%	2.5%	+10.0% (2011)
China (renminbi)*	0.8%	4.8%	+8.6% (2022)
Taiwan (dollar)	-0.5%	5.1%	+10.8% (2022)
India (rupee)	3.1%	8.1%	+23.3% (2008)
Indonesia (rupiah)	1.9%	8.7%	+26.2% (2013)
Korea (won)	2.1%	11.4%	+41.2% (2008)
Mexico (peso)	0.3%	11.6%	+26.6% (2008)
South Africa (rand)	0.6%	15.7%	+37.8% (2008)
Brazil (real)	3.3%	19.1%	+49.3% (2015)

Vietnam and China are overtly managed; their low volatility is a policy choice, not a market outcome. Source: Kakde Research EM-FX statistics (S2, derived; RBI-corroborated). Values are AMBER-classed derived figures corroborated against RBI's own characterisation.

Two things fall out of the table. First, the rupee's *rate* of depreciation is middling, close to Brazil's, higher than the managed currencies, lower than the won over five years. It is a depreciator; that is not in dispute. Second, and this is the finding that matters, its *volatility* is exceptionally low for a currency that floats. Among the genuinely market-determined currencies in the set, the rupee is the calmest. Its 8.1% annualised volatility sits at less than half of Brazil's and South Africa's, and its worst single year, the 23% move in the 2008 crisis, is mild next to the won's 41%, the rand's 38%, or the real's 49%. The Reserve Bank of India's own characterisation, that the rupee has been "the least volatile" of the major EM currencies since 2023, corroborates the series rather than contradicting it. This is a managed float that has been managed well, in the narrow sense of variance suppression.

The allocator's translation is precise and it is not the reassurance it first sounds like. Low volatility is a risk statistic, not a return statistic. A currency can be extremely steady and still cost you money every year, and the rupee does exactly that: it debits the dollar investor a reliable 3% or so annually, with small standard deviation around that debit. Steadiness makes the loss *predictable*, not absent. For a global allocator sizing an India position, the right read is that the FX drag is a near-certainty to be planned for, not a tail risk to be hedged against. You do not lose sleep over the rupee; you lose return to it.

10.2 Hedging is priced in — do not write "hedge the rupee" as if it were free

This is the single most important correction in the section, and it is where the instinct of many India bulls goes wrong. The temptation is to say: the rupee depreciates 3% a year and costs me return, therefore I should hedge it and recover that 3%. The forward market has already thought of this.

~2.8% a year — the twelve-month USD/INR forward premium as of August 2026. That premium is the cost of hedging the rupee. Because it sits close to the trend depreciation of ~3.1%, hedging removes roughly as much cost as it saves.

The mechanism is covered interest parity, and it is not a market inefficiency waiting to be arbitrated; it is the market working. The forward price of the rupee embeds the interest-rate differential between India and the United States. Indian rupee rates sit structurally above dollar rates because Indian inflation and nominal growth run higher. To hedge a rupee exposure, you sell the rupee forward, and the forward is priced at a discount to spot by almost exactly that rate differential, which in turn tracks the expected depreciation. So when you pay the forward premium to hedge, you are pre-paying the depreciation you were trying to avoid. Over the trend, hedged and unhedged Indian equity land in roughly the same place in dollar terms: the unhedged investor eats the ~3.1% spot depreciation, the hedged investor pays the ~2.8% forward premium, and the two nearly cancel.

What hedging *does* change is the path, not the destination. It converts a 3%-a-year drag delivered with 8% annual volatility into a ~2.8% certain cost delivered with almost no volatility. For an allocator who marks to market and cares about tracking error, that variance reduction is worth something, especially in a crisis year when the rupee can move 20-plus percent and dominate the equity return. For a long-horizon investor who can sit through the volatility, hedging is close to a wash that also burns operational cost and rolls forward contracts every quarter. Neither investor should describe the hedge as "recovering" the depreciation. It does not. It swaps a volatile expected loss for a smooth near-certain loss of similar size.

The correct framing, then, is: - **Do not** treat the rupee's depreciation as a hedgeable free lunch. The forward market prices it; hedging pays for it up front. - **Do** treat hedging as a *volatility* tool: it strips the currency variance out of the equity return, at a cost that approximately equals the trend depreciation. - **Do** remember that both hedged and unhedged investors carry the drag. There is no version of owning Indian equity in dollars where the rupee is not a cost. The only question is whether you pay it smoothly (hedged) or lumpily (unhedged).

For most global EM allocators benchmarked in dollars and holding India at a single-digit weight, the pragmatic answer is to leave the position unhedged, accept the low-volatility drag, and size the position knowing the FX cost is baked in. The rupee's steadiness is precisely what makes that acceptable: an 8% -volatility drag is livable inside an equity position that itself carries 17% volatility. The FX leg is the small, quiet, reliable tax on the position, not the thing that will blow it up.

10.3 The transmission chain: oil, the current account, inflation and the RBI

The rupee's low volatility is not luck; it is the product of a policy framework built around a specific external vulnerability. India imports the overwhelming majority of the crude oil it consumes, and oil is the single largest line in its import bill. That makes the oil price the dominant exogenous driver of the currency, and the chain from crude to the rupee runs through the current account and monetary policy in a way worth tracing, because it explains both why the rupee depreciates persistently and why the central bank works so hard to keep the depreciation smooth.

The chain runs: **oil price up → import bill up → current-account deficit widens → more dollars demanded to pay for imports → downward pressure on the rupee → imported inflation (fuel, then second-round through transport and food) → RBI leans against it.** Each link is real and each is a channel the RBI actively manages.

Start with the current account. India runs a structural goods-trade deficit, financed by services exports (IT, business process work), remittances (among the largest in the world), and capital inflows (FDI plus portfolio flows into exactly the equity and bond markets this report studies). When oil is cheap, the current-account deficit is manageable and the rupee is calm. When oil spikes, the import bill balloons, the deficit widens, and the currency comes under pressure precisely when the country can least absorb it. The rupee's worst years line up with external shocks: 2008, 2013 (the "taper tantrum," when a widening deficit met a sudden stop in flows), 2018 and 2022 (oil spikes). The +23% move in 2008 and the depreciation clusters in 2013 and 2022 are the fingerprints of this oil-to-current-account channel.

Then inflation. A weaker rupee raises the domestic-currency cost of imported oil, which feeds directly into fuel prices and then, with a lag, into transport, food and the broad price level. India's monetary policy framework, a flexible inflation target centred on 4% with a tolerance band, makes the RBI a de facto currency manager even though the currency is not the formal target. The central bank cannot let the rupee slide disorderly, because a disorderly slide imports inflation and unanchors expectations. So it leans against sharp moves.

The tools are the third link. The RBI manages the rupee through foreign-exchange reserves (among the largest in the world, a war chest built precisely to smooth these shocks), through the policy rate (higher rates defend the currency by widening the rate differential and attracting flows), and through direct intervention in the spot and forward markets. The reserves are the reason the volatility is so low: the central bank spends them to cap the upside in USD/INR during stress and rebuilds them when flows return. This is the mechanical source of the "least volatile since 2023" observation. The rupee floats, but it floats inside a band the RBI polices with reserves.

This is the trilemma in action, and it is worth naming because it constrains what the RBI can promise. The impossible trinity says a central bank cannot simultaneously run an open capital account, a fixed exchange rate, and an independent monetary policy; it can have two. India has chosen a *partially* open capital account, a *managed* (not fixed) float, and a genuinely independent, inflation-targeting monetary policy. The managed float is the release valve: the rupee is allowed to depreciate on trend, which relieves the pressure that a hard peg would otherwise transmit into either reserves or domestic rates, while intervention keeps the *pace* of depreciation slow. Low volatility with persistent trend depreciation is exactly what you would predict from this corner of the trilemma. The RBI is not trying to hold the rupee flat, which would be futile against an inflation and growth differential; it is trying to let it slide smoothly, and it succeeds.

It is worth being concrete about the two episodes that define the rupee's risk profile, because they are the template for the next one. In 2013, the "taper tantrum," the mere signal that the Federal Reserve would slow its bond purchases triggered a sudden stop in portfolio flows into EM. India, then running a wide current-account deficit financed by exactly those flows, was among the "Fragile Five" hit hardest; the rupee fell sharply and the RBI had to raise rates and open special swap windows to defend it. The lesson the central bank drew was structural: build a reserve buffer large enough that the next sudden stop can be absorbed without a disorderly move. The 2022 episode, a genuine oil-and-dollar double shock as crude spiked and the Fed hiked aggressively, was the test of that lesson, and the rupee's ~10% move that year was orderly by the standard of its own 2013 and 2008 experience. The difference between 2013 and 2022 is the difference between a country that gets run over by a dollar shock and one that leans into it with reserves. That accumulated resilience, a narrower deficit, record remittances, a local bond market now included in global indices and pulling in structural inflows, is why the volatility has compressed and why the "least volatile since 2023" characterisation holds. It is also, precisely, what an allocator is trusting when they leave the position unhedged.

The global dollar system sits underneath all of this. India is a dollar-short economy: it must earn or borrow dollars to pay for oil and to service external liabilities, and the price and availability of those dollars is set in a system it does not control. When the Federal Reserve tightens and the dollar strengthens globally, the rupee weakens regardless of anything happening in India, because the whole EM complex is on the other side of a stronger dollar. The 2013 and 2022 episodes were as much US-rate stories as India stories. This is the sense in which the currency is "the outside world's" leg of the India trade: it imports a global monetary condition that no amount of domestic reform can insulate against. India's reserves and its improved external metrics (a narrower deficit, larger remittances, a deeper local bond market now in global indices) have made it more resilient to these shocks than it was in 2013. But resilient is not immune, and the structural dollar-shortness means the rupee will keep depreciating on trend as long as India grows and inflates faster than the United States, which is to say, for the foreseeable horizon of any allocation decision.

10.4 The currency as the FX leg of the Return Bridge

Everything above resolves into one number in the return arithmetic, and it is the cleanest way to hand the currency question to an allocator. The Return Bridge, the single-universe decomposition of what actually produced India's dollar equity return over twenty years, breaks the ~7.1% annualised USD return into four legs: earnings growth, valuation change, dividends, and currency.

-3.8% a year — the FX leg of the twenty-year Return Bridge. The full decomposition is **+9.7% EPS growth – 0.1% P/E change + 1.3% dividend – 3.8% FX ≈ 7.1% USD**. The rupee is the second-largest term in the bridge and the only large negative one.

Read that decomposition and the currency's role becomes unmistakable. Indian companies grew earnings at a genuinely world-class ~9.7% a year in rupee terms over two decades, the engine of the whole story. But the currency handed back roughly 3.8 percentage points of that every year, converting a spectacular rupee earnings record into a merely good dollar return. Note that the FX drag in the Bridge (-3.8%/yr over 20 years) is somewhat larger than the trailing-ten-year depreciation rate (~3.1%/yr), because the twenty-year window includes the higher-depreciation earlier decade; the currency has, if anything, become a *smaller* drag over time as India's external position strengthened, which is itself part of the resilience story. But over any horizon a dollar investor should plan for, the FX leg is negative, material, and reliable.

This is why the currency cannot be waved away and cannot be wished into an advantage. It is the mechanism that turns India's best-in-class earnings growth into a return that, as the next section shows, has actually *trailed* the EM index over five and ten years in dollar terms. The rupee is not the villain of the India story; the earnings engine is real and the currency's steadiness is a genuine virtue. But it is the reason the local-currency and hard-currency versions of the India thesis are two different investments, and it is the reason a dollar-benchmarked allocator must underwrite India on its dollar return, net of a currency that will, quietly and reliably, take its cut.

ALLOCATION TAKEAWAY — SECTION 10

- **Size the rupee as a cost, not a risk.** The FX drag is a near-certain ~3% a year (higher, ~3.8%, over the full 20-year Bridge), delivered at low ~8% volatility. Plan for it in the base case; do not treat it as a tail to be insured.
- **Do not hedge expecting to recover the depreciation.** The 12-month forward premium (~2.8%) prices in the trend depreciation. Hedged and unhedged are roughly a wash on return; hedging only buys you lower volatility, at a cost that approximately equals what it saves.
- **If you hold India at a single-digit EM weight and mark to a dollar benchmark, leave it unhedged.** The rupee's low volatility makes the drag livable inside a 17%-volatility equity position, and hedging burns cost and operational friction for a variance reduction most long-horizon allocators do not need.
- **Watch oil, the dollar and the RBI's reserves, not the rupee itself.** The currency's next disorderly move, if it comes, will be an oil spike or a Fed-driven dollar surge transmitted through India's structural dollar-shortness. The reserves are the shock absorber; the day they are visibly depleted is the day the low-volatility assumption is at risk.
- **Underwrite the dollar return, always.** India's rupee earnings growth is genuinely world-class. The currency is the reason that does not translate one-for-one into your return. The comparison that follows is conducted entirely in dollars for exactly this reason.

SECTION 11

India Against the Alternatives

Held to one ruler, in one currency, over the same window, India is a mid-table emerging market, not a leader. On MSCI net total-return data in US dollars to 31 July 2026, Indian equity returned 8.03% a year over ten years.

Mid-table


ON ONE RULER, ONE CURRENCY, ONE WINDOW

THE FINDING, UP FRONT. Held to one ruler, in one currency, over the same window, India is a mid-table emerging market, not a leader. On MSCI net total-return data in US dollars to 31 July 2026, Indian equity returned **8.03% a year over ten years**. That is *below* Taiwan (22.5%), Korea (15.3%) and the United States (14.5%); *below* the MSCI EM index itself (9.19%) over the same decade; and only *above* the laggards, China (4.9%) and Indonesia (−3.9%). This is the single most under-used fact in India coverage, and this report promotes it to the front of the comparison: India underperformed the EM index over both five and ten years in dollar terms. The bull case is built on India having the best growth and the best earnings engine in the emerging world, and both may be true. What is *not* true is that this translated into the best returns. The link from GDP growth to equity return is, on this thirteen-country cross-section, statistically absent: the slope is slightly negative, the fit is near zero, and with only thirteen data points no relationship survives a significance test. Growth is not return. This section runs India against its real alternatives, China, Taiwan, the EM ex-China index, on one consistent basis, retires the flawed "capture ratio" heuristic, and closes by refusing to crown a single winner, because the honest answer is that different markets serve different objectives.

EXHIBIT 19

Across 13 markets, faster growth did not buy higher returns

Equity return (MSCI net TR, USD) vs real GDP growth, 2015–2026. Vietnam grew fastest (6.2%); Taiwan returned most.



Source: MSCI factsheets 31-Jul-2026; World Bank WDI [GREEN]. Slope/ R^2 Kakde Research. Kakde Research.

11.1 One ruler: the matched dollar-return table

The discipline that separates this comparison from most India cheerleading is boring and non-negotiable: one index family (MSCI), one return definition (net total return, dividends reinvested, withholding deducted), one currency (US dollars), one end date (31 July 2026), the same windows for every country. No local-currency figures smuggled in for India, no price return compared against total return, no cherry-picked start date. Here is the matched set.

8.03% — MSCI India's ten-year annualised net USD total return to 31-Jul-2026. The MSCI EM index returned **9.19%** over the same decade. India trailed the benchmark it is supposed to lead.

MSCI INDEX (NET USD)	10Y ANN.	5Y ANN.	1Y	TRAILING P/E
Taiwan	22.5%	23.4% 83.9%		31.3
Korea	15.3%	17.1% 150.1%		17.2
United States	14.5%	11.9% 18.7%		27.1
World (DM)	12.7%	11.2% 20.4%		24.3
EM ex-China	11.1%	12.4% 52.2%		19.1
Japan	9.3%	10.0% 32.3%		21.0
EM (benchmark)	9.2%	8.0% 36.4%		17.7
UK	9.1%	13.0% 25.6%		16.5
Saudi Arabia	8.9%	2.3% 3.9%		14.7
India	8.0%	5.3% -6.5%		23.9
Mexico	7.7%	13.3% 35.0%		14.7
Brazil	7.2%	8.1% 44.7%		10.1
South Africa	6.3%	10.8% 28.8%		12.5
Vietnam	6.1%	0.0% 24.9%		18.3
China	4.9%	-2.1% -1.1%		14.1
Indonesia	-3.9%	-6.2% -33.4%		10.7

Source: MSCI factsheets, net total return in USD, annualised, to 31-Jul-2026 (GREEN). The matched set is the discipline: one provider, one currency, one basis, one date.

Read the table without a thesis and three facts stand out. First, India is squarely mid-table on the ten-year window and *below the EM index* on both the five- and ten-year windows, the fact this report refuses to bury. Second, the decade's actual winners, Taiwan and Korea, are semiconductor-and-hardware stories riding the AI capital-expenditure cycle, not consumption-and-demographics stories of the kind the India bull case tells. Third, India's most recent year is *negative* (-6.5%), during a period when the EM index returned +36% and Korea +150%, because this is precisely the window in which India de-rated from its 2024 peak while Taiwan and Korea surged. India is also, tellingly, the *second-most-expensive* market in the table on trailing P/E (23.9, behind only Taiwan), which is the valuation setup of a market that has already been paid for its growth, not one about to deliver a surprise.

The five-year column is if anything more sobering for the bull: India's 5.3% annualised over five years sits below Mexico, Brazil, South Africa, the EM index, EM ex-China and the developed world. The market that "everyone should own" has, on the most recent half-decade in the currency a global investor actually banks, been a laggard. That is not an argument against ever owning India; valuations and cycles turn. It is an argument against owning it *because you assume its growth guarantees its return*. It has not.

11.2 Growth is not return — and the statistics say so

This is the analytical heart of the section, and the place where the analysis must show the object rather than wave its hands at it. The bull syllogism is: India has the fastest large-economy growth, therefore India will deliver the best equity returns. The cross-section of the thirteen major markets, plotting each country's real GDP growth against its dollar equity return, tests that syllogism directly, and it fails.

On the thirteen-country cross-section of GDP growth versus USD equity return, the fitted slope is **slightly negative**, the **R² is near zero**, and with **n = 13** no coefficient clears a conventional significance threshold. The honest statement is not "growth hurts returns" but **"there is no statistically reliable positive relationship between a country's growth and its equity return."**

The precision matters: a scatter published with no slope, no confidence interval, no R² and no sample size is an illustration, not a test. Stating the statistics does not let us claim the opposite of the bull case, that growth is bad for returns. The slope is negative but it is statistically indistinguishable from zero, and thirteen points is far too thin a sample to assert a real negative relationship. What the data supports, and all it supports, is the null: knowing a country's growth rate tells you essentially nothing reliable about its equity return. Taiwan grew slowly and returned 22%; India grew fast and returned 8%; China grew fast and returned 5%; Indonesia grew respectably and lost money. The scatter is a cloud, not a line.

The mechanism behind the null is not mysterious, and it is the through-line of this entire report: GDP growth accrues to an economy, but equity returns accrue to *shareholders*, and the leakages between the two are large and country-specific. Growth can be diluted away by new share issuance (India issues a lot of equity, a genuine drag on per-share earnings). It can be competed away into thin margins. It can be captured by unlisted firms, state enterprises, or foreign strategics rather than by the listed, free-float, foreign-investable companies an allocator can own, which is the "Four Indias" problem of Section 2 restated at the index level. And it can be paid for in advance through a high entry valuation, so that even delivered growth produces a poor return because you overpaid for it, which is India's specific situation at a 23.9 P/E and a CAPE near 1.5× its own history. Fast growth at a rich multiple with heavy dilution is a recipe for exactly the mid-table return the table records. This is Fable #21 stated cleanly: the relationship the bull case assumes is not in the data.

11.3 Retire the capture ratio

A "capture ratio" — an attempt to summarise, in a single multiplier, how much of a country's GDP growth its equity market "captured" as return — is a tempting construction, and this research built and tested one before retiring it. The retirement is deliberate, and the reasons are stated in full so the construct is not quietly resurrected.

Retired: the capture ratio. The single-number "growth-to-return capture" multiplier is not a stable or well-defined statistic. It was built, tested and discarded in the course of this research, and it stays retired — with the reasons stated rather than the construct silently dropped.

The reasons it fails are worth stating so the construct is not resurrected. First, it is a ratio of two noisy numbers over a chosen window, and it is wildly unstable to the start and end dates, exactly the start-point sensitivity that plagues the five-year return window. Second, it presumes the very linear growth-to-return relationship that Section 11.2 just showed does not exist; you cannot compute a meaningful "fraction of growth captured" when the underlying correlation is statistically zero. Third, it invites precisely the false-precision sin this report elsewhere demolishes: a made-up multiplier with a decimal point, dressed as a measurement. The correct replacement is the decomposition

already built, the Return Bridge, which shows *where* the return came from (earnings, valuation, dividends, currency) without pretending there is a stable scalar linking GDP to price. The Bridge is a mechanism; the capture ratio was a mirage. Keeping it retired is not a loss of analytical firepower; it is the removal of a false one.


11.4 EM ex-China: India as a core comparator, not a special case

The most important structural change in how allocators frame emerging markets is the rise of the "EM ex-China" mandate, and it reshapes where India sits. As global investors carved China out into dedicated allocations, the EM ex-China index became a real, large, benchmarked universe in its own right, and it is the one India must now be measured inside.

EXHIBIT 20

Even ex-China, India is the third weight — behind Taiwan and Korea

Country weights, MSCI EM vs MSCI EM ex-China, 31-Jul-2026. India 14.8% ex-China, not the ~30% often assumed.



In MSCI EM ex-China, India's index weight is **14.84%** — third, behind Taiwan (~33.9%) and Korea (~25.9%). Not the ~30% one might assume from India's economic size, and not a dominant position. India is one of three heavyweights, and the smallest of the three.

Source: MSCI factsheets, 31-Jul-2026 [GREEN]. Kakde Research.

This reframes India twice over. On *weight*, India is a core holding in EM ex-China but a junior partner to the two semiconductor economies, the same Taiwan and Korea that dominated the return table. An allocator running an EM ex-China book is, whether they intend it or not, running a large bet on Asian hardware with India as the third leg. On *return*, the EM ex-China index itself compounded at 11.1% over ten years and 12.4% over five, comfortably ahead of India's 8.0% and 5.3%, because Taiwan and Korea carried it. So the index India is now benchmarked against has *out-returned India* over both windows, which means India has been a *drag* on EM ex-China performance over the recent past, not a driver of it. The comparator is no longer the sluggish all-EM index dragged down by China; it is a sharper, hardware-led index that India has trailed. That raises the bar for an India overweight, not lowers it.

The implication for position sizing is concrete. A global allocator no longer decides "India versus China" in isolation; they decide India's weight *within* an EM ex-China sleeve where Taiwan and Korea already dominate and have already delivered. Overweighting India in that context is an active decision to underweight the very markets that produced the returns, and it needs a forward-looking reason, cheaper relative valuation, a turn in the AI capex cycle, a domestic-demand catalyst, that the trailing data does not supply on its own.

11.5 One Dollar, Ten Markets: no single winner

The cleanest way to end the comparison is to refuse the question it invites. "Which emerging market should I own?" has no single answer, because the markets differ not on a single quality axis but on which *objective* they serve. The One Dollar, Ten Markets scorecard makes this explicit: run one dollar into each market and score it against different allocator goals, and no market wins on all of them.

EXHIBIT 21

One Dollar, Ten Markets – no single winner; India leads on growth, lags on price and return

Allocator scorecard, 1 (weak) to 5 (strong), India-relative. Read across a row for a market; down a column for who wins each dimension.

	10y USD return	Real GDP growth	Valuation (P/E)	Div+buyback yield	FX stability	Foreign access	Governance
India	3	5	2	2	4	3	3
China	2	5	4	4	3	2	2
Taiwan	5	4	2	3	5	4	4
S. Korea	5	3	4	3	3	4	3
Indonesia	1	4	4	5	3	3	3
Brazil	4	2	5	5	1	4	3
Mexico	4	1	4	4	2	4	3
S. Africa	4	1	4	4	2	3	3
Saudi	4	3	4	4	5	2	2
Vietnam	3	5	3	3	5	2	2

SCALE 1 2 3 4 5 1 weak · 5 strong

No market sweeps the scorecard. Taiwan and Korea won the last decade's return but carry concentration and geopolitical tail risk and now trade rich. China is cheap and hated. Brazil and Mexico offer value and yield with currency and political volatility. India offers the best earnings *growth* and the deepest domestic *demand story* at the *highest EM valuation* and *below-benchmark realised return*. Each is a different instrument for a different mandate.

Source: Return/growth/valuation/FX from MSCI, World Bank, market data [GREEN]; access/governance = Kakde Research judgment, labelled [INTERP]. Kakde Research.

Read the scorecard as a set of objective-to-market matches rather than a ranking:

- **For realised ten-year return:** Taiwan, then Korea, then the US. India is mid-table. If the objective is backward-looking performance, India did not deliver it.
- **For cheapness / mean-reversion:** China, Brazil, South Africa, Indonesia. India is the *opposite* trade, expensive and already re-rated.
- **For dividend income:** Indonesia (7.7%), Brazil (5.6%), Mexico (3.9%). India yields 1.27%, near the bottom.
- **For a low-volatility currency:** India is genuinely differentiated here (Section 10), a real point in its favour for a risk-managed dollar investor.
- **For domestic-demand, low-export-dependence exposure:** India is the cleanest large expression, the strongest part of the structural case, its high domestic-profit intensity insulating it from the trade and tariff cycle that whips Taiwan, Korea and China.
- **For AI / semiconductor capex exposure:** Taiwan and Korea, decisively. India has almost none of this, which is why it missed the last two years' returns.

Now place India directly against its three real alternatives, on this basis:

China vs India. The purest value-versus-growth pair in EM. China is cheap (14.1 P/E, 1.5 P/B), structurally distrusted, and has returned 4.9% over ten years and *lost* money over five. India is expensive (23.9 P/E), structurally loved, and has returned 8.0% and 5.3%. The allocator's real choice is between a hated cheap market with policy and governance tail risk, and a beloved expensive market with valuation and dilution risk. They are not substitutes; a barbell that holds both is a defensible position, and an EM investor who owns *neither* is making a large active bet.

Taiwan vs India. The last decade's decisive divergence: Taiwan compounded at 22.5% to India's 8.0%, on *slower* GDP growth, because it sits at the centre of the AI hardware cycle and India does not. Taiwan carries the return record and the semiconductor concentration *and* the sharpest geopolitical tail in the EM set. India carries the growth story, the domestic-demand insulation and the rich valuation. This pair is the single clearest refutation of "growth equals return": the slower-growing economy delivered nearly triple the equity return.

EM ex-China vs India. As Section 11.4 established, the benchmark most allocators now use has out-performed India over five and ten years, with India as its third-largest and worst-performing heavyweight. The burden is on the India *overweight* to justify itself against a comparator India has trailed, not the other way round.

There is one more comparison the scorecard forces, and it is India against *itself* across the cycle, because the round-trip in India's own benchmark weight is the sharpest single illustration of growth-is-not-return. India's weight in MSCI EM ran from roughly 9% in 2020 to a peak near 20% in 2024, briefly the single largest country in the broader investable universe, and then *fell back to 11.7% by mid-2026* as Taiwan and Korea surged on the AI cycle and India de-rated. That is not a story of a market steadily compounding its way to dominance; it is a story of a market that over-earned its weight on a wave of domestic enthusiasm and heavy issuance, re-rated to the top of the EM valuation table, and then gave a chunk of it back. The allocator who bought the "India is becoming the largest EM" narrative at the 2024 peak bought the top of a round-trip. This is the empirical shape of the growth trap: the growth was real throughout, but the *return* to a dollar investor depended entirely on the entry valuation, and the entry valuation was set by sentiment that mean-reverted. The lesson is not that India is uninvestable; it is that India is a market where *when you buy dominates that you buy*, precisely because the growth is priced in advance.

The 1-year column of the return table makes the point in miniature and is worth not skipping past. India returned -6.5% over the trailing year while the EM index returned +36% and Korea +150%. A market can have the best structural story in the emerging world and still be the worst performer in a given year, because it entered that year

expensive and de-rating while the money rotated into cheaper, more cyclical, AI-levered names. This is the single-year face of the ten-year truth: India's problem has not been its economy or its earnings; it has been the price you had to pay to participate and the currency that quietly took its cut on the way out.

ALLOCATION TAKEAWAY — SECTION 11

- **Underwrite India on its dollar return, mid-table and below-benchmark, not on its growth rank.** India returned 8.0% over ten years and 5.3% over five, below the EM index on both. The growth story is real; it has not been a return story. Own India for what it forwardly offers, not for a GDP number that the data shows does not translate.
- **Do not resurrect the capture ratio or any single growth-to-return multiplier.** The growth-return relationship across markets is statistically zero (slope slightly negative, R^2 near zero, $n = 13$). Use the Return Bridge decomposition, not a scalar, to reason about where returns come from.
- **Benchmark India inside EM ex-China, where it is the third heavyweight at 14.8% and has been the laggard of the three.** An India overweight is an active underweight of Taiwan and Korea; it needs a forward catalyst (relative valuation, a demand trigger, an AI-cycle turn) the trailing record does not provide.
- **Match the market to the mandate.** For realised return, Taiwan/Korea; for value, China/Brazil; for income, Indonesia/Brazil; for a stable currency and clean domestic-demand exposure, India. No market wins on every axis, so size India for the objective it actually serves, domestic-demand, low-export-beta, low-currency-volatility exposure at a premium price, and not as the all-purpose EM answer.
- **The honest allocator's answer:** India belongs in a diversified EM book at a weight near or modestly above its ~11.7% benchmark *only if* you are buying its forward domestic-demand story and can tolerate paying the highest EM multiple for below-benchmark recent returns. If your objective is the last decade's performance, or cheapness, or income, the data points elsewhere. India is a hold-for-the-structural-story market, not a chase-the-return one.

SECTION 12

The Allocator's Answer

Every section to here has built toward a single decision: given one more dollar, how much India should a global investor hold, and out of what.

Hold, near benchmark

EARN THE OVERWEIGHT — DO NOT ASSUME IT

Every section to here has built toward a single decision: given one more dollar, how much India should a global investor hold, and out of what. This section answers it. It does not manufacture a Buy or a Sell, because the honest answer is not a rating – it is a sizing decision conditioned on price, objective, and funding source. We revisit the return bridge one last time, separate what a foreign dollar *can capture* from what it *cannot own*, subtract what implementation actually takes back, and then size the position by mandate. We close with one canonical falsification dashboard and the Kakde Research view on a single page.

Revisit the bridge: what the dollar is actually buying

The twenty-year, single-universe, USD return decomposition is the spine of the whole report, so start there once more: **+9.7 (EPS) – 0.1 (P/E) + 1.3 (dividend) – 3.8 (FX) $\approx$ 7.1% a year**. Read as an allocator, that sentence says four things at once. The earnings engine is real and large. Dividends are a modest positive. The multiple is a round-trip that nets to nothing across a full cycle. And the currency is a persistent tax that removes roughly 40% of the earnings compounding before it reaches a dollar account.

The single most important reframe in this report follows directly. Because the multiple contributed –0.1 over twenty years, the marginal dollar into India is not a bet on GDP and not a passive bet on re-rating. It is a bet on two things and only two things: **that earnings keep compounding near their double-digit pace, and that the multiple does not compress from today's rich level**. From a CAPE of 36 against a mean near 24, the multiple is more likely to subtract than to add. That means the forward return the bridge implies is *lower* than the historical 7.1% unless earnings do more of the work – which is precisely the claim the allocation is underwriting.

There is an internal tension worth naming because it is the sharpest one in the report. The reverse-DCF that the market's current price implies requires a perpetual growth rate around 11.8% on a single-stage read, which in turn implies a return on equity meaningfully *above* the ~14% India appears to earn today. In plain terms, **the price already assumes India's profitability rises from here**, not merely that it holds. That is a demanding embedded assumption on top of a profit share already at an all-time high. The bull needs profitability to climb from a record level; the bear only needs it to stop climbing. The bridge and the reverse-DCF are telling the same story from two directions: the earnings leg has done the work historically, and the price now asks it to do even more.

What global investors CAN capture

A foreign dollar can capture the earnings leg, net of currency, and it is a genuinely good leg. India's +9.7%/yr earnings contribution is among the strongest in the emerging-market set, built on a formal corporate sector that includes a world-class dollar-earning IT and business-services export industry, a large and profitable banking system, and an energy-and-telecom complex. Dividends add a reliable ~1.3%/yr. And the vehicle to capture it is liquid: India is the world's fifth-largest equity market by capitalisation at ~\$5.05 trillion, with deep single-stock liquidity, functioning index products, and enough investable float that a global allocator can build and unwind a position without moving the market. The capturable object is a ~7% USD compounder historically, at ~17% volatility, in an orderly currency. For a portfolio that needs an EM earnings engine, that is a real and ownable asset.

It is worth being precise about *why* the earnings leg is the one that survives to the foreign investor while GDP does not. Earnings are the point in the transmission chain where the formal, listed, corporate slice of India is measured directly – the IT exporters booking dollar revenue, the banks earning a spread on a formalising economy, the consumer companies capturing the organised share of a fragmented market. That slice grows faster than aggregate GDP in the periods when formalisation is taking share, which flatters the earnings line relative to the economy. So the +9.7 leg is not GDP passed through; it is the *formal-sector share-gain* plus organic growth of the listed cohort, which is a different and, for the investor, better thing. The catch is that this same mechanism is why breadth is

limited: the earnings engine is powerful precisely because it is concentrated in the formal names a foreign dollar can own, and those names are a minority of the economy. The strength of the leg and the narrowness of what is ownable are two faces of the same fact.

What global investors CANNOT own

The gap between India-the-economy and India-the-investable-slice is wide, and it is where the growth leaks before it ever reaches the bridge. Three walls stand between the GDP number and the foreign dollar.

The informal wall. The overwhelming majority of Indian economic activity happens in the informal economy – small enterprises, self-employment, casual labour, agriculture – in businesses that will never list and whose growth never touches a public earnings line. (Widely cited estimates put informal employment near 90% of the workforce; treat that as an illustrative order of magnitude, not a canonical figure, but the fact is not in dispute.) The transmission from GDP growth to *listed* revenue is structurally lossy because so much of the economy is unlistable.

The promoter wall. Promoters hold **50.0%** of the listed market. Half of India's market capitalisation is controlling blocks that are not for sale at any index-level price. Free float is ~\$2.67 trillion against ~\$5.35 trillion of listed cap. A foreign investor competes for the other half with domestic institutions that hold 16.1% and retail that holds 18.5%.

The foreign-investability wall. Even within the float, sector caps, strategic holdings, and the mechanics of foreign-ownership limits mean the foreign-investable cap is smaller than the float. Actual foreign custody is ~\$0.84–0.93 trillion, and – the distinction the report insists on – **that is what foreigners hold, not the maximum they could own.** The remaining foreign room is the free-float foreign-investable cap minus current holdings, a genuinely smaller number than the headline suggests. Held is not ownable.

The concentration of what *is* ownable compounds the problem. The Nifty's top ten names are **52.9%** of the index. A foreign investor buying India is buying a highly concentrated slice – a few financials, an energy-and-telecom conglomerate, IT majors, a handful of consumer names – not a broad claim on the 1.4-billion-person economy. The breadth of the real economy does not translate into breadth of the investable one.

The geographic version of the same problem sits underneath. India is not one economy but several, and the dispersion is extreme: the richest large state (Karnataka, ~₹4.22 lakh per-capita income) out-earns the poorest (Bihar, ~₹0.76 lakh) by roughly **5.6×**, a gap wider than that between many separate countries. The "1.4 billion consumers" framing overstates the addressable market, which is concentrated in a handful of high-income states where the listed consumer companies actually sell. The growth runway is long *if* the lagging states converge, but that convergence has been slow, and the index owns the formal, urban, high-income slice – not the informal, rural base where most of the population and most of the unrealised growth sit. What a foreign dollar can own is narrower than India-the-economy on every axis at once: formal versus informal, floated versus promoter-held, top-ten versus broad, and rich-state versus poor-state.

What implementation takes back

The 7.1% historical USD return is a gross, pre-friction number. A real dollar keeps less, and the report is explicit about each subtraction because they compound.

Currency. The –3.8 leg is already inside the bridge, so it is not double-counted here, but the *decision* about it is not. An unhedged dollar wears the full ~3.1%/yr trend depreciation. A hedged dollar pays ~2.8%/yr (the 12-month forward premium, August 2026) to remove it – close to a wash on the central estimate, but it converts an uncertain currency outcome into a known, smaller cost and removes the tail. For most institutional dollars the hedge decision is worth making deliberately; it moves realised returns by more than most security-selection choices inside the sleeve.

Withholding and tax friction. India levies withholding on dividends paid to foreign investors, and capital-gains treatment for foreign holders carries its own frictions and holding-period sensitivities. The dividend leg is +1.3%/yr gross; a foreign holder keeps less than that after withholding. This is a real, recurring subtraction, not a one-off, and it lands specifically on the one positive non-earnings leg in the bridge.

Vehicle fees. The dollar that accesses India through an EM or India-dedicated ETF or active fund pays an expense ratio, and access products for a single EM country tend to sit at the higher end of the fee range. Active India funds face a steeper hurdle still, because the market is hard to beat: SPIVA India shows **75.0% of large-cap active funds underperformed over one year and 76.3% over ten**. Paying active fees to underperform three-quarters of the time is a subtraction the passive route avoids, but the passive route pays the tracking and access costs instead.

Net the frictions and the point is plain: **the USD investor keeps meaningfully less than the ~7.1% gross bridge return** – after an unhedged currency drag or a near-offsetting hedge cost, after dividend withholding, and after vehicle fees. The forward number, from a rich multiple, is lower still. This is not an argument against owning India. It is the discipline that turns a headline return into a portfolio input.

It is worth stacking the subtractions in one place, because their order matters and each is defensible individually. Start with the ~7.1% gross bridge return. The currency drag is already embedded, so the first live decision is the hedge: leave it unhedged and wear ~3.1%/yr of trend depreciation with its 8.1% volatility, or hedge at ~2.8%/yr and convert it to a known cost – the two are close on the central estimate, which is the report's point that the rupee is a manageable line item rather than a disqualifier. Then subtract dividend withholding, which lands on the +1.3 dividend leg. Then subtract the vehicle fee, higher for a single-country access product and higher still for active. What remains is a real dollar return that is respectable but clearly below the ~9.2% EM average – and the forward version starts lower again because the multiple, which added nothing historically, now begins from a rich level with more room to fall than to rise. None of these is a reason to avoid India. Together they are the reason the position is sized to benchmark and not above it.

Portfolio construction and sizing, by objective

There is no single "India allocation." The right size depends entirely on the mandate, and the report treats four archetypes explicitly.

The benchmark-aware EM fund. India is already **11.66%** of MSCI EM. For this manager the only decision is the deviation from that weight, and every basis point of overweight is an active bet that must be funded by underweighting something else in the same book. Given a below-EM realised return, a rich CAPE, and a multiple that adds nothing, the defensible default is **at or modestly below benchmark**, with the overweight reserved for a clear improvement in price or a confirmed earnings-durability signal. What India *adds* to this book is a high-quality earnings engine and the deepest domestic buyer base in EM; what it *risks* is ~17% volatility and valuation compression; what an overweight *displaces* is whatever else the manager holds – most consequentially Taiwan, which delivered the returns India's price already embedded.

The EM ex-China fund. Here India is structurally larger at **14.84%**, because removing China mechanically raises every other country's weight. A manager running this increasingly popular construction holds materially more India by default and should recognise that the ex-China book is, in part, a concentrated bet on India, Taiwan and Korea. The India sizing decision is therefore *more* consequential in this mandate, not less, and the same valuation discipline applies from a higher starting weight.

The global (ACWI-type) equity fund. India is a low-single-digit share of a global book. For this manager the question is not the deviation but whether India warrants any dedicated active attention at all, given the friction of accessing a single EM country. The defensible answer: India earns a place through its EM sleeve at benchmark, and a *dedicated* global overweight requires conviction on the specific earnings-and-margin bet strong enough to justify the access cost and the ~17% volatility it imports.

The family office / unconstrained pool. This is the only archetype that can express India the way its structure actually rewards. Freed from a benchmark, a long-horizon pool can access single stocks rather than the concentrated top-ten index, can choose its currency posture deliberately, can hold through a -72.6%-type drawdown, and can wait for CAPE to improve before adding. For this investor India is most attractive *selectively* – specific compounders at specific prices – rather than as an index beta. The report's "selectively investable" verdict is aimed squarely here: the index is the least efficient way to own India, and the unconstrained pool is the one investor who does not have to.

Across all four, the same three constraints bind: India adds a genuine earnings engine, it imports ~17% volatility and a below-EM expected return, and it displaces whatever funds it – which is the decision the next subsection isolates.

A word on how much, concretely. Because India already sits at 11.66% of EM and 14.84% of EM ex-China, the *active* decision for a benchmark-aware manager is measured in the deviation, not the level, and the evidence in this report argues for keeping that deviation small and negative-to-neutral until price improves: a below-EM realised return (8.03% vs 9.19% over ten years) and a CAPE at ~1.5× its mean do not support paying up for an overweight. For the unconstrained pool the "how much" is bounded instead by risk: at ~17% volatility with a -72.6% historical tail, a standalone India sleeve of more than a modest single-digit share of a diversified book imports concentration risk that the ~7% expected return does not obviously compensate. In both framings the sizing conclusion is the same – India earns its benchmark weight and has to *prove* an overweight – but the binding constraint differs, tracking error for one investor and absolute drawdown for the other.


Where the dollar comes from

Sizing is half the decision; funding source is the other half, and it is the half most India research ignores. A marginal dollar into India is a marginal dollar out of something specific, and the something matters.

EXHIBIT 22

Where the India dollar comes from matters as much as the India view

Illustrative relative risk of funding a +200bps India overweight from different sources. Funding from Taiwan/China is lower-tracking than from cash.



- **Funded from Taiwan:** you are selling the market that *beat* India — the one that delivered the returns India's valuation had already priced. This is the highest-conviction funding source only if you specifically believe the AI-capex cycle that drove Taiwan is exhausted. Absent that belief, it funds an overweight by selling the winner.
- **Funded from Korea:** you are selling a market that is *cheaper* than India. This trade requires India's earnings-and-margin bet to more than overcome the valuation gap it opens.
- **Funded from cash:** cash now yields something again, so this is a genuine opportunity-cost trade. It raises the book's equity beta and imports ~17% volatility and a -72.6% tail into a position that was previously earning a real risk-free rate. It is the cleanest expression of "I want more India," and the most honest about the risk it adds.
- **Funded from US / developed markets:** the most common funding source for EM overweights in practice, and the most defensible for an India overweight, because it does not require a relative call *within* EM. It is a straightforward "more EM earnings engine, less developed-market beta" decision, and it isolates the India bet from the Taiwan-versus-India and Korea-versus-India questions that the intra-EM funding sources force.

Source: Kakde Research portfolio framework [ILLUSTRATIVE/INTERP – labelled]; correlations pending Bloomberg pull. Kakde Research.

The report's guidance is explicit: **an India overweight is most defensible funded from the developed-market or cash sleeve, and least defensible funded by selling Taiwan or Korea**, because those two are the very markets whose outperformance the return data documents. Funding source is not an afterthought to sizing; for an intra-EM allocator it *is* the trade.

There is a second-order point the funding chart makes concrete. Because India's realised return has been *below* the EM average over both five and ten years (8.03% vs 9.19%, 10y), funding an India overweight from the EM sleeve itself is close to circular: you are selling EM to buy a below-EM performer. The only intra-EM funding sources that survive scrutiny are ones where you hold a *specific* view that the funded market's run is over — a live view on Taiwan's AI cycle, or a valuation view that Korea's discount will not close. Absent such a view, the developed-market or cash

funding source is not merely the most defensible; it is the only one that does not embed a losing relative bet by construction. This is why the report treats "where does the dollar come from" as co-equal with "how big is the position," rather than as a mechanical residual after sizing.

Implementation route: index, active, or single-stock

The vehicle is the last thing that separates the bridge return from the dollar in the account, and it interacts with everything above. The **index route** is cheap and liquid but buys the 52.9%-concentrated top ten and the full ~17% volatility; it is the right route for a benchmark-aware manager expressing a small deviation. The **active route** carries a steep hurdle – 75.0% of large-cap India active funds underperformed over one year, 76.3% over ten – so paying active fees is a bet against three-quarters odds and is hard to justify as a default. The **single-stock route** is the only one that escapes the index's concentration and lets an investor own the specific compounders that carry India's earnings story rather than the aggregate; it is available in practice only to the unconstrained pool with the horizon to hold through drawdowns and the mandate to be selective. The report's "selectively investable" verdict resolves here: India rewards the investor who can choose, and penalises the one forced to buy the whole concentrated index at a full multiple.

One canonical falsification dashboard

A falsification dashboard only works if there is exactly one of it: divergent lists with different counts, inverted labels and inconsistent thresholds are worse than none. That is exactly the false-confidence the report is built to purge. Here is the single reconciled dashboard. Each trigger has a threshold and a direction, and the view moves when a threshold is crossed, not before.

Bull triggers – cross these and India moves toward overweight:

1. **Valuation:** IIMA India CAPE de-rates from 36 toward its long-run mean (~24) – roughly a one-third compression – *while ROE holds*. A cheaper multiple on a stable earnings base is the cleanest bull signal, because it repairs the one leg (the multiple) that has added nothing.
2. **Earnings durability:** two or more consecutive years of double-digit Nifty-50 consolidated EPS growth *with* profit/GDP holding above 5%. This confirms the FY26 all-time-high profit share is a durable level shift rather than a cyclical peak.
3. **Currency:** the rupee's annualised depreciation durably narrows below ~2.5%/yr, or the hedge cost falls decisively below the realised drag – either of which lifts the –3.8 leg toward neutral and lets more of the earnings engine reach the dollar.

Bear triggers – cross these and India moves toward underweight:

1. **Profit reversion:** Nifty-500 profit/GDP falls back below ~4.5%, signalling the post-2020 profit surge was cyclical normalisation after all and the elevated base was not structural.
2. **Currency acceleration:** the rupee's annualised depreciation accelerates past ~5%/yr, or realised FX volatility breaks durably above its ~8% norm toward the mid-teens of its EM peers – either of which turns an orderly toll into a portfolio-wrecking drag.
3. **Multiple and earnings both roll:** CAPE stays elevated (near or above its current ~1.5× mean) while EPS growth slows to mid-single digits – the worst combination, a rich price no longer supported by the earnings that justified it.

The dashboard's discipline is symmetry: the same variables (valuation, earnings/profit-share, currency) drive both the bull and the bear, and the view is a function of *which thresholds are crossed*, not of narrative. As of the report date, none of the bull triggers has fired and none of the bear triggers has fired – which is precisely why the view is "hold near benchmark," not a directional call.

India is **structurally attractive**. The earnings engine is real and among the strongest in the emerging-market set (+9.7%/yr over twenty years), the domestic buyer base is the deepest in EM (DII buying outran FPI selling 4.6× in FY25), and the economy is the world's fifth-largest with a long convergence runway. None of that is in dispute, and none of it is the reason to be cautious.

India is **selectively investable**. Half the listed market is promoter-held and not for sale; the investable index is 52.9% concentrated in ten names; the informal economy that employs most of the country never reaches a public earnings line; and foreign *custody* is not the same as foreign *ownable* room. The economy is far larger and broader than the slice a foreign dollar can actually buy. The index is the least efficient way to own India; the unconstrained investor who can be selective captures the story better than the benchmark-bound one who cannot.

India is **valuation-sensitive**. The multiple added nothing over twenty years, CAPE sits at 36 against a mean near 24, and profit/GDP is at a fresh all-time high. The forward return the bridge implies is *below* its historical 7.1% unless earnings do more of the work, because from here the multiple is more likely to subtract than to add. Price is the swing factor, and the current price is full.

India is **more robust on currency than assumed**. The rupee depreciates ~3.1%/yr, but at 8.1% volatility it is the most orderly depreciator in the EM set, hedgeable at ~2.8%/yr – close to a wash against the drag. The currency is a manageable, quantified headwind that must be sized deliberately, not a disqualifier.

And India is **less investable through the index than the economy implies**. A dollar captures part of India's rise – the earnings part, net of currency and friction – but not the GDP number that dominates the narrative, and less of even the investable part than the headlines suggest.

The resulting position is a sizing decision, not a rating: **hold India at or near its benchmark weight; earn the overweight through a clear improvement in price or a confirmed earnings-durability signal rather than assuming it; fund any overweight from developed-market or cash sleeves rather than from Taiwan or Korea; size to ~17% volatility and a below-EM expected return; and make the currency-hedge decision explicitly**. This is nuance, not vagueness. It commits to a weight, a funding source, a set of falsifiable triggers, and a price at which the view changes. What it refuses to do is convert a market that trails its own benchmark, at a rich multiple, with a genuine earnings engine, into a single word the evidence does not support.

APPENDICES

Appendices & Company Compendium

These appendices are the working underside of the report. Nothing here is written to persuade; everything is written to be checked.

GREEN • AMBER • RED

EVERY FIGURE CARRIES ITS CONFIDENCE CLASS

These appendices are the working underside of the report. Nothing here is written to persuade; everything is written to be checked. Where a number is measured, its source and confidence class travel with it. Where a number is a judgment, it is labelled a judgment. Where the data does not exist, that gap is named rather than filled. A reader who distrusts the main text should be able to reconstruct, or refute, every load-bearing claim from what follows.

APPENDIX A — METHODOLOGY

A.1 The source traffic-light

Every figure in this report carries a confidence class, and the class is assigned before the sentence around it is written – not after, and never to fit the argument. The discipline exists because the characteristic failure of confident research prose is a sentence written at a certainty its underlying joints cannot carry: green anchors, amber joints, and headlines pitched at the certainty of the anchors. This report forbids that by construction. No headline may exceed the confidence class of the number beneath it.

- **GREEN** – a primary figure taken directly from the institution that produces it, current to the stated date, and reproducible by a reader with access to the same public document. MSCI factsheet country weights (31-Jul-2026), NSE ownership tracker percentages, SEBI issuance bulletins, RBI reserve and rate series, and MoSPI GDP prints sit here. A GREEN number can anchor a headline.
- **AMBER** – a figure that is real and sourced but carries one unresolved step: it is derived from a primary series rather than published as such (the Return Bridge legs, the currency-volatility statistics), or it is a single-source primary figure not yet corroborated (the Motilal Oswal profit/ GDP series), or it comes from a credible secondary report (media citing Prime Database on issuance) rather than the primary filing. AMBER numbers may support an argument; they may not headline one alone.
- **RED** – a figure the report wanted and could not obtain to standard: direct return-on-equity for the MSCI India universe, and a properly free-float-and-foreign-inclusion-factor-adjusted foreign-ownable ceiling. Where a RED cell exists, the report states the gap plainly and declines to substitute an identity-derived proxy dressed as the real thing. A RED item can never carry weight; it can only be flagged.

The rule that ends number drift is mechanical: a single canonical table (Appendix B) is the only source every exhibit, sentence, summary line and website figure draws from. No number appears in two places with two values. If the canonical table is wrong, everything is wrong in the same direction, which is the correct property for an audit trail to have.

A.2 The four evidence types

Claims are tagged by *what kind of thing they are*, because conflating them is how false precision enters. Four types are used throughout.

1. **FACT** – a measured quantity with a source and a date. "India's MSCI EM weight is 11.66% at 31-Jul-2026." It can be checked against a document.
2. **ESTIMATE** – a quantity built by a stated calculation from facts, where the method is disclosed and the inputs are FACTs. "The single-universe 20-year USD return decomposes to roughly +7.1%/yr." An estimate is only as good as its disclosed method, so the method is always shown.
3. **EMPIRICAL** – a statistical relationship computed from a dataset, reported with its sample size and uncertainty. The cross-country growth-versus-return scatter is EMPIRICAL: it is reported with n, slope, and the honest conclusion that it supports no statistically reliable positive relationship – not a tidy negative one either.

4. **INTERPRETATION** – a judgment the analyst makes, explicitly labelled as such and rated for how well the evidence supports it. "The forward asymmetry is unattractive" is an INTERP. The premiumisation read of the consumer is an INTERP rated well-supported. Labelling judgment as judgment is what lets a reader disagree with a conclusion while still trusting the facts.

The taxonomy's whole purpose is to stop one statistic doing two jobs. A level is not a change; a custody figure is not an ownable ceiling; a self-assigned tick is not a measurement. Where a single number tries to migrate across these boundaries, the report recomputes the right object for each stage.

A.3 The Return Bridge – single-universe method

The Return Bridge decomposes long-run equity return into the economic legs that produced it: earnings growth, re-rating (the change in the price-to-earnings multiple), dividends, and the currency translation from rupees to dollars. The identity is standard –

$$\text{USD total return} \approx \text{EPS growth} + \text{P/E change} + \text{dividend yield} - \text{currency depreciation}$$

– but it is fragile to one specific abuse: mixing universes. If the EPS leg is computed on one index and the price leg on another, the residual absorbs the mismatch and can be mistaken for a missing economic factor. This report builds the bridge on one universe end to end (Nifty to Nifty). Rebuilt this way, the 20-year decomposition reads **+9.7 EPS – 0.1 P/E + 1.3 dividend – 3.8 FX ≈ 7.1% USD**, and the published Nifty TRI reconciles to within roughly 1.8 percentage points. That residual is explained by anchor-and-endpoint mismatch – the earnings series runs on December year-ends while the report closes at 31-Jul-2026 – and not by any unmodelled driver. Mixed-universe bridges can produce an alarming "residual equals the no-re-rating leg" coincidence; it dissolves once the universe is held constant – an artifact of the join, not a discovery.

The standalone-to-consolidated EPS break (April 2021) caveat. This is the single most important footnote in the entire quantitative apparatus, and it is flagged at every use. In April 2021, NSE switched the Nifty earnings basis from **standalone to consolidated** accounts. The effect is visible and large in the valuation history: derived EPS jumps from ~₹364 (Dec-2020, standalone) to ~₹720 (Dec-2021, consolidated), and trailing P/E drops from 38.5 to 24.1 across the same break – a discontinuity driven by a definitional change, not by earnings doubling in a year. Any EPS-growth or P/E-change leg that *spans* April 2021 is contaminated unless it is computed on a consistent basis. The consequence for this report: the 20-year and 10-year bridges are treated as robust because the break is a single identifiable step inside them and can be reasoned around; the **5-year window is treated as the least robust of the three** and is never headlined. It begins at a December-2021 local P/E peak of 23.7, which makes it start-point-sensitive independent of the accounting break. The report states the 5-year USD figure (~3.8%) and immediately labels it the weakest leg in the exhibit.

A.4 The multi-stage reverse-DCF

A reverse discounted-cash-flow model asks not "what is the stock worth" but "what does today's price already assume." A single-stage version of the model produces a characteristic defect: an implied perpetual growth rate (~11.8%) that, divided by the assumed retention, implies a return-on-equity near 16.9% – a figure that then sits unreconciled against the ~14% observed ROE. A single-stage model forces a company's entire future into one growth number and one payout, which is where the error originates.

This report uses a multi-stage structure instead: an explicit high-growth phase, a fade phase in which growth and profitability decay toward a terminal state, and a terminal phase anchored to a defensible long-run nominal growth and cost of equity. The output is not a point estimate of "fair value" but a statement of *embedded expectations*: the

growth-and-return path the current index price requires in order to justify itself. The model's central finding is a tension it makes explicit rather than hiding – **the price embeds a rising ROE**. The implied return profile is only reconcilable with observed profitability if margins and returns hold at or above their already-elevated level. That is the core asymmetry the report rests on, and the multi-stage model surfaces it as the answer rather than tripping over it as a contradiction. Scenarios for 2035 are rebuilt on this model, not on a hard-coded "residual stays constant" assumption.

A.5 The DRI / DPI framework – end-market versus profit geography

"How much of this company is India" is two questions, and the standard shortcut collapses them, deriving a company's India-share as 100% minus its export ratio – which counts a refinery's domestic output, labour, profit and tax as "not India" merely because the barrels sail abroad. This report splits the question into two measures that are allowed to disagree.

- **Domestic Revenue Intensity (DRI)** – the share of a company's *end-market revenue* earned from customers located in India. This is a demand-geography measure. A software exporter has low DRI; a private bank has ~100%.
- **Domestic Profit Intensity (DPI)** – where the *profit, production, cost base, assets and tax* actually sit. This is a value-geography measure, and it can diverge sharply from DRI. A company can sell abroad (low DRI) while producing, employing and paying tax at home (high DPI); or sell at home (high DRI) while exporting the profit to a foreign parent through royalties and dividends.

The two measures together answer the question an allocator actually has – not "where are the customers" but "where does the value land, and can I own the entity that captures it." The framework requires that every DRI/DPI figure state its universe, its weighting and its method, and that revenue geography and profit geography are never silently swapped. A worked correction: the widely repeated "Airtel 43/57" split confuses an EBITDA ratio with a revenue ratio; corrected, Airtel's India end-market share is ~77%.

A.6 Honest limitations

Four gaps are load-bearing enough to name here rather than bury.

1. **The direct-ROE gap (RED).** The report could not obtain a clean, direct return-on-equity for the MSCI India universe from a primary provider this cycle. It therefore does *not* assert ROE parity with EM to the basis point (the oft-quoted "13.9 versus 13.8" parity is a P/B-over-P/E identity masquerading as a measurement). Where profitability is discussed, the report states it as a P/B-plus-growth-persistence premium and flags the direct number as unobtained.
2. **The FIF-adjusted foreign-ownable ceiling (RED/AMBER).** Foreign *holdings* (an NSDL custody figure) are not the same as the foreign-*ownable* maximum, which requires each stock's free float and MSCI foreign-inclusion factor. The report builds the ownership ladder as far as the public data allows and explicitly declines to present a holdings figure as a ceiling.
3. **Single-source profit dependence (AMBER).** The entire profit-to-GDP architecture traces to one Motilal Oswal series. A second source is used for corroboration where available, but the dependence is real and is flagged at every use of the 5.2%/5.7% figures rather than assumed away.
4. **The correlation-matrix gap.** The report makes an allocation argument but does not present a full cross-asset or cross-country correlation matrix for India within an EM or global portfolio. Diversification claims are therefore made narrowly (return, volatility, currency behaviour) and the report does not claim to have measured India's marginal contribution to a specific multi-asset portfolio's risk. That is a genuine boundary of the work.

APPENDIX B — DATA

The canonical table below is the single source of truth (audit §5). Every exhibit and sentence in the report draws from it; no figure appears elsewhere with a different value. Data to 31-Jul-2026 unless noted. Class per Appendix A.1.

B.1 Canonical numbers

QUANTITY	CANONICAL VALUE	SOURCE	CLASS
India MSCI EM weight (now / 2024 peak / 2020)	11.66% / ~20% / 9.2%	MSCI EM factsheet 31-Jul-26; media/factsheet history	GREEN / AMBER / AMBER
Taiwan / China / Korea MSCI EM weight	26.63 / 21.38 / 20.33%	MSCI EM factsheet 31-Jul-26	GREEN
India MSCI EM ex-China weight	14.84%	MSCI EM ex-China factsheet	GREEN
India global market-cap rank / value	5th / ~\$5.05tn	exchange/media aggregate	AMBER
MSCI India 10y USD net return / MSCI EM	8.03% / 9.19% (India below EM 5y & 10y)	MSCI index returns	GREEN
Nifty single-universe USD return (20y/10y/5y)	~7.1% / ~9.5% / ~3.8% (5y least robust)	derived, Return Bridge	AMBER (from GREEN)
Return Bridge 20y legs (EPS / P / E / div / FX)	+9.7 / -0.1 / +1.3 / -3.8 %/yr	derived	AMBER
India nominal GDP FY25	₹330.7 lakh cr ≈ \$3.9tn (not \$4.1tn)	MoSPI	GREEN
Profit/GDP FY26 (Nifty-500 / broad listed)	5.2% (all-time high) / 5.7%	Motilal Oswal 11-Jun-26	GREEN (single-source — flag)
Corporate profit vs GDP FY26	profit +15.6% vs nominal GDP +8.9%	Motilal Oswal / MoSPI	GREEN
Four Indias (Econ/Corp/Listed/FreeFloat/Foreign-held)	\$3.9tn / ~\$0.20tn / \$5.35tn / ~\$2.67tn / ~\$0.84-0.93tn	composite	GREEN + AMBER legs
Promoter / FPI / DII / retail ownership (NSE-listed)	50.0 / 15.8 / 19.6 / 18.7%	NSE ownership tracker (Q4FY25)	GREEN
India valuation: P/E / P/B (Nifty) / CAPE / div yld	20.5-23.9 / 2.95 / 36.0 (mean ~24) / 1.27%	NSE/IIMA	GREEN
India ROE (direct) vs derived	direct = GAP; derived ≈14.4% ≈ EM ~13.8%	— / identity	RED / AMBER
Reverse-DCF implied growth (single-stage, to be replaced)	~11.8%	derived	AMBER
INR depreciation 10y / volatility / hedge cost	3.11%/yr / 8.12% (lowest-vol depreciator) / ~2.8%	derived; RBI-corroborated	AMBER
MSCI India vol (3y) / Sharpe / max drawdown	17.3% / 0.12 / -72.6%	MSCI	GREEN
DII vs FPI FY25 / SIP (Jun-26)	+₹6.07tn vs -₹1.31tn (≈4.6×) / ₹31,781 cr/mo	exchange / AMFI	GREEN / AMBER
	₹37.3 / ~57 / 24.8 / 16 lakh cr	AMFI / regulator disclosures	AMBER/GREEN

QUANTITY	CANONICAL VALUE	SOURCE	CLASS
Domestic pools (MF equity / LIC / EPFO / NPS)			
Equity supply FY26 / FY25 (IPO+QIP+OFS)	₹3.05 / 3.71 lakh cr	Prime Database / SEBI	AMBER
SPIVA India active underperformance (1y/10y large- cap)	75.0% / 76.3%	S&P DJI SPIVA	GREEN
Nifty top-10 concentration	52.9%	NSE	GREEN
State income gap (Karnataka / Bihar per-capita)	~5.6× (₹4.22L / ₹0.76L)	MoSPI/state accounts	GREEN

B.2 Return Bridge legs (single-universe, Nifty→Nifty, USD)

WIN-DOW	EPS GROWTH	P/E CHANGE	DI-VIDEND	FX (INR DEP.)	≈ USD RE-TURN	ROBUSTNESS
20-year	+9.7	-0.1	+1.3	-3.8	~7.1%	strong (break is one step inside)
10-year	—	—	—	-3.1	~9.5%	strong
5-year	—	—	—	—	~3.8%	weakest — starts at Dec-21 P/E peak (23.7)

The Apr-2021 standalone→consolidated EPS break sits inside every window; the 20y/10y are reasoned around it, the 5y is not headlined. Reconciliation to published Nifty TRI within ~1.8pp is endpoint mismatch (Dec year-end vs 31-Jul-2026), not a missing factor.

B.3 Currency statistics (10-year, annualised; derived, RBI-corroborated)

CURRENCY	DEPRECIATION / YR	VOLATILITY (σ OF ANNUAL CHANGE)	WORST YEAR	NOTE
INR	3.11%	8.12%	2008 (+23.3%)	lowest-vol depreciator ex-managed pegs
CNY	0.77%	4.76%	2022	managed
VND	1.39%	2.54%	2011	managed
Korea (KRW)	2.07%	11.41%	2008 (+41.2%)	
Mexico (MXN)	0.34%	11.57%	2008	
Brazil (BRL)	3.30%	19.12%	2015 (+49.3%)	free-floating, high vol
South Africa (ZAR)	0.61%	15.70%	2008	

The rupee's signature: a persistent, low-volatility depreciator. 12-month forward premium ≈ 2.8%/yr (Aug-2026) is the hedge cost. This retires the "bottom-tier rupee" framing.

B.4 Domestic pools and flows

POOL / FLOW	VALUE	CLASS
Equity mutual-fund AUM	~₹37.3 lakh cr	AMBER

POOL / FLOW	VALUE	CLASS
LIC total assets	~₹57 lakh cr	GREEN
EPFO corpus (~9.5% in ETFs)	~₹24.8 lakh cr	AMBER
NPS + APY	~₹16 lakh cr	AMBER
SIP run-rate (Jun-2026)	₹31,781 cr/month	AMBER
DII net buy CY2025 (record)	+₹7.88 lakh cr	GREEN
FPI net CY2025 (record outflow)	-₹1.66 lakh cr (~\$18.9bn)	GREEN
DII net buy CY2024	+₹5.27 lakh cr	GREEN
Equity supply FY25 (record)	₹3.71 lakh cr	AMBER
Equity supply FY26	₹3.05 lakh cr	AMBER

The SIP run-rate is a slice of the domestic pool, not its whole; equity supply (IPO+QIP+OFS) is a real sponge that partly off-sets the domestic bid, which is why "floor and ceiling from one savings pool" is the correct framing.

B.5 Valuation history (Nifty-50, year-end; earnings basis break flagged)

DATE	CLOSE	P/E (TRAILING)	P/B	DIV YLD	EPS (DERIVED)	BASIS
2019-12-31	12,168	28.30	3.75	1.24	430.0	standalone
2020-12-31	13,982	38.45	3.96	1.14	363.6	standalone
2021-12-31	17,354	24.11	4.37	1.19	719.8	consolidated ← break
2023-12-29	21,731	23.17	3.81	1.28	937.9	consolidated
2025-12-31	26,130	22.75	3.55	1.28	1,148.6	consolidated
2026-07-31	24,384	20.78	2.99	1.22	1,173.4	consolidated

The EPS jump 363.6→719.8 across 2021 is the accounting basis change, not a doubling of earnings – the reason the P/E leg of any bridge spanning it is contaminated. Note the divergence at 31-Jul-2026: P/E 20.8 sits ~12% below its 10-year average while CAPE (36.0) sits ~1.5× its long-run mean, because the 2021 reset lowered headline P/E. The report shows both rather than choosing.

APPENDIX C – BIBLIOGRAPHY & RESEARCH LITERATURE

This is a verifiable list, split into (I) primary data and institutional sources actually used to build the exhibits, and (II) the academic foundation whose findings frame the argument. Primary sources are the ones a reader checks the numbers against; academic sources are the ones a reader checks the reasoning against.

Part I – Primary data & institutional sources

Index and benchmark providers. - MSCI Inc. – MSCI Emerging Markets, MSCI EM ex-China, and MSCI India factsheets (country weights, index net returns in USD, annualised volatility, Sharpe, and drawdown), 31-Jul-2026 vintage. *Primary.* The single most load-bearing external source: India's benchmark weight, its EM-relative return, and its risk statistics all originate here. - NSE Indices Ltd (niftyindices.com) – Nifty-50 and Nifty-500 closing levels, trailing P/E, P/B and dividend-yield history; Nifty Total Return Index; top-10 concentration; sectoral composition. *Primary.*

Note the Apr-2021 standalone→consolidated earnings-basis change originates in NSE methodology and is documented in their factsheets. - S&P Dow Jones Indices – SPIVA India Scorecard. Active-versus-passive underperformance rates for Indian large-cap funds over 1- and 10-year horizons (75.0% / 76.3%). Primary. Used to price the cost of the "just pick stocks" alternative. - IIM Ahmedabad – India CAPE (cyclically-adjusted P/E) series (Prof. and research-cell publication of the Nifty-500 CAPE). Jul-2026 reading ~36.0 against a long-run mean near 24. Primary/academic hybrid – an academically-constructed index maintained institutionally.

Market microstructure, ownership and flows. - **National Securities Depository Ltd (NSDL)** – FPI custody/AUC data and foreign-holding levels. *Primary. Used with the explicit caveat that custody ≠ foreign-ownable ceiling.* - **NSE Market Pulse & NSE ownership tracker** – quarterly promoter/FPI/ DII/retail ownership shares for the NSE-listed universe; category-wise net equity flows (FPI, DII, individuals) by calendar and fiscal year. *Primary.* - **Securities and Exchange Board of India (SEBI)** – monthly bulletins on equity mobilisation (IPO, QIP/IPP, preferential allotment, rights issues); FPI aggregate flow statements. *Primary.* - **Association of Mutual Funds in India (AMFI)** – equity mutual-fund AUM and monthly SIP inflow run-rate. *Primary.* - **Prime Database** (via secondary reporting, Business Standard / Tribune) – IPO and public-equity issuance totals FY21–FY26. *Secondary-primary; classed AMBER because accessed through media citation rather than the underlying subscription database.*

Macro and monetary. - **Reserve Bank of India (RBI)** – USD/INR reference rates and history; foreign-exchange reserves; policy rates; the household financial-savings series (correctly denominated as % of GNDI, not GDP); listed private-corporate sales sample. *Primary.* - **Ministry of Statistics & Programme Implementation (MoSPI)** – nominal and real GDP, the FY25/FY26 prints and the base revision; state per-capita income (Karnataka/Bihar gap). *Primary.* - **Motilal Oswal (institutional research)** – the corporate-profit-to-GDP series (Nifty-500 5.2% / broad listed 5.7% in FY26; profit +15.6% vs nominal GDP +8.9%). *Primary but single-source – flagged AMBER at every use; corroboration sought against RBI and NSE aggregates where possible.* - **World Bank – World Development Indicators (WDI)** and **IMF – World Economic Outlook / IFS.** Cross-country GDP-growth and GDP-level comparators used in the growth-versus-return scatter and the GDP base reconciliation. *Primary. The IMF CY2025 ~\$4.1tn GDP is a different base/period from the MoSPI FY25 ₹330.7 lakh cr ≈ \$3.9tn figure; the two are never mixed inside one funnel.* - **Schroders – Emerging Markets Lens** (quarterly EM valuation and profitability chartbook). Used as a cross-check on EM-relative valuation and profitability, and cited as the intended external source for a direct ROE that could not be obtained to standard this cycle. *Primary/institutional.* - Currency peer statistics (Brazil, Mexico, Korea, Taiwan, South Africa, Indonesia, China, Vietnam) derived from central-bank and market spot histories, corroborated against RBI for the rupee.

Part II – Academic foundation (the underlying literature themes)

These are the findings that frame the argument; the report's contribution is empirical India-specific measurement, not new theory. Author, year and role are given so each can be located and weighed.

- **Jay R. Ritter (2005, 2012), "Economic Growth and Equity Returns."** *Foundational.* The empirical result that cross-country economic growth is uncorrelated – even weakly negatively correlated – with real equity returns. This is the intellectual spine of the entire report: growth rank ≠ return rank. The report's own n=13 scatter is a fresh test of Ritter's claim on current EM data, and it survives.
- **Elroy Dimson, Paul Marsh & Mike Staunton – Triumph of the Optimists (2002) and the annual Credit Suisse/UBS Global Investment Returns Yearbook.** *Foundational.* The long-run cross-country evidence on equity returns, dilution, and the wedge between GDP growth and per-share value. Underpins the treatment of net issuance/dilution as a leg of the Return Bridge.

- **Marianne Bertrand, Paras Mehta & Sendhil Mullainathan (2002), "Ferretting out Tunneling."** *Mechanism*. The canonical evidence on value diversion within Indian business groups. Frames the promoter-control sections of the compendium: ownership concentration is not automatically alignment; it is also a channel through which minority value can leak.
- **Dani Rodrik (2016), "Premature Deindustrialization."** *Structural*. Developing economies peaking in manufacturing employment at lower income levels than their predecessors. Frames the scepticism toward a straight-line "India as the next China" manufacturing-led listed-earnings story, and the China+1 section of the compendium.
- **Raghbendra Jha / Nirvikar Singh and the crowding-in literature; Bahal, Raissi & Tulin (2018, IMF WP), "Crowding-Out or Crowding-In? Public and Private Investment in India."** *Policy*. The empirical finding that Indian public investment has, in the relevant period, crowded *in* private investment – the evidence base for taking the public-capex ramp seriously as a macro driver while still separating it from a listed-profit claim.
- **Barry Eichengreen & Poonam Gupta (2017, and related RBI/NIPFP work), on inflation targeting in India.** *Monetary*. Frames the treatment of the rupee as a managed, low-volatility depreciator: the currency's behaviour is a policy outcome, not an accident, which is why its volatility is an EM outlier.
- **Aswath Damodaran (NYU Stern, ongoing), country risk premia, cost-of-equity and ERP datasets.** *Valuation method*. The reference architecture for the reverse-DCF's discount-rate and equity-risk-premium assumptions, and a candidate external source for direct ROE.
- **Robert J. Shiller (Yale), CAPE methodology.** *Valuation method*. The cyclically-adjusted P/E construct that the IIMA India CAPE applies to the Nifty-500; the basis for reading 36.0 against a ~24 mean as a mean-reversion headwind.
- **Franklin Allen, Rajesh Chakrabarti, Sankar De et al. – financial-development and market-structure literature on India.** *Structural*. Context for the formalisation-and-financialisation channel (the funnel from informal activity to foreign-ownable listed profit).
- **Viral Acharya and co-authors – Indian banking, credit cycles and financial-sector fragility.** *Sectoral*. Frames the credit-cost normalisation that drove part of the FY20→FY26 profit surge and the treatment of financials in the compendium.

Primary-versus-academic note: every number in the exhibits comes from Part I; every framing in the prose is anchored to Part II. Where the report makes a claim the literature does not license (e.g., a specific 2035 path), it is labelled INTERPRETATION, not evidence.

APPENDIX D – COMPANY COMPENDIUM

Twenty-one companies, each chosen to make one mechanism visible. Ownership percentages are real, from exchange filings for the quarter ended 30-Jun-2026 (promoter / FPI / DII / public; classed AMBER – exchange-filed but not independently re-audited this cycle). India-revenue shares, where stated, are the report's DRI figures. The organising question throughout is not "is this a good company" but "where does the value land, and can a foreign allocator own the entity that captures it." Structure is deliberately varied; there is no house template.

Group 1 – Domestic compounders: the value lands where you can buy it

These are the cleanest case for owning India: revenue, profit, production and tax all sit at home, and the float is genuinely investable. The mechanism they illustrate is the one the whole report is looking for – high Domestic Revenue Intensity and high Domestic Profit Intensity – and the catch is not geography but price and float composition.

HDFC Bank – the reference domestic compounder. Ownership tells the story before the P&L does: **promoter 0.00%, FPI 41.82%, DII 41.75%, public 16.27%**. There is no promoter; the bank is a widely-held institution, and foreign and domestic institutions hold it almost exactly evenly. India-revenue is ~100% – a rupee-in, rupee-out lender whose profit, deposits, credit book and tax are entirely domestic. This is DRI and DPI both near their maximum. Who captures the profit? Minority shareholders, in proportion, with no controlling family to share it with – which is precisely why the free float is so large and so foreign-accessible. The takeaway: this is what an unambiguous claim on Indian domestic growth looks like, and the only real question an allocator faces is valuation and the FPI-crowding risk, since a 42% foreign holding is itself a source of flow volatility when the benchmark weight moves.

ICICI Bank – the same mechanism, tilted domestic. Promoter 0.00%, **FPI 33.79%, DII 42.32%**, public 23.68%. Note the difference from HDFC Bank: here domestic institutions out-hold foreigners by nearly nine points, and retail/public ownership is higher. India-revenue ~100%. The mechanism is identical – a fully domestic profit pool with no promoter to share it – but the ownership mix makes it less hostage to foreign flow. When FPIs sold India through 2025–26, banks like this were cushioned by exactly the DII wall the flows section documents. Takeaway: among the domestic compounders, the ones with a thicker DII base are structurally steadier holdings for a foreign investor who worries about being in a crowded foreign trade.

Bajaj Finance – promoter-backed, but the profit still lands at home. Here the structure changes: **promoter 54.68%**, FPI 20.20%, DII 16.30%, public 8.64%. A controlling group owns just over half, so the free float is roughly the other half – smaller than the zero-promoter banks. But the group is aligned with minorities in the ordinary way (dividends and capital appreciation, not tunnelling), and India-revenue is ~100%. The value lands at home and is shared pro-rata. Takeaway: a high promoter stake is not a red flag by itself; it shrinks the float and concentrates control, and the question is always whether control is used to build value or to divert it. Here the record reads as the former – which is why the market pays a premium multiple that itself becomes the risk.

Titan – the Tata-and-state hybrid, pure domestic demand. Promoter **52.90%** (a Tata-group and Tamil Nadu state-investment combination), FPI 15.39%, DII 15.15%, public 16.31%. India-revenue ~100%: jewellery and watches sold to the top of the Indian income distribution. This is the premiumisation thesis in a single name – its revenue base *is* the affluent cohort at the top of the K, which is why it has out-earned mass-market consumer categories. Who captures the profit: the Tata group and the state as co-promoters, and minorities alongside them. Takeaway: Titan is a clean way to own Indian premium consumption, and it carries the premiumisation-thesis risk directly – it thrives if the top of the K keeps separating from the bottom, and disappoints if mass incomes finally catch up and the growth broadens away from premium.

Avenue Supermarts (DMart) – high control, thin float, domestic to the core. **Promoter 74.48%** – one of the tightest promoter holdings in the large-cap universe – FPI 9.22%, DII 8.67%, public 7.55%. India-revenue ~100%, value retail to the mass and middle market. The mechanism worth noting is the float: at ~25% free float, foreign and domestic institutions are competing for a small ownable slice of a fully domestic profit pool, which is a recipe for a persistently rich multiple regardless of the underlying return on capital. Takeaway: scarcity of float is itself a valuation factor, and DMart shows how a genuinely domestic, well-run business can trade expensive not because growth is mispriced but because there is simply not enough of it to buy.

Eternal (formerly Zomato) – new-economy, no promoter, institutionally held. Promoter 0.00%, **FPI 29.09%, DII 39.21%**, public 26.99%. A professionally-managed, widely-held platform business with no controlling family – closer in ownership shape to a US technology listing than to a traditional Indian promoter company. India-revenue ~100%: food delivery and quick-commerce to Indian consumers. The takeaway is about ownership evolution as much

as the business: the newest cohort of large Indian companies is arriving *without* promoters, which over time lifts the free-float share of the market and makes more of India's growth genuinely foreign-ownable. This is a structural offset to the promoter-concentration story that dominates the older index.

Hindustan Unilever – sells to India, banks the profit in London. This one is placed among the domestic compounders to make a point by contrast. **Promoter 61.90% – but the promoter is Unilever plc, a foreign parent.** FPI 9.50%, DII 16.92%, public 11.60%. India-revenue is ~100%: HUL is one of the most domestically-facing revenue lines in the entire market, soaps and staples sold to Indian households. So DRI is maximal. But DPI is where the framework earns its keep: a meaningful share of the economic value flows *out* – to the foreign parent via royalty and brand fees and via majority dividends. The end-market is India; a slice of the profit geography is the Netherlands and the UK. Takeaway: a company can be 100% "India revenue" and still route part of its value abroad by ownership. HUL is why DRI and DPI must be reported separately – read either one alone and you misjudge who is actually getting paid.

Group 2 – Global earners: Indian-listed, foreign-facing revenue

Here the framework inverts. These companies are listed in Mumbai and often controlled by Indian groups, but their customers – and a large share of their profit – sit abroad. Owning them is owning a claim on global IT budgets, US generic-drug pricing, or European car demand, denominated in rupees. Low DRI; the DPI question is subtler, because the production and jobs are often still Indian even when the customer is not.

COMPANY	PROMOTER	FPI	DII	PUBLIC	INDIA REVENUE (DRI)	WHAT YOU ARE REALLY BUYING
TCS	71.77% (Tata Sons)	9.07%	13.41%	5.69%	~5.9%	Global IT services demand, billed mostly in USD
Infosys	13.82%	27.09%	42.78%	15.88%	~2.9%	Same, with almost no promoter and a huge float
Sun Pharma	54.48%	14.54%	22.06%	8.81%	~32.5%	US + EM generics/specialty, one-third India
Dr Reddy's	26.63%	20.67%	31.75%	20.76%	~16.5%	US generics primarily; India a minority
Tata Motors (JLR)	42.51%	17.11%	16.97%	23.09%	~28%	Global luxury autos (JLR) + Indian CV/PV

TCS is the archetype: a ~72% Tata-controlled float so tight that the ownable slice is small, wrapped around a business that earns barely a twentieth of its revenue in India. The rupee here is a *tailwind*, not a drag – costs are Indian, revenues are dollar – which flips the currency leg of the return bridge for this name relative to the domestic compounders. **Infosys** makes the same bet with the opposite ownership: only 13.82% promoter, 42.78% DII and 27.09% FPI, one of the most institutionally-owned and freely-floating large caps in India, and the *lowest* India-revenue share (2.9%) in this table. If you own Infosys you are, in substance, a rupee-denominated holder of Western enterprise IT spending. The DPI nuance: the profit accrues to Indian-listed shareholders and the jobs are Indian, so the *production* geography is domestic even though the *demand* geography is not – a case where DRI is low but the value still lands partly at home.

Sun Pharma and Dr Reddy's show the same mechanism in pharma with different intensities: Sun keeps about a third of revenue in India, Dr Reddy's about a sixth, and the balance is US and other-market generics – a business whose margins are set in Washington and by US generic price deflation, not in Delhi. Note Dr Reddy's ownership: a modest 26.63% promoter and a combined ~52% institutional holding make it one of the more foreign-accessible

names here. **Tata Motors** is the most vivid: ~28% India revenue, and the swing factor is Jaguar Land Rover – European and Chinese luxury-car demand booked through a Mumbai listing. Takeaway for the group: these are the names that make "India revenue" and "Indian equity" diverge most, and they are the reason a top-down "buy India's domestic growth" thesis and a bottom-up index purchase are not the same trade. A large chunk of the index earns its living outside the country whose growth story is being sold.

Group 3 – State-owned: the government is the promoter

Where the promoter is the state, three things change at once: the float is capped by the government's stake, capital-allocation answers partly to policy rather than only to return, and the alignment between the controlling owner and a minority foreign shareholder is weakest in the whole compendium. The businesses can be excellent; the ownership structure is the risk.

State Bank of India – promoter **55.47% (Government of India)**, FPI 10.81%, DII 26.52%, public 7.06%. India-revenue ~100%; the largest lender in the country, a pure domestic-profit machine. But the controlling shareholder is the sovereign, and its objectives include financial inclusion, priority-sector lending and occasional recapitalisation of the system – objectives that do not always maximise minority return. The DII holding (26.52%) is notably high, reflecting LIC and domestic funds owning the national bank; foreign holding is correspondingly modest. Takeaway: SBI is a genuine claim on Indian credit growth at a structurally cheaper multiple than the private banks, and the discount is not a mistake – it is the market pricing the state's non-financial objectives and the dilution risk of future government-driven capital raises.

Hindustan Aeronautics and **Bharat Electronics** – the defence pair, and a study in how the state-as-promoter can be a *tailwind* when policy and profit align.

	HAL	BEL
Promoter (Govt)	71.64%	51.14%
FPI	9.34%	18.02%
DII	11.90%	21.02%
Public	7.03%	9.82%
India revenue	~100% (domestic defence)	~100% (domestic defence)

Both earn almost entirely from the Indian defence budget – a customer that is also, through the promoter stake, the controlling owner. That is an unusual alignment: indigenisation policy (the drive to build defence hardware at home rather than import it) directly grows their order books, so here the state's non-financial objective and the minority's financial interest point the same way. The float is tight, especially HAL's (~28%). Takeaway: state ownership is not uniformly a discount factor; when the government's policy agenda is the company's revenue growth, as in defence indigenisation, the state-promoter structure can be the thesis rather than the caveat. The risk is the mirror image of SBI's – order flow, pricing and even export permissions are policy-determined, so the earnings are as political as they are commercial.

Group 4 – Promoter-controlled conglomerates: control, and who it serves

Half the Indian market by value sits behind a controlling promoter. Concentration is not automatically misalignment – the tunnelling literature (Bertrand–Mehta–Mullainathan) documents that it can be, which is exactly why each case must be read on its own conduct rather than assumed good or bad. The common thread is a shrunk free float and a controlling owner whose incentives the minority must underwrite.

Reliance Industries – one ticker, several economies. Promoter 50.48%, FPI 17.19%, DII 21.10%, public 11.05%. The India-revenue figure is where naive measures go wrong and the DRI/DPI split earns its place: on a *standalone* basis Reliance earns roughly **49%** of revenue in India, because the refining-and-petrochemical export book is large – but that export revenue is produced at Jamnagar, by Indian labour, on Indian assets, paying Indian tax. So its DRI is moderate while its DPI is high: the customer is often abroad, the value is overwhelmingly at home. Layered on top are two almost-entirely-domestic businesses, Jio (telecom) and Retail, whose DRI is ~100%. Who captures the profit: the Ambani promoter group controls just over half, and the consolidated entity is really a portfolio – energy exporter, domestic telecom, domestic retail – inside one 50%-held float. Takeaway: never assign Reliance a single India-share; it is a refiner, a telco and a retailer wearing one ticker, and the honest read requires disaggregating both the revenue geography and the segment mix.

Adani Ports & SEZ – infrastructure, tight control, domestic cash flows. Promoter 66.03%, FPI 15.58%, DII 13.60%, public 4.79%. India-revenue ~100%: port throughput and logistics for Indian trade. The distinguishing feature is the ~66% promoter stake and a very thin ~34% float, the tightest control in this group. The business is a genuine domestic-infrastructure monopoly-adjacent cash generator; the ownership is the reason it trades with a governance-and-leverage risk premium that the cash flows alone would not command. Takeaway: with a two-thirds promoter holding, the minority is a small passenger in a founder-driven vehicle – the assets are real and domestic, but the terms on which minority value is protected are set almost entirely by the controlling group.

Larsen & Toubro – the professionally-run conglomerate with no promoter. Promoter 0.00%, FPI 19.13%, DII 43.01%, public 37.60% – the highest public holding in the compendium. L&T is the counter-example to the concentration story: India's flagship engineering-and-construction group is board-controlled and widely held, with domestic institutions owning 43% and no family behind it. India-revenue is high but not total (a meaningful Middle-East order book). Who captures the profit: a broad base of institutions and public shareholders, with professional management as the steward rather than a promoter. Takeaway: L&T shows that the promoter model is not the only Indian model – and that the most direct listed proxy for the domestic capex cycle happens to be one of the most cleanly-governed, widely-floated names in the market.

Mahindra & Mahindra – a light-touch promoter over a sprawling group. Promoter 18.44% – strikingly low for a group of its scale – **FPI 34.89%, DII 32.58%**, public 10.47%. The Mahindra family controls the group with well under a fifth of the equity, and foreign and domestic institutions together own two-thirds. India-revenue is high (autos and farm equipment sold largely at home). Takeaway: control and ownership are not the same quantity. An 18% stake exercising group control over a large float is common in older Indian houses, and it raises the standard governance question in reverse – the minority owns most of the economics while the promoter holds most of the control, so alignment depends on conduct and structure rather than on skin-in-the-game arithmetic.

Group 5 – China+1 and the manufacturing bet

The last group is the one the market most wants to believe and the report treats most sceptically, framed by Rodrik's premature-deindustrialisation caution. These are the names sold as India-replaces-China: electronics contract manufacturing, and the metals-and-mobility complex that a manufacturing build-out would feed. The mechanism to watch is margin – contract manufacturing captures assembly, not the high-value chip or brand – and, in metals, the fact that the profit is set by a global commodity price, not by Indian demand.

Dixon Technologies – the poster child, and the margin question. Promoter 28.55%, FPI 17.87%, DII 28.38%, public 25.20%. India-revenue ~100% by billing – Dixon assembles phones, appliances and electronics, much of it for global brands, on Indian soil. This is the single most direct listed China+1 proxy. But the DPI framework flags the catch: contract assembly captures the *thin* slice of the value chain (labour, integration), while the *thick* slice (chips, IP, brand) accrues to the foreign principal. So the revenue is domestic and the volume story is real, yet the margin is

structurally low and the multiple the market pays discounts years of the volume ramp in advance. Takeaway: owning the China+1 thesis through an assembler means owning volume growth at assembly margins at a growth multiple – a demanding combination that requires the ramp to arrive exactly as hoped.

Kaynes Technology – the same bet, one rung up the value chain. Promoter 53.46%, FPI 5.82%, DII 11.78%, public 28.93% (note the very low FPI – a name foreign investors have barely entered). Kaynes sits slightly higher in the electronics value chain than pure assembly (design-led manufacturing, and a semiconductor-packaging ambition), which is the part of the China+1 story that would actually capture margin if it succeeds. India-revenue ~100%. Takeaway: within the manufacturing bet, the distinction that matters is *where on the value chain* the profit is captured – Kaynes is a bet that an Indian firm can climb from assembly toward higher-value work, which is the version of the thesis worth paying for, and also the version most exposed to execution and capital-intensity risk.

Tata Steel – domestic listing, global commodity price. Promoter 32.94% (Tata), FPI 18.90%, DII 26.44%, public 21.52%. India-revenue is substantial but the profit is set by the *global* steel price and by its European operations, which have been a persistent drag. The mechanism: a metals producer's earnings are a leveraged bet on a world commodity cycle, only loosely coupled to Indian demand in any given year. Takeaway: the metals complex is where "India growth" and "the P&L" decouple through the commodity price rather than through geography – you can be right about Indian steel demand and still lose money if the global price and the European book move against you.

InterGlobe Aviation (IndiGo) – domestic demand, foreign-priced inputs. Promoter 41.57%, FPI 20.32%, DII 31.85%, public 6.23%. India-revenue ~100%: it is the dominant domestic airline, a direct claim on Indian mobility and the travelling top-of-the-K consumer. But two of its three biggest costs – fuel and aircraft leases – are dollar-priced, so a rupee depreciation that helps the IT exporters *hurts* IndiGo. Takeaway: IndiGo is a clean domestic-demand story with a currency mismatch buried in the cost line, and it sits in this group as the reminder that even a 100%-DRI domestic champion can have its margins set abroad – here through input costs rather than through revenue or ownership. The direction of the rupee leg depends entirely on which side of a company's P&L the dollars sit.

Closing note on the compendium. The five groups are five different answers to one question – where does the value land, and can you own the entity that captures it. The domestic compounders answer cleanly and charge you for it in the multiple. The global earners answer "abroad, mostly," which severs Indian-equity ownership from the Indian growth story. The state-owned names answer "at home, but the owner has other objectives." The promoter conglomerates answer "at home, behind a controlling family whose conduct you are underwriting." And the China+1 names answer "at home by revenue, but at assembly margins or at a global commodity price." No single India-share, no single ownership template, and no single takeaway fits all five – which is the entire reason the report refuses to headline the market with one number.



*Independent India and emerging-markets research
for global capital allocators*

INDIA & EMERGING-MARKETS STRATEGY

inquiries@kakderesearch.com

kakderesearch.com